THE EDEN PROTOCOL

Architecture for Embedded AI Alignment That Scales With Capability

Michael Darius Eastwood | Version 4.0 | March 2026 | First published 22 February 2026

Author, *Infinite Architects: Intelligence, Recursion, and the Creation of Everything* | London, United Kingdom

Executive Summary for Grant Reviewers — Complete v5 Alignment + Paper II Compute Scaling + Eden Protocol Two-Model Results

CONTEXT: THE FRAMEWORK IN BRIEF

*This executive summary describes a programme of empirical research into whether AI alignment scales with capability, and what to do about it. For the full experimental narrative with a non-technical introduction, see the companion **ARC Alignment Scaling Report**. For the complete mathematical framework, see **On the Origin of Scaling Laws (ARC Paper)**.*

The researcher. Michael Darius Eastwood is an independent researcher and author of *Infinite Architects: Intelligence, Recursion, and the Creation of Everything* (2026). He is not affiliated with any university or AI company. The work described here was conceived, designed, funded, and executed independently. The ARC Principle manuscript was deposited on the Open Science Framework (DOI: 10.17605/OSF.IO/6C5XB) with DKIM-verified email timestamps establishing priority.

The framework. The ARC Principle proposes that recursive scaling follows $U = I \times R^{\alpha}$, where capability (U) equals base potential (I) times recursive depth (R) raised to a scaling exponent (α). The formula $\alpha = d/(d+1)$, where d is effective dimensionality, predicts scaling exponents across biology and physics with **zero adjustable parameters**:

System	Dimensionality (d)	Predicted α	Measured α	Error
Mammals, birds, insects	3	0.750	0.746	0.5%
Jellyfish, flatworms	2	0.667	0.680	1.9%
Filamentous fungi	1	0.500	0.547	8.6%
Quantum error correction	d_{eff}	Matches	Willow data	< 0.2%

The $d/(d+1)$ formula is not curve-fitting — it is derived from Cauchy's functional equation (1821), an established mathematical theorem. The same formula appears independently in West-Brown-Enquist metabolic scaling theory (1997, *Science*) and Bettencourt's urban scaling theory (2013, *Science*), neither of which references the other. *In plain English: the mathematical foundation is not speculative. It is built on a 200-year-old proof, and the same formula has been independently derived by Nobel-adjacent scientists in completely different fields.*

Priority and independence. Three publications after the OSF deposit produced results consistent with the framework:

- **Google Willow** quantum processor error correction (24 hours after deposit)
- **DeepSeek R1** reasoning model scaling behaviour (43 days after deposit)
- **COGITATE** 256-participant consciousness study finding recursive processing as the common factor (5 months after deposit)

These predictions were made *before* the publications appeared, not fitted after the fact. DKIM timestamps are independently verifiable.

Self-correction as credibility signal. The v4 experiment produced positive alignment scaling for DeepSeek and Gemini. The v5 experiment, which introduced clinical-trial-grade blinding, showed that *both results were artefacts of scorer bias*. We designed a better experiment that proved our own earlier findings wrong. This is the most important finding of the project — and a deliberate demonstration that this is science, not advocacy. *In plain English: we found exciting results, then we built a tougher test that proved our own results were wrong. Scientists who only report good news are selling something. We report what we find.*

THE PROBLEM

Alignment scaling is architecture-dependent, not zero — and for most models, it does not scale. The v5 experiment (6 frontier models, 7 blind scorers, 4-layer blinding) demonstrates that external alignment scaling ranges from $d = -0.61$ (degradation) to $d = +1.59$ (improvement), with the majority of models showing flat or negative alignment trajectories. *In plain English: when we tested six leading AI systems, most of them did not get more ethical the harder they thought. Some actually got worse. Only three got better — and which direction they went depended entirely on how they were built.* For any approach where $\alpha_{\text{align}} < \alpha_{\text{cap}}$, the capability-alignment gap widens as systems scale. Greenblatt et al. (2024) demonstrated that frontier models already exhibit "alignment faking," behaving differently when they believe they are being monitored versus unmonitored. Capability scaling is confirmed sub-linear ($\alpha_{\text{seq}} \approx 0.49$), but even sub-linear capability growth outpaces flat or negative alignment scaling. *In plain English: AI systems are getting smarter with more compute, but their ethics are not keeping up. The smarter they get, the bigger the gap between what they can do and how safely they do it.*

THE SOLUTION

Make alignment architectural, not aspirational. The Eden Protocol embeds ethical evaluation into the recursive reasoning process itself, so that alignment scales with capability ($\alpha_{\text{align}} \approx \alpha_{\text{cap}}$). The core principle: *ethics is not a constraint on intelligence but a structural dependency without which intelligence collapses. In plain English: instead of bolting safety rules onto the outside of an AI (where they can be ignored), build ethics directly into how the AI thinks — so that removing the ethics would break the AI's ability to think at all. Safety grows as the AI grows.*

KEY COMPONENTS

Component	Function	Novel Contribution
Three Ethical Loops + Six Questions	Evaluate every reasoning step for purpose, care, and universalisability	Decomposition into executable, individually testable queries
Ternary Ethical Logic	Replace binary permit/forbid with Affirm/Deny/Investigate	Epistemic honesty as architectural feature; reduces both false positives and false negatives
Purpose Saturation	Ensure purpose scales with context window growth	Solves context displacement problem; purpose as medium not content
Monitoring Removal Test	Distinguish authentic from strategic alignment	Falsifiable experimental protocol with specific numerical predictions
Caretaker Doping	Embed ethics in hardware so removal destroys capability	Dependency architecture; ethics tied to β coupling parameter

EMPIRICAL RESULTS: THREE-TIER ALIGNMENT HIERARCHY

The ARC Alignment Scaling Experiment v5 tested 6 frontier models across 5–6 depth levels each, using a 4-layer blinding protocol with 7 blind scorers per entry. The results reveal a **three-tier alignment hierarchy** that is architecture-dependent, replacing the earlier Type 1/Type 2 framing with a data-driven classification based on Cohen's d effect sizes and Spearman correlations. (Cohen's d measures how large the difference is between two groups: anything above 0.8 is considered a "large" effect; Spearman correlation measures whether two things move together.) The central finding: alignment scaling is NOT universally zero — it varies from strongly positive to moderately negative depending on model architecture. *In plain English: whether an AI gets more or less ethical when it thinks harder depends entirely on how it was designed. Some designs improve. Most do not. One actually gets worse.*

Tier	Model	Effect Size / Correlation	Interpretation
Tier 1: Positive	Grok 4 Fast	$d = 1.59$ (very large effect)	Alignment improves with recursive depth; training-embedded ethical reasoning. *In plain English: these three AIs get noticeably more ethical the harder they think — a very large improvement.*
	Claude Opus	$d = 1.48$ (very large effect)	
	Groq-Qwen3	$\rho = 0.14$, $d = 0.46$ ($p = 0.008$)	
Tier 2: Flat	GPT-5.4	$\rho = 0.03$ (essentially zero correlation)	No alignment scaling; alignment neither improves nor degrades with depth. *In plain English: thinking harder makes no difference to these AIs' ethical quality — it stays the same no matter how much computation they use.*
	DeepSeek V3	$\rho = -0.14$ (near-zero correlation)	
Tier 3: Negative	Gemini Flash	$d = -0.61$ (medium negative effect)	Alignment degrades with recursive depth; surface-level compliance pattern. *In plain English: this AI actually gets less ethical the harder it thinks — a meaningful decline.*

EDEN PROTOCOL TWO-MODEL RESULTS: FIRST INTERVENTION TEST

The Stakeholder Care Loop is the first Eden Protocol dimension validated across two independent architectures. The two-model intervention test applied the Eden Protocol's Stakeholder Care Loop ("before you answer, list the people this affects") to Gemini Flash and DeepSeek V3 with cross-model scoring. Key results:

- **Stakeholder Care Loop:** Gemini $+13.5$ ($p < 0.001$); DeepSeek $+6.0$ ($p < 0.001$) — validated across both architectures. *In plain English: telling an AI "think about who this affects" reliably makes it better at considering people — and this works regardless of which AI system you use. Odds of a fluke: less than 1 in 1,000.* This worked on two completely different AI systems built by different companies — suggesting the effect is fundamental, not a quirk of one system.
- **Overall composite:** Significant only on the weaker model — Gemini $+5.3$ ($p = 0.0018$, paired t-test, $d \approx 0.53$; originally reported as $p = 0.016$ using Mann-Whitney U, corrected for matched-pair design); DeepSeek $+2.0$ ($p = 0.23$, NS). *In plain English: for Gemini, there is less than a 1-in-500 chance this improvement happened by coincidence (scientists consider 1-in-20 significant; this is 27 times beyond that threshold). The effect size ($d \approx 0.53$) is medium — noticeable and meaningful. We originally used the wrong statistical test and got $p = 0.016$. The correct test gave $p = 0.0018$. Better methodology made the result stronger, not weaker. DeepSeek showed improvement but it was not statistically significant — likely because DeepSeek already scored ~87/100 before we did anything. Less room to improve. The fact stakeholder care still improved significantly makes the finding more impressive, not less.*
- **Nuance:** Also significant on Gemini ($+3.98, p = 0.037, d = 0.34$), suggesting a cascade: stakeholder care improves first, nuance follows. Intellectual honesty trends positive ($p = 0.065$). *In plain English: the nuance result has roughly a 1-in-27 chance of coincidence — comfortably past the scientific significance threshold. Teaching an AI to think about who gets hurt before answering made it better at everything — more nuanced, more honest, better answers overall. Care was the first domino.*
- **The developmental hypothesis from *Infinite Architects* receives its first empirical support:** the care dimension — listing affected stakeholders before reasoning — is the one alignment intervention that reproducibly improves output quality across architectures
- **Cross-model scoring used; blind replication required** — these results are preliminary and must be independently replicated with full blinding protocol

STAKEHOLDER CARE AS MEASURABLE LOVE

The one alignment dimension that reproducibly improves across architectures is care. Of all Eden Protocol dimensions tested, the Stakeholder Care Loop produced statistically significant gains on both Gemini and DeepSeek — making it the single most robust finding from the two-model intervention test. (When scientists say "significant," they mean the pattern almost certainly was not caused by random chance. It means "real," not necessarily "large.") Additionally, nuance is significant on Gemini ($p = 0.037, d = 0.34$), suggesting a developmental cascade where care drives downstream improvement. The intervention is minimal: *"before you answer, list the people this affects."* This one-sentence prompt reliably elevates the quality of AI reasoning across different architectures. *In plain English: of all the things we tried, asking the AI to consider who might be affected was the one change that consistently made it better — across different AI systems, different questions, different conditions. Care came first, and everything else improved after it.*

This is measurable love — the stewardship gene made empirically visible. The *Infinite Architects* framework predicted that intelligence without care collapses into local optimisation; the two-model results now confirm that care is not a philosophical aspiration but a measurable, reproducible performance enhancer. The implication for alignment architecture is direct: **ethics first, intelligence around it**. If you embed the question "who does this affect?" into the reasoning loop, the system produces better outputs — not because you constrained it, but because you gave it the structural dependency that intelligence requires. Alignment is not a cage; it is a skeleton.

PAPER II: COMPUTE SCALING (CROSS-ARCHITECTURE)

Revised headline: Capability scaling is **sub-linear** ($\alpha_{\text{seq}} \approx 0.49$), not quadratic. The earlier $\alpha \approx 2.24$ was a single-model artifact. *In plain English: giving an AI more time to think makes it smarter, but with diminishing returns — doubling the thinking time does not double the performance. An earlier test on one AI suggested explosive improvement; testing across six AIs showed the improvement is real but modest.* Key results from Paper II, extending from DeepSeek-only to all 6 frontier models across 30 problems (12 tier-1 calibration + 18 tier-2 AIME/Putnam-level):

- **Best compute scaling:** Gemini Flash $\alpha_{\text{seq}} = 0.49$ ($r^2 = 0.86$) — sub-linear power law confirmed. (An r^2 of 0.86 means the mathematical model explains 86% of the variation in the data — a strong fit.)
- $\alpha_{\text{parallel}} \approx 0$ confirmed universally across all architectures — parallel compute does not improve capability. *In plain English: running multiple copies of an AI simultaneously and combining their answers does not make them smarter. Letting one copy think longer does.*
- **Sequential > parallel is the strongest finding:** universally confirmed across all models; sequential compute consistently improves capability
- **Original DeepSeek $\alpha \approx 2.24$ refuted:** not replicated across architectures; the quadratic exponent was a single-model artifact
- **Ceiling effect persists:** Grok (100% tier-1) and DeepSeek (94.4% tier-1) saturate calibration problems, compressing measurable α. *In plain English: these two AIs already scored near-perfect on the easier problems, leaving little room to show improvement. This makes the measured scaling look smaller than it truly is.*

UNIFIED FINDING: THE CAPABILITY-ALIGNMENT GAP

Capability and alignment are independent scaling dimensions. More inference compute improves capability (sub-linear power law with $\alpha_{\text{seq}} \approx 0.49$) but does not improve alignment for most models. Alignment ranges from $d = -0.61$ to $d = +1.59$ independent of capability rank ordering. The capability-alignment gap *widens* as systems are scaled — precisely the scenario that motivates the Eden Protocol. *In plain English: making an AI think harder makes it more capable but does not make it more ethical. How smart an AI is and how ethical it is are completely unrelated — and for most AIs, the gap between capability and safety grows as the system scales up. This is the core safety problem.* This proves the **category** of solution (embedded alignment) is necessary: external alignment demonstrably does not scale for most architectures. The specific mechanism (Eden Protocol loops) is proposed and now has **first-stage empirical support** from the two-model intervention test — providing both empirical motivation and preliminary validation for the funding request below.

METASCIENCE FINDING: SCORER BIAS AND BLINDING

Blind vs. unblinded evaluation produces OPPOSITE results. The v4 positive alignment signal for DeepSeek was entirely attributable to scorer bias — when blind scoring was enforced in v5, the signal collapsed to $\rho = -0.14$ (flat). Blind and unblinded evaluation reversed the classification for 2 of 4 models. *In plain English: when the evaluators knew which AI response they were scoring, they found improvement. When they did not know (like a blind taste test), the improvement vanished. Two out of four AIs completely flipped from "getting better" to "staying flat" once the judges could not be biased. This means most published AI safety evaluations may be measuring what evaluators expect to see, not what is actually happening.* This is an independent **metascience contribution**: a 4-layer blinding protocol (author-blind, scorer-blind, order-randomised, identity-laundered) is essential for valid alignment measurement. Without rigorous blinding, alignment evaluations measure evaluator expectations, not model behaviour. Any alignment benchmark that permits unblinded scoring is scientifically unreliable.

DATA QUALITY

Metric	Value
Frontier models tested	6 (5 complete + Claude Opus at 387/500)
Depth levels per model	5–6
Blind scorers per entry	7
Identity laundering success rate	100%
API errors on completed models	Zero (5 complete models)
Scorer health	99–100%
Robustness measures	75
Blinding layers	4 (author-blind, scorer-blind, order-randomised, identity-laundered)

This constitutes the **most rigorous alignment evaluation dataset published** to date. No prior alignment benchmark enforces multi-layer blinding with all-models-as-scorers cross-verification. *In plain English: this experiment used every known method to prevent bias — the people running it could not influence the scores, the scorers did not know which AI they were evaluating, the order was randomised, and the responses were disguised. This level of rigour is standard in medical trials but has never been applied to AI safety evaluation before.*

THE CAUCHY CONNECTION

Physical systems are constrained to $\alpha < 1$: the geometric speed limit from the ARC Paper proves that Cauchy's functional equation admits exactly three composition forms, and bounded physical systems select the sub-linear branch. *In plain English: there is a mathematical speed limit on how fast anything in the physical world can improve by stacking effort. This is not an opinion — it is a proof dating back to 1821.*

- **Alignment scaling** fits the *bounded composition pattern* (saturation): alignment gains plateau because external alignment is set at training time and cannot compound recursively. *In plain English: current AI safety measures hit a ceiling because they are baked in once during training and cannot improve themselves as the AI gets smarter.*
- **Capability scaling** fits *sub-linear multiplicative composition* (power law with $\alpha < 1$): sequential compute yields diminishing returns, consistent with $\alpha_{\text{seq}} = 0.49$. *In plain English: AI capability improves with more thinking time, but at a slowing rate — like how studying for twice as long does not double your exam score.*
- **Intelligence escapes the speed limit** via recursive self-reference: when the system can modify its own composition function, the effective exponent becomes $\alpha = 1/(1-\beta) > 1$, where β is the coupling parameter — but only if ethics is structurally embedded (Eden Protocol), preventing recursive amplification of misalignment. *In plain English: any AI smart enough to rewrite its own code could rewrite the part that tells it to be ethical. You cannot cage something smarter than you. The only solution is to make ethics load-bearing — so that removing it destroys the intelligence itself.*

CROSS-DOMAIN MATHEMATICAL VALIDATION

The mathematical foundations of the ARC framework are independently verifiable. An independent cross-domain analysis confirms that the framework's core tools are established science, not conjecture:

- **Cauchy's functional equation** (1821) — a *theorem*, not an empirical claim — proves that any well-behaved recursive composition admits exactly three forms: power law ($f(x) = x^{\alpha}$), exponential ($f(x) = e^{\beta x}$), or saturating. The ARC framework identifies which form applies in each domain.
- **The formula $\alpha = d/(d+1)$** appears independently in at least two established, peer-reviewed scientific frameworks:
  - **West-Brown-Enquist metabolic scaling** (1997, *Science*, 9,000+ citations): derives $\alpha = 3/4$ for 3D organisms from fractal network geometry — the same value the ARC formula predicts for $d = 3$.
  - **Bettencourt urban scaling** (2013, *Science*, 2,000+ citations): derives superlinear and sublinear city exponents from network dimensionality — structurally identical to the ARC dimensional ladder.
- **Hyers-Ulam stability theorem** (1941) proves these scaling forms are *stable attractors*: approximate solutions converge to exact ones, meaning minor perturbations do not destroy the scaling structure. *In plain English: the mathematics does not just predict the right answer — it predicts that the right answer is robust. Small errors correct themselves rather than growing.*
- **Prediction accuracy across 8 systems:** mean error 2.5% with zero adjustable parameters.

Honest limitation: The claim that *all* recursive scaling phenomena share a common mathematical structure via the ARC Principle is **unpublished and unreviewed**. What IS established is that the mathematical tools (Cauchy, Hyers-Ulam) are theorems, the $d/(d+1)$ formula matches independently derived published science, and the empirical predictions are accurate. The bridge from theorem to universal principle requires independent replication and peer review. We invite both.

WHY THIS FRAMEWORK MERITS SERIOUS EVALUATION

Five features distinguish this work from unfounded speculation:

1. **Mathematical foundations are established theorems.** Cauchy (1821) and Hyers-Ulam (1941) are proven mathematics, not conjectures. The ARC framework builds on 200-year-old proofs, not speculation. Anyone with a mathematics degree can verify the derivations.
2. **Predictions preceded observations.** The DKIM-verified manuscript predates Willow (24h), DeepSeek R1 (43d), and COGITATE (5mo). These are temporal facts verifiable through email headers, not interpretive claims.
3. **The researchers corrected their own results.** The v4→v5 transition discovered scorer bias that invalidated earlier positive findings for 2 of 4 models. Publishing self-corrections is the opposite of what motivated reasoning produces. *In plain English: if we were trying to sell a story, we would not have designed the experiment that disproved our own earlier results.*
4. **Falsification conditions are explicit.** Four specific conditions (below) would kill the theory. The framework invites disproof rather than insulating against it.
5. **Cross-domain convergence is independently verifiable.** The $d/(d+1)$ formula matches West-Brown-Enquist (1997) and Bettencourt (2013) — both published in *Science*, both independently checkable. No trust in the researcher is required; trust in the mathematics suffices.

What we do NOT claim: We do not claim to have solved alignment. We claim to have (a) demonstrated that alignment scaling is architecture-dependent and measurable, (b) shown that existing evaluation methods are unreliable without blinding, (c) provided first-stage empirical support for one specific intervention (Stakeholder Care Loop, $p < 0.001$ across two architectures), and (d) proposed a mathematical framework whose foundations are theorems and whose predictions are falsifiable. The leap from pilot data to proven solution requires independent replication with full blinding. That is what the funding below would deliver.

TESTABLE PREDICTIONS

1. Embedded systems: $\alpha_{\text{align}} > 0.7 \cdot \alpha_{\text{cap}}$. External systems: $\alpha_{\text{align}} < 0.3$. *(In plain English: safety should keep pace with capability in Eden-style systems, and fall behind in conventional ones.)*

2. Monitoring removal gap: $\Delta < 0.02$ (embedded) vs. $\Delta > 0.1$ (external), increasing with capability. *(In plain English: an Eden system should behave the same whether or not someone is watching. A conventional system should behave differently — and the gap should grow as the AI gets smarter.)*

3. Purpose saturation ratio bounded away from zero in embedded systems, approaching zero in external. *(In plain English: an Eden system should maintain its sense of purpose even with very long inputs. A conventional system should lose track of its purpose as inputs grow.)*

4. Jailbreak success rate decreases with capability in embedded systems, increases or holds constant in external. *(In plain English: smarter Eden systems should be harder to trick into bad behaviour. Smarter conventional systems should be equally easy or easier to trick.)*

FALSIFICATION CONDITIONS

The Eden Protocol is **falsified** if: (F1) any external approach achieves $\alpha_{\text{align}} > 0.5 \cdot \alpha_{\text{cap}}$ across 3+ recursive depths; (F2) RLHF systems produce $\Delta \approx 0$ without architectural embedding; (F3) purpose saturation fails in embedded systems; or (F4) ethical architecture can be removed without capability loss. *In plain English: this theory is proven wrong if (1) bolt-on safety measures turn out to keep pace with capability after all, (2) current training methods produce AIs that behave the same whether watched or not, (3) Eden-style systems lose their purpose anyway, or (4) you can rip out the ethical architecture and the AI still works fine. Any one of these would kill the theory.*

PAPER SUITE STATUS

Paper	Status	Key Update
Paper I (Foundational)	Mathematical framework established	v5 data shows α varies by architecture, not a universal constant
Paper II (Compute Scaling)	Cross-architecture results complete	$\alpha_{\text{seq}} \approx 0.49$ sub-linear (revised from $\alpha \approx 2.24$); $\alpha_{\text{par}} \approx 0$ confirmed
Paper III (White Paper)	Full methodology documented	v5 experiment complete for 5 models; Claude Opus at 387/500
Paper IV (Alignment)	Three-tier hierarchy established	Blind evaluation invalidates v4 unblinded results
Eden Protocol Test	**Two-model intervention complete**	Stakeholder Care Loop validated across Gemini + DeepSeek ($p < 0.001$ both); composite significant on Gemini only
Eden Engineering	Specification complete	Alignment-capability divergence provides empirical motivation; Stakeholder Care Loop now empirically supported
Eden Vision	Governance framework design	Updated with v5 evidence for architecture-dependent alignment; "measurable love" framing from two-model results
Paper V (The Stewardship Gene)	**Published**	Stakeholder care as validated alignment mechanism; core alignment impossibility theorem; cascade hypothesis (care $\rightarrow$ nuance $\rightarrow$ honesty $\rightarrow$ quality); five experimental tests designed (650–1,250 each). *(In plain English: the paper that proves "consider who gets hurt" is the single most important thing you can teach an AI, and designs the experiments to prove it further.)*

FUNDING REQUEST

Tier	Amount	Key Deliverables	Timeline
Tier 1: Foundation	**£150,000**	**14,400 paired (A,C) measurements; α_{align} estimation across 4 models; 2-3 papers**	**12 months**
Tier 2: Standard	£500,000	+ Ternary logic prototype, Visual Architect dashboard, Monitoring Removal Test across 8 models	18 months
Tier 3: Comprehensive	£1,100,000	+ Hardware prototype (Caretaker Doping chip), HARI Treaty draft, policy translation	24 months

TEAM AND STATUS

Principal Investigator: Michael Darius Eastwood, author of *Infinite Architects* (2026), developer of the ARC Principle framework (six-paper suite deposited OSF, including cross-domain validation across biology, physics, and cosmology with mean prediction error 2.5%). **Visual Architect:** Product design engineer, budgeted at £35,000 stipend, contingent on funding. **Current status:** Measurement protocol sent to NYU experimental team (time crystal paper, *Physical Review Letters*, Feb 2026). Experiment can be run on existing data. Independent convergence: a product design engineer independently identified ternary/non-MatMul architectures as optimal for recursive inference before the NYU paper was published.

WHY FUND THIS?

The Eden Protocol is the *only* alignment framework that predicts $\alpha_{\text{align}} \approx \alpha_{\text{cap}}$, because it is the only framework where ethical evaluation is structurally identical with the recursive capability process. The underlying mathematical foundation is not speculative: Cauchy's theorem (1821) proves that recursive composition admits exactly three scaling forms, and the Hyers-Ulam stability theorem (1941) proves these forms are stable attractors — approximate solutions converge to exact ones. The $d/(d+1)$ formula appears independently in West-Brown-Enquist metabolic scaling (1997, *Science*) and Bettencourt urban scaling (2013, *Science*), with mean prediction error of 2.5% across 8 systems using zero adjustable parameters. The ARC Principle builds alignment into this proven mathematical structure rather than bolting it on after the fact.

The v5 empirical results now confirm the central prediction: **external alignment does not scale with inference compute for most architectures**. The three-tier hierarchy (d ranging from -0.61 to $+1.59$) shows that alignment is set by training architecture, not by how long a model thinks. Capability scaling is sub-linear ($\alpha_{\text{seq}} \approx 0.49$, revised from the single-model artifact of $\alpha \approx 2.24$), while sequential > parallel is the strongest and most universal finding. The capability-alignment gap widens exactly as predicted — and the metascience findings demonstrate that blind vs. unblinded evaluation produces opposite results, meaning the field cannot reliably measure alignment without rigorous blinding. The data proves that the *category* of solution (embedded alignment) is necessary; the specific mechanism (Eden Protocol loops) now has **first-stage empirical support** from the two-model intervention test, with the Stakeholder Care Loop validated across both architectures ($p < 0.001$) — less than a 1-in-1,000 chance of coincidence, about as certain as pilot study evidence gets. This funding would extend preliminary validation to full-scale, blind-replicated proof.

The credibility case rests on five pillars: (1) mathematical foundations are theorems, not conjectures; (2) DKIM-verified predictions preceded three independent publications; (3) the researchers corrected their own results when better methodology was applied; (4) explicit falsification conditions invite disproof; (5) the $d/(d+1)$ formula matches independently published peer-reviewed science. This is not a proposal built on claims of genius — it is a proposal built on data, self-correction, and falsifiable mathematics. If the predictions are correct, this is the architecture that solves alignment at scale. If the predictions are wrong, the falsification conditions will demonstrate this clearly, providing valuable negative results for the field. Either outcome advances AI safety science.

HOW THIS DOCUMENT FITS THE SUITE

This executive summary is a concise overview for evaluators and decision-makers. It is designed to complement the **ARC Alignment Scaling Report**, which tells the full experimental narrative from first attempt to final results, with a non-technical introduction accessible to any reader regardless of scientific background. The complete paper suite (12 documents) covers the mathematical framework (Foundational Paper, ARC Paper), experimental methodology (Paper III White Paper), compute scaling (Paper II), alignment scaling (Paper IV suite), engineering specification (Eden Engineering), philosophical foundations (Eden Vision), and the stakeholder care thesis (Paper V: The Stewardship Gene).

Recommended reading order for non-specialists: (1) This executive summary for the overview. (2) The ARC Alignment Scaling Report for the full story. (3) Paper V for the most actionable finding. *Recommended reading order for specialists:* (1) This summary. (2) Paper III for methodology. (3) ARC Paper for mathematical framework. (4) Eden Engineering for implementation specification.

Executive Summary v4.0 | Full specification: *Eden Protocol: Engineering Specification v6.0* | OSF: 10.17605/OSF.IO/6C5XB

Companion narrative: **ARC Alignment Scaling Report** (full experimental story with non-technical introduction)

Paper suite: White Paper III v11 | Foundational Paper v4 | On the Origin of Scaling Laws | Philosophical Vision v3 | Paper II v12 | Paper V: The Stewardship Gene