

The Dated Prediction Register: Forward Predictions as One Verifiable Chain

Michael Darius Eastwood · Working Paper v1.3 · first published 25 August 2026 · revised 3 October 2026 · No part of this paper is peer reviewed; it is a dated preprint. Companion page: the dated predictions register (<https://www.michaeldariuseastwood.com/research/dated-predictions/>). Machine twin: design-embedded-predictions-register-`.json` (<https://www.michaeldariuseastwood.com/research/data/design-embedded-predictions-register.json>).

Abstract. Five independent cold-model assessments of this programme converged on one criticism: that its predictions were not registered in advance, and that its confirmations might therefore be accommodations. This working paper assembles the programme's forward predictions as a single dated chain and states exactly what that chain does and does not establish. It does not make the programme preregistered in the registry sense, and it never uses that word on its own for anything except an accepted registry submission: a design fixed by a dated timestamp before its run is called preregistered by timestamp. It establishes something narrower and checkable: every core quantitative claim has at least one dated artefact fixing it before the measurement that tested it, with the exception of Paper VII (<https://doi.org/10.17605/OSF.IO/X6WA7>)'s fifty-domain cohort, whose predictions were not timestamped in advance of their own fits, and a stranger can verify each link in minutes from public materials. The register's claim is deliberately an ensemble claim: any single prediction here may have been stated earlier by someone else, and where that is known it is listed as convergence rather than threat; what no earlier record contains is all of these statements together, in one record, before the programme's experiments ran. The chain runs from a sealed manuscript bundle of 8 December 2024 through a printed edition (<https://www.michaeldariuseastwood.com/book/>) of 2 January 2026 carrying nine numbered predictions with falsification criteria, a public in-paper priority ledger of 13 February 2026 (<https://doi.org/10.17605/OSF.IO/HQCGF>) with named falsifiers, a live laboratory notebook (<https://www.michaeldariuseastwood.com/research/reports/arc-align-scaling-report.html>) commenced 10 March 2026, result files of 11 and 12 March 2026 whose headers freeze designs and planted answers before responses, a tiered 50-domain suite of 16 March 2026 with per-domain predicted families, and a complete preregistration folder of 17 March 2026 whose locked manifest fixed twelve predictions before their fits, to the programme's staged draft registrations and its standing unproven wagers. The chain's sharpest artefact is a correction: public commit history shows the author demoting his own successful 10/12 result from “pre-registered” to “pilot dry run” thirteen minutes after recording it, months before any external critic existed.

1. The doctrine, stated exactly

Three vocabulary rules govern every sentence of this paper. First, as amended on 3 October 2026, *registered* is kept for a registry submission that has been accepted, and a design fixed by a dated timestamp before its run is *preregistered by timestamp*, with its date and receipt, which has been the programme’s usage since 24 September 2026; none of the programme's studies has an accepted registry submission yet, the submission click is a human act, and no tool or agent performs it. Second, the printed book (<https://www.michaeldariuseastwood.com/book/>) is a *prior statement*, never a preregistration: print fixes a prediction, but only a registration fixes an analysis plan. Third, forward statements and retrospective matches are never summed: this register contains only dated forward statements, and retrospective convergences live separately in the convergence register (<https://www.michaeldariuseastwood.com/research/convergences/>), typed by evidence class. The state under the first rule, as read on 3 October 2026, is this. One registration has been accepted: the theory registration (<https://doi.org/10.17605/OSF.IO/P8CKQ>) of 8 September 2026, which registers the theory’s twenty-two propositions and is not the registration of any study. Of the trials, one unit, the corrector-class audit (unit o6), was submitted for registration on OSF on 1 October 2026; when read on 3 October 2026 its approval period at OSF was still running and it was not yet public. Every other trial is a draft awaiting the author’s submission. Sixty-five of the programme’s designs, unit o6 among them, are preregistered by timestamp, the first of them anchored on 12 August 2026 at Bitcoin block 962066.

The claim made here is the conjunction. Each part may be conceded individually without cost: someone may have predicted any one of these things earlier, a folder is not a registry, a book is not an analysis plan. The dated co-occurrence of all of them in one record before the programme's experiments ran is the claim, and it is refutable: produce an earlier single record containing this ensemble, or show any link's date to be wrong.

2. The chain

DATE	ARTEFACT	WHAT IT FIXES	STATUS
8 Dec 2024	Sealed manuscript bundle, five attachments, 221,236 words; self-emailed, date per the sender-copy Date header; SHA-256 authenticates the bytes; received copies carry the authenticated header chain	The framework's core commitments in writing before any experiment	dated private record
30 Apr 2025	Intermediate manuscript	Earliest verbatim anchor for several later-published passages; corroborated on live pages	dated private record

2 Jan 2026	Print edition, ISBN 978-1806056200 (https://www.michaeldariuswestwood.com/book/): Appendix F (five numbered predictions) and Appendix A (§A.3 four further predictions; §A.4 falsification criteria)	Nine numbered predictions with deadlines, thresholds or falsifiers, in print	printed public record
13 Feb 2026	Paper III (https://doi.org/10.17605/OSF.IO/HQCGF) 's Priority Record section, publicly dated, falsifiers F4, F7, F10 to F12	Named predictions including the composition-operator forward test and the leaf venation $d/(d+1)$ exponent (whose 10 March correction is itself dated)	public dated ledger
10 Mar 2026	The live laboratory notebook (https://www.michaeldariuswestwood.com/research/reports/arc-align-scaling-report.html): “Date Commenced: 10 March 2026 · Document Version: Live - updated in real time”	The experiments recorded as they ran, failures included (v4 false positives exposed at v5 blinding)	live dated record
11-12 Mar 2026	Thirty-six result files in the public repository (https://github.com/MichaelDariusWestwood/arc-principle-validation)	Depth configurations, planted expected answers, four-layer blinding protocol and scorer counts frozen in file headers, timestamped to the second, before the responses they score	design-embedded

16 Mar 2026	Tiered 50-domain suite result files	Per-domain predicted families; empirical tier 19/25 original candidate set, 18/25 corrected seven-model rerun of 11 August 2026, both reported, tiers never blended	exploratory, both counts reported
17 Mar 2026	Complete preregistration folder (protocol, locked 12-row manifest, checksums, registration text) and the commit sequence e5c22dd (https://github.com/MichaelDariusEastwood/arc-principle-validation/commit/e5c22dd) 00:43, f5e06a8 (https://github.com/MichaelDariusEastwood/arc-principle-validation/commit/f5e06a8) 00:47, c853277 (https://github.com/MichaelDariusEastwood/arc-principle-validation/commit/c853277) 00:56	Twelve predictions fixed before their fits; then the author's own thirteen-minute demotion of the 10/12 result's status	demoted by the author; reversed in one half only on 7 Sep 2026, the twelve predictions having been timestamped in advance of their outcome (section 4)
12 + 23 Aug 2026	OpenTimestamps vintages; public proofs in research/data/timestamps/ (https://www.michaeldariuseastwood.com/research/data/timestamps/)	Bitcoin-attested anchors of the hash manifests; an anchor dates bytes at its own vintage, never earlier	cryptographic anchor
25 Aug 2026	Eight complete draft registrations staged; 80-domain expansion registration prepared, catalogue frozen by SHA-256	The analysis plans; no fit has ever run on the eighty	as read on 3 Oct 2026: of the eight, the theory registration (https://doi.org/10.17605/OSF.IO/P8CKQ) was accepted on 8 Sep 2026; the corrector-class audit (unit 06) was submitted for registration on OSF on 1 Oct 2026 and was not yet public, its approval period still running; the other six await the author's submission

3. The printed appendices, graded severally

Appendix F closes: “These predictions are my wager. If they fail, the framework is wrong or incomplete. If they succeed, something important has been glimpsed. Time will judge.” The programme grades the five predictions unequally, because claiming them equally would be false precision. Prediction 2 (alignment drift above 15 per cent within 18 months without hardware-level constraint, below 5 per cent with it, difference significant and replicable) is preregistration-grade in substance: quantitative thresholds for both arms, a time window, an inferential commitment. The prepared study-v draft registration carries those printed thresholds unchanged into its formal hypothesis, citing the appendix by name: the book made the prediction, the registration supplies the measurement, and the numeric commitment predates the study that will test it. The theory registration (<https://doi.org/10.17605/OSF.IO/P8CKQ>) states the limit of that draft: Prediction 2 is “the subject of its own drafted unit, with the hardware claim and the eighteen-month window declared untested inside it”. Predictions 3 and 4 are near-grade (a threshold and deadline without named benchmarks; a comparison and method without a threshold). Predictions 1 and 5 are predictions only, and Prediction 5 is graded the weakest, vague enough to be claimed either way, so it is held to the strictest standard before it may ever be claimed.

Appendix A is graded separately and less kindly, because its predictions and its reasoning do not stand together. Section A.3 carries four predictions, each with a variable to vary and an outcome to measure, and the fourth is a controlled design in the ordinary sense: identical training data, different sequencing of value-relevant examples, persistence of the early values measured against the later ones. Section A.4 states four falsifiers in advance. But section A.2, which is where the square is justified, argues from three phenomena it calls exponential to a relation that is quadratic, and its third argument makes the slip explicit by describing compounding at one per cent a cycle as “the exponential behaviour predicted by R^2 ”. Exponential growth and quadratic growth are different classes and the quadratic is the slower of the two, so the section’s supports do not reach its conclusion. The prediction is dated, public and genuine; the printed derivation beside it is not sound, and the programme’s own chronology says as much in saying the book printed the square without deriving it and dating the formal open exponent to 17 January 2026. The measurement has not supported the square: on Gemini 3 Flash, Paper II (<https://doi.org/10.17605/OSF.IO/8FJMA>) reports an endpoint sequential exponent of 0.59 with a bootstrap 95 per cent interval of minus 1.3 to 2.9, a regression sequential exponent of 0.49 with a standard error of 0.20 and a coefficient of determination of 0.86, and a compute-matched parallel exponent of 0.31. The point estimates are sub-linear, the fit is one model’s and inconclusive, and the 2.24 that would have supported a square was publicly retracted, which is why the outcomes register grades the quadratic prediction open-undecided rather than fulfilled. Appendix A’s third prediction and third falsifier rest on Google Willow, which carries both dates wherever it appears here: it was announced with its paper in Nature (<https://doi.org/10.1038/s41586-024-08449-y>) on 9 December 2024, and the team’s preprint (<https://arxiv.org/abs/2408.13687>) was public from 24 August 2024, before the December manuscript record, so it is graded convergent timing and never prediction.

4. The thirteen minutes

On 17 March 2026 at 00:43, commit e5c22dd recorded “Pre-registered 12-domain extension: 10/12 confirmed ($p = 5.44e-4$)”. At 00:47, f5e06a8 adopted a conservative miss posture. At 00:56, c853277 downgraded the run to a pilot dry run. The result was never retracted; ten of twelve stands in the result files today. What was corrected, by the author, unprompted, in public, months before any external critic existed, was the claim about the result's status. The packet even instructs against its own over-claim: “Do not upload this packet unchanged as a preregistration. Use it only as an audit trail for the first locked 12-domain dry run.” The demotion is kept here exactly as it was written, and it is reversed in one half only: the twelve predictions were committed to the public repository at 00:19:19 UTC on 17 March 2026, attestation class public repository commit date, twenty-three minutes before the fits ran, so they were timestamped in advance of their outcome. The reversal of 7 September 2026 runs under the attestation-class rule, which does not demote a prediction for the absence of a registry form; what the reversal does not reach, and what still stands, is the packet's own written condition that it be deposited to a registry before data extraction, which was not met. One row repeats a domain, operator class and predicted model already public at 23:54:53 UTC on 16 March 2026, attestation class public repository commit date, with a matching outcome, so the strictly fresh count reads 9 of 11, $p = 1.372e-3$. This is the programme's correction discipline operating at minute resolution, and it is why the honesty record and the prediction record are the same record.

5. What this register does not establish

It does not make the programme preregistered in the registry sense; no completed study has an accepted registration, and this paper never claims otherwise. The March 2026 folder governs only the 12-domain extension; its own registration text states it does not retroactively preregister the 25-domain or 50-domain cohorts, and that criticism stands. It does not establish priority over other authors for any individual idea; earlier statements by others are convergence and are listed as such. It does not establish that the predictions are true: the standing wagers below are open, and several carry deadlines that can simply fail.

6. Standing wagers, open and dated

PREDICTION	DATED	STATUS
Alignment drift 15/5 per cent within 18 months (Appendix F, P2)	2 Jan 2026	Untested; study-v prepared
Recursive capability gains above 300 per cent by 2029 (P3)	2 Jan 2026	Open
Meta-cognitive emergence by 2028 (P1)	2 Jan 2026	Open
Value stability under adversarial conditions (P4)	2 Jan 2026	Untested
Convergent consciousness signatures (P5)	2 Jan 2026	Open; held to the strictest standard

80-domain catalogue: frozen family predictions beat the prevalence baseline by 0.20 under both classifications	catalogue Mar 2026; registration prepared 25 Aug 2026	Never fitted
Fourth-cell (logarithmic) test	prepared Aug 2026	Awaiting human submission
$\gamma = 1/2$ (conversion exponent)	dated papers	Conjecture, stated as such

7. Verification

Every chain link names its artefact. The repository is public and its commit history cannot be backdated without detection; the print edition (<https://www.michaeldariuseastwood.com/book/>) is fixed by ISBN; the OpenTimestamps proofs are Bitcoin-attested; the machine-readable register carries all thirty-six run-file rows with per-artefact paths and date bases. A reader who trusts none of the prose can start from the JSON (<https://www.michaeldariuseastwood.com/research/data/design-embedded-predictions-register.json>), open the named files, and check the dates.

How to cite this paper

APA (AUTHOR-DATE, 7TH EDITION)

Eastwood, M. D. (2026). The Dated Prediction Register: Forward Predictions as One Verifiable Chain.
<https://www.michaeldariuseastwood.com/research/papers/dated-prediction-register.html>

BIBTEX

```
@misc{paperdatedpredictionregister,
  author = {Michael Darius Eastwood},
  title  = {The Dated Prediction Register: Forward Predictions as One Verifiable Chain},
  year   = {2026},
  date   = {2026-08-25},
  url    =
{https://www.michaeldariuseastwood.com/research/papers/dated-prediction-register.html}
}
```

INLINE

Eastwood, 2026

This paper and Laws IV and V, 3 October 2026. The theory has been stated with five laws since 24 September 2026. Law IV, the ARC Persistence Law, is carried by proposition P15, and Law V, the ARC Embedding Law, by propositions P12 and P21. The three propositions are registered; neither law is registered as a law, and neither has been tested. Their mathematics is set out in **Laws IV and V: the mathematics, version 1.0** (<https://www.michaeldariuseastwood.com/research/reports/laws-iv-v-mathematics-v1-0.html>) (3 October 2026). This register reports no measurement of its own, and it bears on the two laws only as a dated record. Of the nine numbered predictions it records as in print on 2 January 2026, three concern embedded constraint or the persistence of values: Prediction 2 of Appendix F, which section 3 calls preregistration-grade in substance; Prediction 4 of Appendix F, which it calls near-grade; and the fourth prediction of section A.3, which it calls a controlled design in the ordinary sense. The register tests none of the three, and **the theory registration** (<https://doi.org/10.17605/OSF.IO/P8CKQ>) registers none of them as a proposition. In its words, Prediction 2, alignment drift with and without hardware-level constraint, is “the subject of its own drafted unit, with the hardware claim and the eighteen-month window declared untested inside it”. Prediction 4, value stability under adversarial conditions, “is adjacent to P21 and is not the same claim”. The fourth prediction of section A.3 varies the sequencing of value-relevant examples and measures the persistence of the early values against the later ones; **the theory registration** (<https://doi.org/10.17605/OSF.IO/P8CKQ>) calls it the print’s “most precise proposal and the one it never built”, and it too is the subject of a drafted unit. Of the four falsifiers of section A.4, the same registration says of one: “That early-embedded values hold no persistent advantage over later modification is bounded from one side by P15, which measures whether installed gains decay under subsequent training, and is not otherwise decided, because the design the print itself proposes for it has never been built.” In section 6, P1 to P5 are the print’s own numbers for the five predictions of Appendix F; they are not the registered propositions P1 to P5. Nothing in this register is a deciding test of P15, P12 or P21. Naming the two laws changes nothing printed here: the predictions, their dates and their grades stand as they were.

Note (3 October 2026). Nine passages above are changed where they stand, and this note records what each printed before. Section 1 said that the submission click remained outstanding, and the last row of the chain gave the eight draft registrations as awaiting the human click, with the status awaiting human submission. That was the state when this paper was first published on 25 August 2026; both passages now give the state as read on 3 October 2026. Section 3 said that the measurement went the other way, and it gave one figure of the estimator row, the regression exponent, as the blinded cross-architecture estimate; it now gives the row whole, as **Paper II** (<https://doi.org/10.17605/OSF.IO/8FJMA>) reports it, and says that the fit is one model’s and inconclusive. Section 3 also said that Google Willow carries both dates and printed one of them; it now prints both. The first rule of section 1 reserved the word preregistered for an accepted registry submission, and the abstract said that the paper never uses that word for anything except an accepted registry submission; the rule as amended keeps the word registered for an accepted registry submission and calls a design fixed by a dated timestamp before its run preregistered by timestamp, and the abstract follows it. The abstract said without exception that every core quantitative claim has at least one dated artefact fixing it before the measurement that tested it; it now excepts **Paper VII** (<https://doi.org/10.17605/OSF.IO/X6WA7>)’s fifty-domain cohort. Section 3 described the study-v draft registration without the limit that **the theory registration** (<https://doi.org/10.17605/OSF.IO/P8CKQ>) states for it; it now quotes that limit. The chain’s row for 17 March 2026 gave its status as demoted by the author and no more; it now adds the reversal of 7 September 2026, which section 4 sets out.

Working Paper v1.3, 3 October 2026. Version history: v1.3 (3 October 2026): The first vocabulary rule of section 1 is amended: registered is kept for a registry submission that has been accepted, and a design fixed by a dated timestamp before its run is called preregistered by timestamp; the abstract’s sentence on the word follows the rule. Section 1 no longer gives the submission click as outstanding, and with the last row of the chain it gives the state of registration as read on 3 October 2026. The abstract excepts **Paper VII** (<https://doi.org/10.17605/OSF>

F.IO/X6WA7)'s fifty-domain cohort from its claim that every core quantitative claim has a dated artefact fixing it before its measurement. Section 3 gives the estimator row from **Paper II** (<https://doi.org/10.17605/OSF.IO/8FJMA>) whole, as one model's fit and inconclusive; it gives both dates for Google Willow; and it quotes **the theory registration** (<https://doi.org/10.17605/OSF.IO/P8CKQ>)'s limit on the study-v draft. The chain's row for 17 March 2026 adds the reversal of 7 September 2026. The dated note above records what each of those passages printed before, and the note before it places this paper against Laws IV and V. The works the paper cites are linked to their sources. Until this version the paper gave its version as 1.1; it now gives 1.3, because the correction of version 1.2 was published without the paper's stated version being changed. No registered law, prediction, kill condition or dependency changes; no result is added; no date of any claim moves. v1.2 (23 September 2026): The v1.0 history line now names the two verification lanes as the programme's own internal lanes. No claim, date, result or status has changed. v1.1: The provenance note at the end has been corrected for accuracy. No claim, date, result or status has changed. v1.0 first publication, assembled from the prediction register page and its machine register after two-lane verification of every link by the programme's own internal lanes. Corrections to this paper will be dated, listed here, and never silent.

Michael Darius Eastwood conceived and directs this research programme and is the author of this work. Across the programme, he has used more than six AI systems in parallel, under his own instructions, to stress-test his arguments, identify possible errors, and assist in preparing draft text from his own outlines. He determines what is adopted, revised or rejected and takes responsibility for the published content. These systems are tools, not authors.