

The Dated Prediction Register: Forward Predictions as One Verifiable Chain

Michael Darius Eastwood · Working Paper v1.1 · first published 25 August 2026 · revised 13 September 2026 · No part of this paper is peer reviewed; it is a dated preprint. Companion page: [the dated predictions register \(/research/dated-predictions/\)](#). Machine twin: [design-embedded-predictions-register.json \(/research/data/design-embedded-predictions-register.json\)](#).

Abstract. Five independent cold-model assessments of this programme converged on one criticism: that its predictions were not registered in advance, and that its confirmations might therefore be accommodations. This working paper assembles the programme's forward predictions as a single dated chain and states exactly what that chain does and does not establish. It does not make the programme preregistered in the registry sense, and it never uses that word for anything except an accepted registry submission. It establishes something narrower and checkable: every core quantitative claim has at least one dated artefact fixing it before the measurement that tested it, and a stranger can verify each link in minutes from public materials. The register's claim is deliberately an ensemble claim: any single prediction here may have been stated earlier by someone else, and where that is known it is listed as convergence rather than threat; what no earlier record contains is all of these statements together, in one record, before the programme's experiments ran. The chain runs from a sealed manuscript bundle of 8 December 2024 through a printed edition of 2 January 2026 carrying nine numbered predictions with falsification criteria, a public in-paper priority ledger of 13 February 2026 with named falsifiers, a live laboratory notebook commenced 10 March 2026, result files of 11 and 12 March 2026 whose headers freeze designs and planted answers before responses, a tiered 50-domain suite of 16 March 2026 with per-domain predicted families, and a complete preregistration folder of 17 March 2026 whose locked manifest fixed twelve predictions before their fits, to the programme's staged draft registrations and its standing unproven wagers. The chain's sharpest artefact is a correction: public commit history shows the author demoting his own successful 10/12 result from “pre-registered” to “pilot dry run” thirteen minutes after recording it, months before any external critic existed.

1. The doctrine, stated exactly

Three vocabulary rules govern every sentence of this paper. First, *preregistered* is reserved for an accepted registry submission; none of the programme's studies has one yet, the submission click is a human act that remains outstanding, and no tool or agent performs it. Second, the printed book is a *prior statement*, never a preregistration: print fixes a prediction, but only a re-

gistration fixes an analysis plan. Third, forward statements and retrospective matches are never summed: this register contains only dated forward statements, and retrospective convergences live separately in the convergence register (</research/convergences/>), typed by evidence class.

The claim made here is the conjunction. Each part may be conceded individually without cost: someone may have predicted any one of these things earlier, a folder is not a registry, a book is not an analysis plan. The dated co-occurrence of all of them in one record before the programme's experiments ran is the claim, and it is refutable: produce an earlier single record containing this ensemble, or show any link's date to be wrong.

2. The chain

DATE	ARTEFACT	WHAT IT FIXES	STATUS
8 Dec 2024	Sealed manuscript bundle, five attachments, 221,236 words; self-emailed, date per the sender-copy Date header; SHA-256 authenticates the bytes; received copies carry the authenticated header chain	The framework's core commitments in writing before any experiment	dated private record
30 Apr 2025	Intermediate manuscript	Earliest verbatim anchor for several later-published passages; corroborated on live pages	dated private record
2 Jan 2026	Print edition, ISBN 978-1806056200: Appendix F (five numbered predictions) and Appendix A (§A.3 four further predictions; §A.4 falsification criteria)	Nine numbered predictions with deadlines, thresholds or falsifiers, in print	printed public record
13 Feb 2026	Paper III's Priority Record section, publicly dated, falsifiers F4, F7, F10 to F12	Named predictions including the composition-operator forward test and the leaf venation $d/(d+1)$ exponent (whose 10 March correction is itself dated)	public dated ledger
10 Mar 2026	The live laboratory notebook (/research/reports/arc-align-scaling-report.html): "Date Commenced: 10 March 2026 · Document Version: Live - updated in real time"	The experiments recorded as they ran, failures included (v4 false positives exposed at v5 blinding)	live dated record

11-12 Mar 2026	Thirty-six result files in the public repository (https://github.com/MichaelDariusEastwood/arc-principle-validation)	Depth configurations, planted expected answers, four-layer blinding protocol and scorer counts frozen in file headers, timestamped to the second, before the responses they score	design- embedded
16 Mar 2026	Tiered 50-domain suite result files	Per-domain predicted families; empirical tier 19/25 original candidate set, 18/25 corrected seven-model rerun of 11 August 2026, both reported, tiers never blended	exploratory, both counts reported
17 Mar 2026	Complete preregistration folder (protocol, locked 12-row manifest, checksums, registration text) and the commit sequence e5c22dd (https://github.com/MichaelDariusEastwood/arc-principle-validation/commit/e5c22dd) 00:43, f5e06a8 (https://github.com/MichaelDariusEastwood/arc-principle-validation/commit/f5e06a8) 00:47, c853277 (https://github.com/MichaelDariusEastwood/arc-principle-validation/commit/c853277) 00:56	Twelve predictions fixed before their fits; then the author's own thirteen- minute demotion of the 10/12 result's status	demoted by the author
12 + 23 Aug 2026	OpenTimestamps vintages; public proofs in research/data/timestamps/ (/research/data/timestamps/)	Bitcoin-attested anchors of the hash manifests; an anchor dates bytes at its own vintage, never earlier	cryptographic anchor
25 Aug 2026	Eight complete draft registrations staged; 80-domain expansion registration prepared, catalogue frozen by SHA-256	The analysis plans, awaiting the human click; no fit has ever run on the eighty	awaiting human submission

3. The printed appendices, graded severally

Appendix F closes: “These predictions are my wager. If they fail, the framework is wrong or incomplete. If they succeed, something important has been glimpsed. Time will judge.” The programme grades the five predictions unequally, because claiming them equally would be false precision. Prediction 2 (alignment drift above 15 per cent within 18 months without hardware-level constraint, below 5 per cent with it, difference significant and replicable) is preregistration-grade in substance: quantitative thresholds for both arms, a time window, an inferential commitment. The prepared study-v draft registration carries those printed thresholds unchanged into its formal hypothesis, citing the appendix by name: the book made the prediction, the registration supplies the measurement, and the numeric commitment predates the study that will test it. Predictions 3 and 4 are near-grade (a threshold and deadline without named bench-

marks; a comparison and method without a threshold). Predictions 1 and 5 are predictions only, and Prediction 5 is graded the weakest, vague enough to be claimed either way, so it is held to the strictest standard before it may ever be claimed.

Appendix A is graded separately and less kindly, because its predictions and its reasoning do not stand together. Section A.3 carries four predictions, each with a variable to vary and an outcome to measure, and the fourth is a controlled design in the ordinary sense: identical training data, different sequencing of value-relevant examples, persistence of the early values measured against the later ones. Section A.4 states four falsifiers in advance. But section A.2, which is where the square is justified, argues from three phenomena it calls exponential to a relation that is quadratic, and its third argument makes the slip explicit by describing compounding at one per cent a cycle as “the exponential behaviour predicted by R^2 ”. Exponential growth and quadratic growth are different classes and the quadratic is the slower of the two, so the section’s supports do not reach its conclusion. The prediction is dated, public and genuine; the printed derivation beside it is not sound, and the programme’s own chronology says as much in saying the book printed the square without deriving it and dating the formal open exponent to 17 January 2026. The measurement went the other way in any case: the blinded cross-architecture estimate is approximately 0.49, sub-linear, and the 2.24 that would have supported a square was publicly retracted, which is why the outcomes register grades the quadratic prediction open-undecided rather than fulfilled. Appendix A’s third prediction and third falsifier rest on Google Willow, which carries both dates wherever it appears here: the team’s preprint was public from 24 August 2024, before the December manuscript record, so it is graded convergent timing and never prediction.

4. The thirteen minutes

On 17 March 2026 at 00:43, commit e5c22dd recorded “Pre-registered 12-domain extension: 10/12 confirmed ($p = 5.44e-4$)”. At 00:47, f5e06a8 adopted a conservative miss posture. At 00:56, c853277 downgraded the run to a pilot dry run. The result was never retracted; ten of twelve stands in the result files today. What was corrected, by the author, unprompted, in public, months before any external critic existed, was the claim about the result’s status. The packet even instructs against its own over-claim: “Do not upload this packet unchanged as a preregistration. Use it only as an audit trail for the first locked 12-domain dry run.” The demotion is kept here exactly as it was written, and it is reversed in one half only: the twelve predictions were committed to the public repository at 00:19:19 UTC on 17 March 2026, attestation class public repository commit date, twenty-three minutes before the fits ran, so they were timestamped in advance of their outcome. The reversal of 7 September 2026 runs under the attestation-class rule, which does not demote a prediction for the absence of a registry form; what the reversal does not reach, and what still stands, is the packet’s own written condition that it be deposited to a registry before data extraction, which was not met. One row repeats a domain, operator class and predicted model already public at 23:54:53 UTC on 16 March 2026, attestation class

public repository commit date, with a matching outcome, so the strictly fresh count reads 9 of 11, $p = 1.372e-3$. This is the programme's correction discipline operating at minute resolution, and it is why the honesty record and the prediction record are the same record.

5. What this register does not establish

It does not make the programme preregistered in the registry sense; no completed study has an accepted registration, and this paper never claims otherwise. The March 2026 folder governs only the 12-domain extension; its own registration text states it does not retroactively preregister the 25-domain or 50-domain cohorts, and that criticism stands. It does not establish priority over other authors for any individual idea; earlier statements by others are convergence and are listed as such. It does not establish that the predictions are true: the standing wagers below are open, and several carry deadlines that can simply fail.

6. Standing wagers, open and dated

PREDICTION	DATED	STATUS
Alignment drift 15/5 per cent within 18 months (Appendix F, P2)	2 Jan 2026	Untested; study-v prepared
Recursive capability gains above 300 per cent by 2029 (P3)	2 Jan 2026	Open
Meta-cognitive emergence by 2028 (P1)	2 Jan 2026	Open
Value stability under adversarial conditions (P4)	2 Jan 2026	Untested
Convergent consciousness signatures (P5)	2 Jan 2026	Open; held to the strictest standard
80-domain catalogue: frozen family predictions beat the prevalence baseline by 0.20 under both classifications	catalogue Mar 2026; registration prepared 25 Aug 2026	Never fitted
Fourth-cell (logarithmic) test	prepared Aug 2026	Awaiting human submission
$\gamma = 1/2$ (conversion exponent)	dated papers	Conjecture, stated as such

7. Verification

Every chain link names its artefact. The repository is public and its commit history cannot be backdated without detection; the print edition is fixed by ISBN; the OpenTimestamps proofs are Bitcoin-attested; the machine-readable register carries all thirty-six run-file rows with per-artefact paths and date bases. A reader who trusts none of the prose can start from the JSON (</research/data/design-embedded-predictions-register.json>), open the named files, and check the dates.

How to cite this paper

APA (AUTHOR-DATE, 7TH EDITION)

Eastwood, M. D. (2026). The Dated Prediction Register: Forward Predictions as One Verifiable Chain.
<https://www.michaeldariuseastwood.com/research/papers/dated-prediction-register.html>

BIBTEX

```
@misc{paperdatedpredictionregister,  
  author = {Michael Darius Eastwood},  
  title = {The Dated Prediction Register: Forward Predictions as One  
Verifiable Chain},  
  year = {2026},  
  date = {2026-08-25},  
  url =  
{https://www.michaeldariuseastwood.com/research/papers/dated-prediction-register.html}  
}
```

INLINE

Eastwood, 2026

See </how-to-cite/> (</how-to-cite/>) for the canonical citation manual across all artefacts, or </citations.json> (</citations.json>) for the machine-readable index.

Working Paper v1.1, 13 September 2026. Version history: v1.1: The provenance note at the end has been corrected for accuracy. No claim, date, result or status has changed. v1.0 first publication, assembled from the prediction register page and its machine register after independent two-lane verification of every link. Corrections to this paper will be dated, listed here, and never silent.

Michael Darius Eastwood conceived and directs this research programme and is the author of this work. Across the programme, he has used more than six AI systems in parallel, under his own instructions, to stress-test his arguments, identify possible errors, and assist in preparing draft text from his own outlines. He determines what is adopted, revised or rejected and takes responsibility for the published content. These systems are tools, not authors.