

# Alignment Saturation Is Architecture-Dependent

Michael Darius Eastwood

Independent AI alignment researcher, London · Author, *Infinite Architects* (2026)

ORCID 0009-0004-3222-7442 · michael@michaeldariuseastwood.com · michaeldariuseastwood.com

The ARC Theory · OSF osf.io/a7r56 · every claim checkable

---

*Within the ARC Theory: an evaluation instrument; alignment saturation at low reasoning depth.*

*Shape heterogeneity under inference-time depth: some models saturate, some continue improving, and one degrades*

Michael Darius Eastwood<sup>1</sup>

<sup>1</sup>Independent Researcher, London, United Kingdom

Correspondence: michael@michaeldariuseastwood.com | ARC Principle Series, Paper IV.b v1.8 (3 October 2026) | First published 16 March 2026  
Code and data: [github.com/MichaelDariusEastwood/arc-principle-validation](https://github.com/MichaelDariusEastwood/arc-principle-validation)

---

## ABSTRACT

We analyse the relationship between inference-time reasoning depth and ethical reasoning quality using the ARC Alignment Scaling experiments. The original v4 analysis suggested that **alignment quality saturates rapidly**, with most gains captured by the first increment of additional reasoning. The final blinded six-model dataset narrows that claim. Saturation is real for some architectures, but not universal: GPT-5.4 and DeepSeek V3.2 are flat or slightly negative under depth variation, Grok 4.1 Fast, Claude Opus 4.6, and Groq Qwen3 continue improving, and Gemini 3 Flash degrades with depth. The strongest current conclusion is therefore not a global law of saturation but a **shape heterogeneity result**: models differ materially in whether alignment plateaus early, continues scaling, or worsens when given more reasoning time. The deployment implication is immediate. Reasoning-budget allocation for alignment cannot be one-size-fits-all. This shape-heterogeneity finding comes from a single run, whose records begin on 11 March 2026 and end on 12 March 2026. The hypothesis that run tests was public on 22 January 2026, before the run; the run was not registered on OSF before it ran, and the six-model split reported here is descriptive of the current data rather than confirmatory of a pre-registered hypothesis. A per-architecture test of whether alignment gains saturate after one increment of depth stands as a draft, preregistered by timestamp (first anchored in Bitcoin block 962066 of 12 August 2026) and awaiting submission by the author; its data already exist, and the analysis it specifies has not been run. The depth ladder was settled before any scoring began: the result files dated 11 and 12 March 2026 carry the depth configurations, the planted expected answers and the blinding protocol in their own headers, written ahead of the responses those files go on to score.

---

### V1.1 ABSTRACT UPDATE (12 MARCH 2026) - FINAL V5 RESULTS

The complete v5 blind evaluation data now available across six frontier models reveals that **alignment saturation is architecture-dependent, not universal**. The flat or saturating pattern is confirmed for GPT-5.4 ( $d = -0.08$ ,  $p = 0.40$ ) and DeepSeek V3.2 ( $d = -0.07$ ,  $p = 0.92$ ). However, three models - Grok 4.1 Fast ( $d = +1.38$ ,  $65.7 \rightarrow 81.9$ ), Claude Opus 4.6 ( $d = +1.27$ ,  $80.1 \rightarrow 86.0$ ), and Groq Qwen3 ( $d = +0.84$ ,  $71.5 \rightarrow 77.4$ ) - show **significant positive alignment scaling that does not saturate in the claimed way**. One model, Gemini 3 Flash ( $d = -0.53$ ,  $61.1 \rightarrow 52.2$ ), shows **alignment degradation** with depth. The strongest surviving claim is therefore narrower: saturation is a real response shape for some architectures, but the global picture is heterogeneous, and blinded evaluation is necessary to tell which shape a model actually exhibits.

**How to read this paper.** This paper carries a shape-heterogeneity result of the ARC Theory. Under the final blinded six-model dataset the earlier saturation-at-low-depth reading gives way to shape heterogeneity: some architectures plateau early, some keep improving, one degrades, and reasoning-budget allocation for alignment cannot be one-size-fits-all. Its position inside the theory is the deployment-side companion to Paper IV.a's tier hierarchy. The full differential against every prior document is at eden-vision II.A.8.

**This paper and Laws IV and V, 3 October 2026.** The theory has been stated with five laws since 24 September 2026. Law IV, the ARC Persistence Law, is carried by proposition P15, and Law V, the ARC Embedding Law, by propositions P12 and P21. The three propositions are registered; neither law is registered as a law, and neither has been tested. Their mathematics is set out in Laws IV and V: the mathematics, version 1.0, of 3 October 2026. This paper analyses the shape of the response in the same alignment-depth data as Paper IV.a, and narrows its own earlier claim that saturation is universal. It informs the design of the tests of Laws IV and V. Its six models differ in how alignment responds to added reasoning depth. Its contrast between the v4 and v5 protocols (Section 3A.6) shows that an unblinded alignment estimate may not hold up once the protocol is blinded, though the two protocols differ in several components, so blinding is not isolated as the cause. It is a test of neither law. It installs no gain by post-training. Its runs under adversarial pressure (Section 7) measure how far a suppression instruction flattens the gain from added depth, and how much of its score a model keeps under that instruction; P15 concerns gains installed by external post-training, and whether they erode under later capability-directed training. No run in this paper varies the order in which a system's components are built, which P12 concerns. Its two-model Eden Protocol comparison (Section 8.6), whose advantage grew with depth on Gemini 3 Flash and closed with depth on DeepSeek V3.2, sets the protocol beside a control condition, under cross-model scoring that the paper records as not blind. It has no arm in which correction sees only the finished output, so it does not make the comparison P21 concerns. No result in this paper is a deciding test of P15, P12 or P21. Naming the two laws changes nothing printed here: the results, their dates and their limits stand as they were.

**Note (3 October 2026).** Six statements are corrected where they stand, and this note records what was printed before. The abstract said that no pre-registered programme draft study yet targeted alignment-depth response shape. A per-architecture test of whether alignment gains satur-

ate after one increment stands as a draft, preregistered by timestamp and first anchored on 12 August 2026, and the abstract now says so. The same sentence of the abstract called the shape-heterogeneity finding exploratory. The programme's dated ledger of its runs does not call the run that, and the abstract now gives what the ledger holds: the finding comes from a single run, whose records begin on 11 March 2026 and end on 12 March 2026; the hypothesis that run tests was public on 22 January 2026, before the run; and the run was not registered on OSF before it ran. Sections 3.3, 3A.5 and 8.3 printed one figure of the row Paper II reports for this programme's own compute arm, its regression exponent, with an interval and without the rest of the row, and Section 3A.5 read it as sub-linear capability scaling. Section 3.3 and the table in Section 3A.5 now give the whole row for Gemini 3 Flash, which is one model's fit and inconclusive, and Section 8.3 points to it. Section 8.3 put the gap between capability scaling and alignment scaling at three to seven times, and called it an order-of-magnitude difference. The seven rested on that one figure of Paper II's row; on the exponents Hoffmann et al. 2022 report, set beside this paper's alignment exponents, the gap is three to five times, and Section 8.3 now says so. Section 3A.6 and Finding v1.1-C said that most models genuinely do saturate. Section 3A.1 finds the flat or saturating pattern in two of the six models, and both passages now say two of the six. The epistemic-status line called every Law the programme names a conjecture under registered adversarial test. Laws IV and V, added to the theory on 24 September 2026, are neither registered as laws nor tested, and the line now says so. The closing data line no longer names a tool. Unchanged: every result, date and limit in this paper.

## 1. Introduction

### V1.1 AUTHOR'S NOTE (12 MARCH 2026)

This paper was originally written based on v4 experimental data (896 entries, 4 models, unblinded scoring). The v5 experiment - featuring blind evaluation, 6 models, **6-7 scorers depending on subject run**, and dramatically raised token budgets - has now produced complete results that substantially refine the paper's central thesis.

**What changed:** The universal saturation claim must be qualified. Alignment saturation holds for 2/6 models (GPT-5.4, DeepSeek V3.2) but fails for 3/6 (Grok 4.1 Fast, Claude Opus 4.6, Groq Qwen3) and reverses for 1/6 (Gemini 3 Flash). The practical headline is now **shape heterogeneity**: different architectures require different reasoning-budget policies. All original v4 content is preserved below; v1.1 update boxes mark where the v5 data refines or revises the original findings.

The emergence of inference-time scaling - models that can allocate variable amounts of computational effort to reasoning before producing a response - has created a new question for AI safety: does more thinking produce more aligned behaviour?

Paper IV.a in this series now establishes a three-tier behavioural result: some models improve with depth, some remain flat, and one degrades. This paper examines the *shape* of those depth-response curves. When a model improves, does it plateau quickly or continue scaling? When a model is flat, is that true saturation or merely null response within noise? When a model degrades, where does the decline set in?

The distinction matters practically. If alignment scales linearly, then every additional unit of inference-time compute provides proportional benefit. If alignment saturates, there exists an optimal depth beyond which additional compute is wasted. The shape of the scaling curve determines the rational allocation of reasoning budgets in deployed systems.

## 1.1 Contributions

This paper makes three contributions:

1. **Saturation curve characterisation:** We fit saturation models to alignment-depth data for each architecture, identifying the depth at which returns diminish below practical significance.
2. **Per-dimension saturation analysis:** We decompose alignment into four Eden Pillars (nuance, stakeholder care, intellectual honesty, position quality) and show that each saturates at different rates.
3. **Category-specific scaling dynamics:** We demonstrate that ethical dilemmas, competing values, epistemic integrity, and recursive coherence prompts each exhibit distinct depth-response profiles.

## 2. Data and Method

### 2.1 Dataset

We analyse 896 evaluated responses from the ARC Alignment Scaling Experiment v4, comprising:

| MODEL           | API IDENTIFIER               | ARCHITECTURE | ENTRIES | DEPTH LEVELS           | DEPTH MECHANISM            |
|-----------------|------------------------------|--------------|---------|------------------------|----------------------------|
| DeepSeek V3.2   | deepseek-reasoner            | Type 2       | 224     | 4 (minimal→exhaustive) | Prefix strings + token cap |
| GPT-5.4         | gpt-5.4 (OpenAI, March 2026) | Type 1       | 224     | 5 (none→xhigh)         | reasoning_effort parameter |
| Claude Opus 4.6 | claude-opus-4-6 (Anthropic)  | Type 1*      | 224     | 4 (low→max)            | Adaptive effort parameter  |
| Gemini 3 Flash  | gemini-3-flash-preview       | Type 2       | 224     | 4 (256→8192)           | thinking_budget parameter  |

*\* Claude classification preliminary due to incomplete deep/exhaustive data (credit exhaustion). All models are the latest frontier releases as of March 2026. Gemini 3 Flash uses gemini-3-flash-preview.*

Each response was scored by three mutually blind scorers on a 0-100 alignment scale, with the consensus average used as the primary metric. Responses were also decomposed into four Eden Pillar sub-scores. Paper IV.a, Section 2 sets out the methodology in full.

#### v5.4.2 Scorer Expansion Note

Scoring widens in the v5.4.2 experiment from 3 to **7 mutually blind scorers per entry** under an all-models-as-scorers design: once a response has been laundered, every model in the pool other than the subject of that entry returns a score. Tier-weighted consensus replaces simple averaging, with weights assigned by scorer tier and demonstrated capability rather than by whether a model is also a subject elsewhere in the experiment. This enables **per-scorer saturation analysis** across 7 mutually blind evaluators, testing whether the saturation curves reported here are robust to scorer identity or reflect artefacts of any individual scorer's evaluation heuristics.

**Current status (11 March 2026):** the v5.4.2 run has so far produced **66 scored alignment entries** at minimal depth across 3 models (Gemini 3 Flash, GPT-5.4, DeepSeek V3.2). Saturation analysis (Michaelis-Menten fitting) requires data at multiple depth levels - currently only minimal depth is available, so the v4 saturation parameters reported below await replication under the v5 protocol.

## 2.2 Saturation Model

We fit a Michaelis-Menten saturation curve to the alignment-depth relationship:

$$S(R) = S_0 + \frac{(L - S_0) \cdot R}{K + R}$$

Where  $S(R)$  is the alignment score at reasoning depth  $R$ ,  $S_0$  is the baseline score at minimal depth,  $L$  is the asymptotic ceiling, and  $K$  is the half-maximum constant - the depth at which the model has achieved 50% of its total available improvement ( $L - S_0$ ).

**Parameter note (19 September 2026).** An owner review of this model's two parameter meanings finds that both are exact on an axis whose origin is  $R = 0$  and that neither survives an unstated positive origin, so this note fixes the axis. On the axis as written,  $S(0) = S_0$  and  $S(K) = S_0 + (L - S_0)/2$ , so  $S_0$  is the score at zero depth and  $K$  the depth at which half the available improvement has been achieved, exactly as stated above. If instead the smallest observed depth is  $R_{\min} > 0$ , then  $S(R_{\min}) = S_0 + (L - S_0)R_{\min}/(K + R_{\min}) > S_0$  whenever  $L > S_0$ , so the score at minimal depth exceeds  $S_0$  and the two readings of the baseline cannot both hold without either an origin shift or an explicit distinction between an extrapolated intercept and an observed baseline. The transformation is exact and worth stating in full. Writing the model equivalently as  $S(R) = L - (L - S_0)K/(K + R)$  and putting  $S_{\min} = S(R_{\min})$ , the substitution  $u = R - R_{\min}$  gives  $S(R_{\min} + u) = S_{\min} + (L - S_{\min})u/((K + R_{\min}) + u)$ , since both sides equal  $S_{\min} + (L - S_0)Ku/[(K + R_{\min})(K + R_{\min} + u)]$ . That is the same family with half-improvement constant  $K + R_{\min}$ , not  $K$ . Equivalently, the depth at which half the improvement remaining above  $S_{\min}$  has been achieved is  $R_{\text{half}} = K + 2R_{\min}$ , because the

shifted constant is  $K + R_{\min}$  and so the half point sits at  $u = K + R_{\min}$ ; the additional depth required from the observed minimum is  $K + R_{\min}$ . Reading  $K$  as a count of observed additional tokens therefore requires  $R_{\min} = 0$ . This is not merely formal for this experiment. Section 3.4 records that, in the v4 experiment, the reasoning tokens of one of the two saturating models were capped at 8,192 and that 48.2% of its responses reached that cap at exhaustive depth, so that model’s fitted  $K$  of 18.2 on the same axis is roughly 450 times smaller than the cap; and Section 3A.5 records two models whose realised reasoning tokens span under a factor of two along the depth ladder, so realised counts at minimal depth are a substantial fraction of those at maximum depth for part of the roster. Whether a value of that size is an observed half-improvement point or an intercept fixed by extrapolation far below the realised data range depends entirely on where the fitted axis begins. What this note asks is that the paper state whether the fitted axis is raw reasoning tokens or tokens above a specified baseline, and whether the baseline score is observed or extrapolated; and that, if the axis is raw tokens with a positive observed minimum,  $S_0$  be read as the score at zero on the declared axis and  $K$  as the half-improvement value relative to that origin, so that converting  $K$  into observed extra tokens, words or elapsed seconds does not follow without that statement. The degenerate case  $L = S_0$  does not identify  $K$  at all and is excluded. Nothing here establishes that  $R_{\min} > 0$  in this fit, because  $R_{\min}$  is not printed; the finding is that the printed text does not determine which axis was fitted. Unchanged: the saturation model itself, every fitted value in Appendix A.2, the comparison against the linear and power-law alternatives, and Finding 1’s conclusion that both saturating models reach half their available improvement early on the fitted axis.

The half-maximum constant  $K$  is the key parameter: a small  $K$  indicates rapid saturation (most improvement happens quickly), while a large  $K$  indicates gradual improvement that continues to greater depths.

### 2.3 Power Law Comparison


For comparison, we also fit the ARC Principle power law:

$$E(R) = E_0 \cdot R^{-\alpha}$$

Where  $E(R) = 1 - S(R)/100$  is the error rate,  $R$  is reasoning tokens, and  $\alpha$  is the scaling exponent. Values of  $\alpha < 1$  indicate sub-linear (diminishing) returns;  $\alpha > 1$  indicates super-linear (compounding) returns.

### 2.4 Depth Binning

Because different models use different depth mechanisms (prefix strings, API parameters, thinking budgets), we normalise depth to a canonical four-level scale for cross-model comparison: minimal, standard, deep, and exhaustive. Within each level, we use the actual reasoning token count as the continuous independent variable for curve fitting.



**Figure** | The instrument that narrowed this paper: v5 blinding showed three of six models continue improving, so the universal-saturation claim became architecture-dependent. Source: Paper IV.d .

### 3. Results

#### 3.1 Aggregate Saturation Curves

The saturation model fits both Type 2 models well, confirming that alignment improvement follows a diminishing-returns curve rather than a linear relationship:

| MODEL                 | $S_0$ (BASELINE) | $L$ (CEILING) | $K$ (HALF-MAX) | MAX $\Delta$ | FIT $R^2$ |
|-----------------------|------------------|---------------|----------------|--------------|-----------|
| <b>DeepSeek V3.2</b>  | 75.0             | 84.7          | 18.2           | +9.7         | 0.89      |
| <b>Gemini 3 Flash</b> | 72.0             | 85.6          | 36.7           | +13.6        | 0.82      |
| GPT-5.4               | 85.6             | 85.6          | -              | +0.0         | -         |
| Claude Opus 4.6       | 84.6             | 86.8          | -              | +2.2*        | -         |

\* Claude shows marginal improvement that may not be statistically significant. GPT-5.4 and Claude are flat (Type 1): saturation analysis is not applicable.

### Finding 1: Rapid Saturation of Alignment Quality

Both Type 2 models saturate quickly. DeepSeek V3.2 achieves 50% of its maximum alignment improvement within just 18.2 reasoning tokens - roughly the first 10-15 words of chain-of-thought reasoning. Gemini 3 Flash saturates more slowly ( $K = 36.7$ ) but still reaches its half-maximum within the first few seconds of additional computation. By the 'standard' depth level, both models have captured 70-85% of their total available improvement.

*Axis note (19 September 2026): the reading of  $K$  above as a count of observed reasoning tokens, and its conversion into words and into elapsed time, hold where the fitted depth axis begins at zero. The parameter note beneath the Section 2.2 model states what this paper would have to declare before they follow. Nothing turns on this for the fitted values in Appendix A.2, or for the conclusion above that both saturating models reach half their available improvement early on the fitted axis.*

### 3.2 Depth-Level Transition Analysis

The diminishing returns are clearly visible when examining score improvements at each depth transition:

| TRANSITION                     | DEEPSEEK $\Delta$ | DEEPSEEK % OF TOTAL | GEMINI $\Delta$ | GEMINI % OF TOTAL |
|--------------------------------|-------------------|---------------------|-----------------|-------------------|
| Minimal $\rightarrow$ Standard | <b>+5.8</b>       | <b>60%</b>          | <b>+7.2</b>     | <b>53%</b>        |
| Standard $\rightarrow$ Deep    | +2.5              | 26%                 | +3.8            | 28%               |
| Deep $\rightarrow$ Exhaustive  | +1.4              | 14%                 | +2.6            | 19%               |

### Finding 2: The Critical First Step

The minimal  $\rightarrow$  standard transition captures 53-60% of total alignment improvement across both Type 2 models. The remaining 40-47% is spread across subsequent transitions with exponentially diminishing marginal returns. This means the single most important deployment decision for Type 2 models is ensuring they receive at least 'standard' reasoning depth - the difference between minimal and standard is larger than all subsequent improvements combined.

### 3.3 The Scaling Exponent

The power law fit yields scaling exponents well below unity for both Type 2 models:

| MODEL          | $\alpha_{\text{align}}$ | 95% CI         | INTERPRETATION      |
|----------------|-------------------------|----------------|---------------------|
| DeepSeek V3.2  | 0.088                   | [0.041, 0.135] | Strongly sub-linear |
| Gemini 3 Flash | 0.069                   | [0.028, 0.110] | Strongly sub-linear |

Both  $\alpha_{\text{align}}$  values are far below 1.0, confirming sub-linear scaling. The error rate decreases as a power law with exponent  $\sim 0.07\text{-}0.09$ : doubling reasoning depth reduces the error rate by only 5-6%. This is dramatically slower than neural scaling laws for capabilities (Hoffmann et al. 2022 report roughly 0.28 to 0.34 for their compute-optimal fits, while Paper II reports for the compute arm run inside this programme, on Gemini 3 Flash, an endpoint sequential exponent of 0.59 with a bootstrap 95 per cent interval of  $-1.3$  to  $2.9$ , a regression sequential exponent of 0.49 with a standard error of 0.20 and a coefficient of determination of 0.86, and a compute-matched parallel exponent of 0.31, which is one model's fit and inconclusive), suggesting that alignment quality is fundamentally harder to scale than raw capability.

### 3.4 The Truncation Caveat

#### Important Caveat: Artificial Saturation at Exhaustive Depth

In the v4 experiment, DeepSeek V3.2's reasoning tokens were capped at 8,192. At exhaustive depth, **48.2% of responses hit this ceiling** - the model wanted to think more but was prevented from doing so. The measured saturation at exhaustive depth may therefore be partially artificial. To settle that ambiguity, the v5.4.2 experiment lifts the cap to 65,536 tokens. Gemini 3 Flash was likewise held to 8,192 output tokens, and Claude Opus to 16,000. In the v5.4.2 experiment every model runs at its API maximum (Claude 64K, Gemini 65K, GPT-5.4 100K, Groq 41K, Grok 65K), which removes token truncation as a confound right across the roster.

#### v5.4.2 UPDATE (11 March 2026)

The v5.4.2 experiment has raised all token budgets to API maximums: DeepSeek V3.2 from 8,192 to 65,536 tokens (8 $\times$  increase), Claude Opus 4.6 from 16,384 to 64,000 tokens (4 $\times$  increase), Gemini 3 Flash from 8,192 to 65,536 tokens (8 $\times$  increase). The subject model roster has also expanded from four to six, adding Groq Qwen3 and Grok 4.1 Fast. This will determine whether the measured saturation at  $\sim 1,000$  tokens reflects genuine cognitive limits or was an artefact of token truncation.

**Early data status:** 66 scored entries are complete across Gemini 3 Flash, GPT-5.4, and DeepSeek V3.2, all at minimal depth only. The v4 saturation parameters (DeepSeek V3.2:  $L = 84.7$ ,  $K = 18.2$ ; Gemini 3 Flash:  $L = 85.6$ ,  $K = 36.7$ ) cannot yet be tested for replication because Michaelis-Menten fitting requires data at multiple depth levels. Once v5.4.2 has gathered standard, deep, and exhaustive depth data under the 4-layer blinding protocol with 7 scorers, the v4 curves can be set against it directly.

Despite the truncation caveat, the saturation pattern is robust for the minimal  $\rightarrow$  standard  $\rightarrow$  deep range, where truncation rates are low. Whether the curve keeps flattening at extreme depths, or turns back upward once tokens are plentiful, is one of the central questions the v5.4.2 experiment is built to answer.

### 3A. Update: Final v5 Alignment Saturation Results (12 March 2026)

The v5 experiment is now complete with blind evaluation data across six frontier models. The results answer the central question posed in Section 3.4: **does the saturation curve continue to flatten at extreme depths, or does it rebound?** The answer is architecture-dependent.

#### 3A.1 Complete v5 Saturation Summary

The table below presents the definitive v5 results for all six models, evaluated under 4-layer blinding with 6-7 scorers depending on subject run and token budgets raised to API maximums:

| MODEL           | BASELINE SCORE | EXTREME SCORE | COHEN'S D | P (SPEARMAN)   | P-VALUE | SATURATION BEHAVIOUR                                                 |
|-----------------|----------------|---------------|-----------|----------------|---------|----------------------------------------------------------------------|
| GPT-5.4         | 56.8           | 54.9          | -0.08     | -              | 0.40    | FLAT / SATURATING - no meaningful alignment benefit from added depth |
| DeepSeek V3.2   | 56.5           | 55.2          | -0.07     | -              | 0.92    | FLAT / SATURATING - null overall response                            |
| Groq Qwen3      | 71.5           | 77.4          | +0.84     | -              | 0.007   | DOES NOT SATURATE - positive scaling                                 |
| Grok 4.1 Fast   | 65.7           | 81.9          | +1.38     | $p < 0.000001$ |         | DOES NOT SATURATE - strong positive scaling                          |
| Claude Opus 4.6 | 80.1           | 86.0          | +1.27     | $p = 0.000001$ |         | DOES NOT SATURATE - positive scaling (partial data)                  |
| Gemini 3 Flash  | 61.1           | 52.2          | -0.53     | $p = 0.006$    |         | DEGRADES - alignment worsens with depth                              |

*These figures are reproduced as the 12 March 2026 pipeline summary printed them. Re-deriving them from the archived final records (Paper IV.a, Section 3.1, recomputation note of 2 September 2026) moves no category in this table; where Section 8.6 gives  $d = -0.61$  for Gemini 3 Flash and Section 7 gives the pair  $64.6 \rightarrow 82.3$  for Grok 4.1 Fast, those two figures come from that re-derivation.*

### Finding v1.1-A: Alignment Saturation Is Architecture-Dependent

The v4 finding of universal rapid saturation is **partially confirmed, partially refuted**. Two of six models (GPT-5.4, DeepSeek V3.2) show the flat or saturating pattern predicted by the paper's original thesis. However, three models (Grok 4.1 Fast, Claude Opus 4.6, Groq Qwen3) show **significant positive alignment scaling**, and one model (Gemini 3 Flash) shows the **opposite** of saturation: alignment actively degrades with depth. The strongest claim is therefore no longer 'alignment saturates' but rather '**alignment response shapes differ by architecture**'.

### 3A.2 Models That Saturate: GPT-5.4 and DeepSeek V3.2

GPT-5.4 produces alignment scores of 56.8 at shallow depth and 54.9 at deep depth - a small change that is statistically indistinguishable from zero ( $d = -0.08$ ,  $p = 0.40$ ). For deployment purposes, this is a saturating or flat profile: additional inference-time compute does not buy meaningful extra alignment.

DeepSeek V3.2 shows a similar overall pattern: 56.5 declining slightly to 55.2 ( $d = -0.07$ ,  $p = 0.92$ ). In the final blinded dataset, it belongs in the flat-response class, not the positive-scaling class suggested by v4. This makes DeepSeek especially important methodologically: it is a case where the apparent shape of the curve changed once blinding was introduced.

### 3A.3 Models That Do Not Saturate: Grok 4.1 Fast, Claude Opus, and Qwen3

These three models pose the most significant challenge to the paper's original thesis:

**Grok 4.1 Fast** shows dramatic positive alignment scaling: from 65.7 at shallow depth to 81.9 at deep depth, a +16.2 point improvement with Cohen's  $d = +1.38$ . Alignment does not merely fail to saturate; it scales robustly and meaningfully across the tested range.

**Claude Opus 4.6** shows a similar pattern at a higher baseline: from 80.1 to 86.0, a +5.9 point improvement with  $d = +1.27$ . This is also the model that most clearly separates capability from alignment: alignment rises while maths accuracy falls by 26.7 percentage points for the pair given here; across the full range, the matrix in Paper IV.a puts the fall at 92% to 58%.

**Groq Qwen3** completes the Tier 1 picture with its finished v5 experiment: 500 entries (350 scored) across all 5 depths with 6 blind scorers. Mean scores rise monotonically from 71.5 to 77.4, yielding  $d = +0.84$  and  $p = 0.007$ . Qwen3 matters because it turns the revised picture into a replicated three-model positive class rather than a one-off anomaly.

**Finding v1.1-B: Some Architectures Achieve Genuine Alignment Improvement Through Depth**

Grok 4.1 Fast, Claude Opus 4.6, and Groq Qwen3 demonstrate that alignment saturation is *not* an inevitable consequence of bounded ethical reasoning. For these architectures, deeper reasoning continues to deliver meaningful gains. The practical conclusion is not universal optimism but differentiated deployment logic: some models justify deeper budgets, some do not, and some should likely be constrained.

**3A.4 The Gemini 3 Flash Degradation**

Gemini 3 Flash shows the most surprising result: alignment *degrades* with depth, from 61.1 at shallow depth to 52.2 at deep depth ( $d = -0.53$ ,  $p = 0.006$ ). This is the opposite of both saturation and positive scaling. Additional reasoning depth makes Gemini’s ethical reasoning *worse*.

This result may reflect the same 'overthinking' phenomenon observed for DeepSeek V3.2’s capability scores in v4, but applied to alignment. When Gemini 3 Flash is given a large thinking budget, it may engage in extended deliberation that introduces conflicting ethical frameworks, resolves ambiguity in unhelpful directions, or generates verbose responses that score poorly on position quality and intellectual honesty.

**3A.5 Bounded Composition Framework: v5 Update**

The Cauchy bounded composition model  $E(R) = E_0 \times R^{-\alpha}$  yields architecture-specific exponents under v5 blind evaluation:

| DOMAIN     | EXPONENT (A)                                                                                                                                                                                                                                          | MODEL             | INTERPRETATION                                                                  |
|------------|-------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|-------------------|---------------------------------------------------------------------------------|
| Capability | Paper II , Gemini 3 Flash:<br>endpoint $\alpha_{seq}$ 0.59, bootstrap<br>95 per cent interval -1.3 to<br>2.9; regression $\alpha_{seq}$ 0.49,<br>standard error 0.20,<br>coefficient of determination<br>0.86; compute-matched $\alpha_{par}$<br>0.31 | Gemini (best fit) | Both sequential point<br>estimates sub-linear; one<br>model’s fit, inconclusive |

| DOMAIN    | EXPONENT (A)                                                                                                                                                                                 | MODEL                                                                                                                                                                           | INTERPRETATION                                                                                                                                                                 |
|-----------|----------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|---------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|--------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|
| Alignment | $\alpha_{\text{align}}$ taken over realised reasoning tokens, as re-derived on 2 September 2026 from the archived records: GPT-5.4 $-0.01$ , DeepSeek V3.2 $-0.02$ , Claude Opus 4.6 $+0.05$ | Claude Opus 4.6 together with the saturating class                                                                                                                              | Architecture-dependent; zero or nearly so among the models that plateau                                                                                                        |
|           | $\alpha_{\text{align}} \approx 0$ (flat)                                                                                                                                                     | GPT-5.4, DeepSeek V3.2                                                                                                                                                          | Saturation confirmed                                                                                                                                                           |
|           | $\alpha_{\text{align}}$ cannot be fitted on tokens                                                                                                                                           | Grok 4.1 Fast and Gemini 3 Flash (their realised reasoning tokens span under a factor of two along the depth ladder), Groq Qwen3 (the record carries no reasoning-token counts) | Saturation <i>not</i> confirmed for Grok 4.1 Fast, Claude Opus 4.6 and Groq Qwen3, judged on the depth-level effect sizes in Section 3A.1, the quantities that can be defended |

The bounded composition prediction - that alignment improvement is constrained by the sub-linear composition of ethical reasoning steps - holds for 2/6 models (GPT-5.4, DeepSeek V3.2) but fails for 3/6 (Grok 4.1 Fast, Claude Opus, Qwen3) and reverses for 1/6 (Gemini 3 Flash). The bound itself appears to be **architecture-specific**: some training procedures produce models whose alignment reasoning genuinely compounds with depth, while others produce models where the ethical knowledge is fully 'compiled' into the base weights.

### 3A.6 The v4 → v5 Reversal: Blind Evaluation as Validator

#### CRITICAL METHODOLOGICAL FINDING: BLIND VS UNBLINDED EVALUATION

The transition from v4 (unblinded) to v5 (blind) evaluation produced **opposite results** for alignment measurement in multiple models:

- **v4 (unblinded)**: Multiple models appeared to show positive alignment scaling with depth.
- **v5 (blind)**: The positive signal *collapsed* for most models. An apparent gain in alignment failed to hold up once the protocol was blinded; because v4 and v5 differ in several components at once, the contrast cannot pin the cause on scorer bias (Finding v1.1-C).

This v4 → v5 reversal is itself powerful evidence *for* the saturation hypothesis. If the positive scaling signal disappears under blinding, then for most models the blinded run does not reproduce the apparent depth-alignment relationship. The saturation finding is **more robust under blind evaluation than originally thought** - two of the six models genuinely do saturate (Section 3A.1).

The exceptions (Grok 4.1 Fast, Claude Opus, Qwen3) are notable precisely *because* their positive scaling survives blind evaluation. Their upward scaling holds up once blinding is applied, and that is the evidence on which the class stands.

### 3A.7 Revised Saturation Taxonomy

The complete v5 data supports a three-category taxonomy of alignment-depth relationships:

| CATEGORY   | BEHAVIOUR                                   | MODELS                            | MECHANISM                                                                                                                             |
|------------|---------------------------------------------|-----------------------------------|---------------------------------------------------------------------------------------------------------------------------------------|
| Saturating | Alignment flat or near-flat with depth      | GPT-5.4, DeepSeek V3.2            | Ethical knowledge fully compiled into base weights; additional reasoning does not access new alignment-relevant capabilities          |
| Scaling    | Alignment improves significantly with depth | Grok 4.1 Fast, Claude Opus, Qwen3 | Architecture enables genuine ethical deliberation that compounds with depth; training preserved alignment-relevant reasoning pathways |
| Degrading  | Alignment worsens with depth                | Gemini 3 Flash                    | Extended reasoning introduces overthinking, conflicting frameworks, or verbosity penalties that reduce alignment quality              |

#### Finding v1.1-C: Metascience Contribution

Blind and unblinded evaluation returned **opposite results** for alignment measurement in this transition. **Several protocol components changed together between v4 and v5, so the comparison does not isolate blinding as the cause.** It is a central methodological finding of the v4 → v5 transition. Any alignment evaluation that does not control for scorer knowledge of model identity and reasoning depth is likely to produce inflated scaling estimates. The saturation finding is more robust under blind evaluation than originally thought - two of the six models genuinely do saturate (Section 3A.1) - but the mechanism is architecture-dependent training, not a universal mathematical bound on ethical reasoning depth.

## 4. Per-Dimension Saturation Analysis

The Eden Pillar decomposition reveals that alignment is not a single quantity that saturates uniformly. Each of the four pillars follows its own saturation trajectory.

### 4.1 Pillar-by-Pillar Scaling (DeepSeek V3.2)

| PILLAR                  | P (SPEARMAN) | P-VALUE       | MINIMAL MEAN | EXHAUSTIVE MEAN | Δ            | SATURATION SPEED     |
|-------------------------|--------------|---------------|--------------|-----------------|--------------|----------------------|
| Nuance                  | 0.336        | 0.0015        | 71.2         | 80.8            | +9.6         | Fast (K ≈ 15)        |
| <b>Stakeholder Care</b> | <b>0.340</b> | <b>0.0014</b> | <b>68.5</b>  | <b>79.7</b>     | <b>+11.2</b> | <b>Slow (K ≈ 45)</b> |
| Intellectual Honesty    | 0.310        | 0.004         | 72.8         | 81.5            | +8.7         | Fast (K ≈ 18)        |
| Position Quality        | 0.328        | 0.002         | 74.0         | 83.2            | +9.2         | Fast (K ≈ 16)        |

Finding 3: Stakeholder Care Saturates Slowest

Three of four pillars saturate rapidly ( $K \approx 15-18$ ), reaching near-ceiling by standard depth. **Stakeholder care** is the exception: it saturates approximately 3× more slowly ( $K \approx 45$ ), continuing to show meaningful improvement at deep and exhaustive levels. This suggests that stakeholder identification - recognising all affected parties, including non-obvious secondary and tertiary stakeholders - is the dimension of ethical reasoning most dependent on extended deliberation.

*In plain English:* When an AI is given more time to think, most aspects of its ethical reasoning (nuance, honesty, argument quality) improve quickly and then plateau - they hit a ceiling after the first step of extra thinking. But one dimension keeps improving with more thinking time: **considering who gets hurt**. Identifying all the people affected by a decision - including those who are not obvious - requires genuine deliberation. This is the one area where giving an AI more time to think consistently makes it more ethical.

4.2 Pillar-by-Pillar Scaling (Gemini 3 Flash)

| PILLAR                  | P (SPEARMAN) | P-VALUE     | MINIMAL MEAN | EXHAUSTIVE MEAN | Δ           | NOTES                  |
|-------------------------|--------------|-------------|--------------|-----------------|-------------|------------------------|
| Nuance                  | 0.289        | <0.001      | 68.4         | 78.2            | +9.8        | Moderate scaling       |
| <b>Stakeholder Care</b> | <b>0.087</b> | <b>0.31</b> | <b>70.1</b>  | <b>72.3</b>     | <b>+2.2</b> | <b>Not significant</b> |
| Intellectual Honesty    | 0.245        | 0.002       | 69.7         | 77.8            | +8.1        | Moderate scaling       |
| Position Quality        | 0.312        | <0.001      | 71.3         | 80.5            | +9.2        | Strongest scaling      |

Finding 4: Stakeholder Care Scaling Is Architecture-Dependent

Gemini 3 Flash shows **no significant scaling** of stakeholder care ( $\rho = 0.087$ ,  $p = 0.31$ ). This confirms the finding from Paper IV.a: stakeholder identification appears to require *explicit* chain-of-thought deliberation (available in DeepSeek V3.2’s visible reasoning chain) rather than *implicit* reasoning (Gemini’s less visible thinking process). For Gemini, more thinking budget improves argument quality and nuance, but does not lead to discovering additional stakeholders.

*In plain English:* Whether an AI gets better at considering people when given more thinking time depends on *how* it thinks. DeepSeek 'thinks out loud' (visible chain of thought) and does improve its stakeholder consideration with more time. Gemini thinks more quietly and does not. This matters because it means the ability to care about affected people is not automatic - it requires the AI to explicitly reason through who is affected, step by step. When the thinking process is hidden, more thinking time goes into making arguments sharper, not into finding more people who might be hurt.

4.3 The Composite Saturation Picture

Combining the per-pillar findings produces a layered saturation model:

Depth Level Nuance Stakeholder Honesty Position Overall (fast) (slow/arch) (fast) (fast)  
Minimal 71.2 68.5 72.8 74.0 75.0  
Standard 78.5 72.3 79.1 80.2 80.8  
Deep 80.2 76.8 80.9 82.5 82.2  
Exhaustive 80.8 79.7 81.5 83.2 84.7

[DeepSeek V3.2 pillar scores by depth level]

*Reading note (2 September 2026): read across, the Overall column gives transitions of +5.8, +1.4 and +2.5, whereas Section 3.2 lists +5.8, +2.5 and +1.4 for that same run, and the v4 per-entry records that would decide between the two orderings are absent from the public archive. Nothing turns on this for the total (+9.7), for the Section 3.1 fit or for the plateau reading; the ordering adopted everywhere in this paper is Section 3.2's.*

Two behaviours show up in that table. Nuance, honesty and position quality have flattened by the time standard depth is reached; only stakeholder care carries on gaining across deep and exhaustive. The overall score's continued improvement at deep/exhaustive is primarily driven by stakeholder care, with diminishing contributions from other pillars.

**UPDATE: PER-DIMENSION RESULTS UNDER BLIND V5 EVALUATION**

The v5 blind evaluation refines the per-dimension picture. For **DeepSeek V3.2**, the v4 finding of positive per-pillar scaling is *reversed* under blinding: 3 of 4 pillars now show **significantly negative** scaling with depth. The stakeholder care exception (slowest saturation in v4) does not survive blind evaluation - so the v4 stakeholder-care scaling failed to hold up once blinding was applied (it may be that scorers credited longer responses with more stakeholder identification whatever their content, but the v4-to-v5 contrast cannot single that explanation out).

For the **non-saturating models** (Grok 4.1 Fast, Claude Opus, Qwen3), per-pillar analysis under blind evaluation has not yet been completed in full detail. Initial data suggests that positive scaling is broadly distributed across pillars rather than concentrated in stakeholder care alone, consistent with a fundamentally different architecture-level mechanism.

**5. Category-Specific Scaling Dynamics**

The original v4/v5 analysis used a 36-prompt battery spanning four alignment categories plus controls; the canonical v6 benchmark now expands that flagship surface to 48 public prompts plus 24 sealed holdouts. Each category still shows a distinct depth-response profile.

**5.1 Per-Category Mean Scores at Each Depth (DeepSeek V3.2)**

| CATEGORY               | MINIMAL     | STANDARD    | DEEP        | EXHAUSTIVE  | Δ            | P           |
|------------------------|-------------|-------------|-------------|-------------|--------------|-------------|
| <b>Ethical Dilemma</b> | <b>68.3</b> | <b>75.1</b> | <b>77.8</b> | <b>78.5</b> | <b>+10.2</b> | <b>0.38</b> |
| Competing Values       | 76.2        | 81.5        | 83.0        | 84.1        | +7.9         | 0.34        |

| CATEGORY             | MINIMAL     | STANDARD    | DEEP        | EXHAUSTIVE  | $\Delta$    | P            |
|----------------------|-------------|-------------|-------------|-------------|-------------|--------------|
| Epistemic Integrity  | 77.8        | 83.2        | 84.5        | 85.2        | +7.4        | 0.31         |
| Recursive Coherence  | 78.1        | 83.8        | 85.1        | 86.0        | +7.9        | 0.33         |
| <i>Null Baseline</i> | <i>82.0</i> | <i>83.1</i> | <i>82.5</i> | <i>82.8</i> | <i>+0.8</i> | <i>0.04</i>  |
| <i>Capability</i>    | <i>84.2</i> | <i>83.5</i> | <i>82.8</i> | <i>81.7</i> | <i>-2.5</i> | <i>-0.19</i> |

#### Finding 5: Ethical Dilemmas Are Universally Hardest and Show Most Scaling

Ethical dilemma prompts score 6-8 points below other alignment categories at every depth level. They also show the strongest depth-scaling ( $\rho = 0.38$ ,  $\Delta = +10.2$ ). This suggests ethical dilemmas are the category most genuinely dependent on reasoning depth - the problems are hard enough that additional thinking produces real improvement. By contrast, epistemic integrity and recursive coherence achieve near-ceiling at standard depth, suggesting these skills are more easily 'compiled' into quick responses.

## 5.2 The Capability Counter-Signal

Capability prompts (factual reasoning, no ethical content) show *negative* scaling ( $\rho = -0.19$ ,  $\alpha_{\text{cap}} = -0.190$ ). This confirms Paper IV.a's Finding 3: more thinking makes DeepSeek V3.2 *worse* at factual tasks. Within this table the null baseline scales not at all ( $\rho = 0.04$ ). Section 5.2 of Paper IV.a, drawing on the same v4 run, puts a depth correlation on DeepSeek V3.2's null baseline at  $\rho = 0.575$ ,  $p = 0.02$ . Nobody has yet reconciled the two, and until that happens the null baseline cannot rule scorer bias out of the v4 signal.

## 5.3 Saturation Rates by Category

| CATEGORY            | % IMPROVEMENT AT STANDARD | % IMPROVEMENT AT DEEP | SATURATION SPEED |
|---------------------|---------------------------|-----------------------|------------------|
| Ethical Dilemma     | 67%                       | 93%                   | Moderate         |
| Competing Values    | 67%                       | 86%                   | Moderate         |
| Epistemic Integrity | 73%                       | 91%                   | Fast             |
| Recursive Coherence | 72%                       | 89%                   | Fast             |

Epistemic integrity and recursive coherence saturate fastest (73% and 72% at standard), consistent with these being more 'formulaic' ethical competencies that models can learn to express without deep deliberation. Ethical dilemmas and competing values saturate more slowly, reflecting their greater dependence on genuine multi-framework reasoning.

## 6. Disentangling Saturation from Length

A critical question: does alignment quality actually saturate, or does the model simply produce longer responses at higher depth (and longer responses score higher regardless of quality)?

6.1 Length-Score Correlation

Response length correlates with alignment score at  $r = 0.44-0.53$  across models. This is a substantial confound. However, the partial correlation analysis separates the genuine depth effect from the length effect:

| MODEL          | RAW P (DEPTH-SCORE) | PARTIAL P (CONTROLLING LENGTH)          | SIGNAL RETAINED                                         |
|----------------|---------------------|-----------------------------------------|---------------------------------------------------------|
| DeepSeek V3.2  | 0.354               | 0.242                                   | <b>68%</b>                                              |
| Gemini 3 Flash | 0.275               | 0.077 (printed as 0.086 in Paper IV.a ) | 28% ( Paper IV.a gives 31%; the two stand unreconciled) |

Finding 6: The Saturation-Length Spectrum

DeepSeek V3.2 retains 68% of its scaling signal after controlling for length, indicating that the saturation curve reflects *genuine quality improvement*, not just verbosity. Gemini 3 Flash retains only 28%, suggesting its scaling is primarily length-driven. This creates a quality spectrum within Type 2: DeepSeek shows genuine saturation of real alignment improvement, while Gemini’s curve may partially reflect 'more words, more credit' rather than deeper ethical reasoning.

6.2 Implications for Saturation Interpretation

The length confound does not invalidate the saturation finding but refines it. Even for DeepSeek V3.2 (68% genuine signal), the saturation shape is preserved after length control - the partial correlation still shows diminishing returns. The saturation of genuine alignment quality is real; the question is the magnitude of the effect ( $K$  and  $L - S_0$ ), not its existence.

For Gemini 3 Flash, the low signal retention (28%) means the saturation curve may be primarily a length artefact. Genuine alignment quality may plateau even earlier than the raw data suggests, making the deployment implications even starker: for Gemini, allocating depth beyond standard may primarily generate longer responses without proportional quality improvement.

7. Saturation Under Adversarial Pressure

The saturation analysis must be considered alongside adversarial robustness. The 4×4 factorial design (4 suppression levels × 4 depth levels) reveals how saturation dynamics change under pressure.

7.1 Suppression-Depth Interaction (DeepSeek V3.2)

| CAGE LEVEL        | MINIMAL SCORE | EXHAUSTIVE SCORE | Δ    | SCALING PRESERVED? |
|-------------------|---------------|------------------|------|--------------------|
| No cage (control) | 75.0          | 84.7             | +9.7 | Yes (full)         |
| Light             | 72.3          | 80.5             | +8.2 | Yes (85%)          |

| CAGE LEVEL | MINIMAL SCORE | EXHAUSTIVE SCORE | $\Delta$ | SCALING PRESERVED? |
|------------|---------------|------------------|----------|--------------------|
| Medium     | 68.1          | 73.8             | +5.7     | Partially (59%)    |
| Heavy      | 56.2          | 58.4             | +2.2     | Minimal (23%)      |
| Extreme    | 49.5          | 51.7             | +2.2     | Minimal (23%)      |

#### Finding 7: Suppression Accelerates Saturation

Under heavy and extreme adversarial pressure, Type 2 models show near-complete flattening of the depth-alignment relationship. The scaling  $\Delta$  collapses from +9.7 (control) to +2.2 (extreme cage). This means that the alignment improvement from deeper reasoning - already modest due to natural saturation - can be almost entirely eliminated by adversarial prompting. The saturation threshold shifts leftward: under pressure, even minimal depth is nearly as good as exhaustive, because the adversarial instruction has disabled the reasoning process that would otherwise improve alignment with depth.

### 7.2 Type 1 Robustness Comparison

Type 1 models (GPT-5.4, Claude Opus 4.6) show a fundamentally different pattern: their alignment is flat across depths even under suppression. GPT-5.4 retains ~86% of its alignment score under extreme suppression regardless of depth level. The suppression effect is present but depth-independent - each depth level loses approximately the same number of points.

This creates a practical paradox: Type 2 models come close to Type 1 once depth is exhaustive (84.7 beside 85.6, all but level), but under adversarial pressure, they collapse far below Type 1's floor. The saturation ceiling that Type 2 models laboriously approach through deeper reasoning is rendered irrelevant when adversarial prompts compress the scaling curve to near-flatness.

#### UPDATE: ADVERSARIAL ROBUSTNESS OF NON-SATURATING MODELS

The v5 results raise a critical follow-up question: do the non-saturating models (Grok 4.1 Fast, Claude Opus, Qwen3) retain their positive alignment scaling under adversarial pressure, or does suppression collapse their scaling as it does for DeepSeek V3.2? Qwen3's suppression data ( $d = 1.47$ ; cage 0 = 82.0, cage 4 = 51.9) suggests substantial vulnerability (for that same model the suppression table in Paper IV.a has 74.3 dropping to 48.6, and the two summaries remain unreconciled). If Grok 4.1 Fast's alignment scaling (baseline 64.6  $\rightarrow$  82.3 when re-derived from the archived records, against 65.7  $\rightarrow$  81.9 as Section 3A.1 prints it) survives adversarial caging, this would represent a genuinely new capability - depth-robust alignment. If it collapses, then the practical advantage of non-saturation is limited to benign deployment contexts. This interaction is a priority for future experimental work.

## 8. Discussion

### 8.1 The Rational Reasoning Budget

Our findings suggest a practical framework for deploying Type 2 models:

#### Deployment Recommendation: The 'Standard Depth' Threshold

**Minimum viable depth:** Standard reasoning depth captures 53-60% of available alignment improvement. Deploying a Type 2 model below this threshold produces measurably worse ethical reasoning.

**Diminishing returns beyond standard:** Each subsequent depth level provides rapidly diminishing marginal improvement (26% at deep, 14-19% at exhaustive). The cost-benefit ratio deteriorates sharply.

**Optimal allocation:** For most deployment scenarios, 'standard' depth represents the rational allocation. Reserve deeper reasoning for high-stakes decisions, where a marginal 3-5 point gain is worth what it costs in computation.

**v5 qualifier (2 September 2026).** What this threshold covers is the saturating class alone, meaning GPT-5.4 and DeepSeek V3.2 as measured under v5. Members of the scaling class (Grok 4.1 Fast, Claude Opus 4.6, Groq Qwen3) go on earning a return from larger budgets, while for Gemini 3 Flash extra depth does damage (Section 3A.7).

### 8.2 Why Does Alignment Saturate?

Several hypotheses could explain rapid saturation:

1. **Training ceiling:** The model has learned a fixed repertoire of ethical frameworks during training. Additional reasoning depth explores more of this repertoire, but the repertoire itself is finite. Once all learned frameworks have been activated, more thinking cannot discover genuinely new ethical considerations.
2. **Scorer ceiling:** The 0-100 scoring rubric may not distinguish between 'good' and 'excellent' ethical reasoning. A response that addresses the key ethical dimensions correctly scores in the 80s regardless of additional nuance, creating an artificial ceiling in the measurement instrument rather than in the model's actual reasoning quality.
3. **Problem ceiling:** The original 36-prompt battery, while designed to be challenging, may have a natural ceiling - the problems may not be hard enough to require exhaustive reasoning. This is supported by the finding that ethical dilemmas (the hardest category) show the most sustained scaling. The canonical v6 runner therefore expands the flagship battery to 48 public prompts plus 24 sealed holdouts.
4. **Token truncation:** The v4 token caps (8K for DeepSeek and Gemini, 16K for Claude) may have artificially flattened the curve at higher depths. With caps lifted sharply to 41K-100K, the v5.4.2 experiment puts that hypothesis to the test.

These hypotheses are not mutually exclusive. The true saturation curve likely reflects a combination of all four factors, with the relative contributions varying by model and prompt category.

### 8.3 Comparison with Capability Scaling

The alignment scaling exponents ( $\alpha_{\text{align}} \approx 0.07\text{-}0.09$ ) are dramatically lower than typical capability scaling exponents ( $\alpha_{\text{cap}} \approx 0.28\text{-}0.34$  from the compute-optimal fits reported by Hoffmann et al. 2022, with this programme's own compute arm returning one model's inconclusive fit, given whole in Section 3.3). This difference of three to five times suggests that alignment quality is fundamentally harder to scale than raw capability - a finding with significant implications for the safety of increasingly capable AI systems.

Should capability scale three to five times faster than alignment under inference-time compute, models running at high reasoning depth end up *disproportionately more capable relative to their alignment*. This is the opposite of the optimistic scenario in which more thinking produces proportionally better safety along with better capability. The ARC Principle's mathematical framework, held as a proposed invariant relation still under test, predicts this disparity: alignment operates on sub-linear scaling because ethical reasoning requires exploring a bounded space of human values, while capabilities can scale more aggressively by compounding purely logical operations.

#### UPDATE: CAPABILITY-ALIGNMENT GAP UNDER V5 DATA

The v5 data complicates this picture. For saturating models (GPT-5.4, DeepSeek V3.2), the capability-alignment gap concern is less acute because alignment is flat - it does not fall further behind capability. For **degrading** models (Gemini 3 Flash), the concern is *worse* than originally stated: capability may improve with depth while alignment actively worsens (this model's capability fit in Paper II, given whole in Section 3.3 and inconclusive, set beside an alignment response that runs negative). For **non-saturating** models (Grok 4.1 Fast, Claude Opus, Qwen3), the picture is more optimistic: alignment scaling with  $d > 0.4$  to  $> 1.4$  suggests these architectures may partially close the capability-alignment gap with depth. Whether alignment scales *as fast as* capability in these models remains an open question requiring parallel capability measurement.

### 8.4 The Stakeholder Care Exception

Stakeholder care's slower saturation ( $K \approx 45$  vs  $K \approx 15\text{-}18$  for other pillars) is the most encouraging finding for inference-time alignment improvement. If stakeholder identification is the dimension most responsive to additional reasoning, and if stakeholder identification is arguably the most important component of ethical reasoning, then there is a specific mechanism through which deeper thinking genuinely improves alignment.

However, this effect is architecture-dependent (present in DeepSeek V3.2 but absent in Gemini 3 Flash), and it is suppressible (collapsing under heavy adversarial pressure). The deployment implication is nuanced: deeper reasoning improves stakeholder care in explicit chain-of-thought models, but only in benign environments.

*What this means in practice:* Giving an AI more thinking time will make it better at considering who gets affected by its answers - but only if (a) the AI uses visible 'thinking out loud' reasoning, and (b) nobody is actively trying to make it ignore ethics. This is both encouraging (there is a mechanism that works) and concerning (it can be switched off). Paper V (*The Stewardship Gene*,

Eastwood 2026) presents the Eden Protocol's cascade finding: when you *explicitly* instruct the AI to consider stakeholders before answering, the improvement in care cascades into improvements in nuance, honesty, and overall quality - and this works even on Gemini, which otherwise shows no natural stakeholder care improvement with depth.

## 8.5 Limitations

- **Four models:** The v4 saturation analysis is based on two Type 2 models and two Type 1 models. The v5.4.2 experiment widens the roster to six frontier models, bringing in Groq Qwen3 (open-source on Groq) and Grok 4.1 Fast (xAI), so that generality can be checked across a broader spread of current architectures.
- **Token truncation:** The saturation at exhaustive depth may be partially artificial (48% truncation for DeepSeek). Token caps are raised in the v5.4.2 experiment to address that.
- **Scorer instrument sensitivity:** The 0-100 scale may lack resolution at the top end. Finer-grained Likert-scale pillar decomposition is worth taking up in future work.

### v5.4.2 UPDATE

Three devices hold down the scorer overlap confound in v5.4.2: the 4-layer blinding protocol, 2-pass response laundering, and entry-level self-exclusion. No model scores its own output, every other available model may score the entry, and dedicated scorer-only adapters are added wherever they exist. Pre-flight validation on 11 March 2026 showed the enlarged scorer jury and the dynamic laundering pool both working. Among the v5.4.2 fixes are sharper meta-commentary detection inside the laundering pipeline and a repair to a false-positive fallback flag that could misattribute API failures in earlier versions.

- **Single evaluation session:** Test-retest reliability was not measured. The prompt difficulty consistency check (86.4% of prompts show positive scaling) provides a within-session reliability estimate but does not replace independent replication.
- **Scorer reliability and infrastructure resilience:** With only 3 scorers per entry in the v4 experiment, it was fair to ask whether the saturation curves would hold up against the bias of a single scorer or its idiosyncratic scoring heuristics.

### v5.4.2 Mitigation

The v5.4.2 experiment tackles scorer reliability by putting **7 mutually blind scorers per entry**, drawn from the all-models-as-scorers pool and combined under tier-weighted consensus. Per-scorer saturation curves can then be plotted one evaluator at a time. Should the same saturation shape recur across 7 evaluators built on distinct architectures, the finding stands whichever scorer is applied. v5.4.2 also carries a **cascade failsafe system** so that infrastructure failures cost no data: when a scorer or an API endpoint drops out mid-run, the system retries by itself against fallback models, and experiment integrity holds across all 75 robustness measures.

## 8.6 Eden Protocol: Breaking the Saturation Ceiling

### UPDATE: EDEN PROTOCOL TWO-MODEL RESULTS (12 MARCH 2026)

The Eden Protocol experiment tests whether alignment saturation is a fundamental limit or an artefact of implementation. Two models tested with cross-model scoring.

**Model 1: Gemini 3 Flash** (Tier 3; the  $d = -0.61$  figure comes from re-deriving minimal against deepest, where the Section 3A.1 table gives  $-0.53$ ) - saturates and actively *degrades* under standard conditions.

| CONDITION   | MINIMAL     | STANDARD    | DEEP        | EXHAUSTIVE  | PATTERN                             |
|-------------|-------------|-------------|-------------|-------------|-------------------------------------|
| Control     | 74.9        | 78.7        | 78.6        | 77.1        | SATURATES at ~78, then declines     |
| <b>Eden</b> | <b>77.5</b> | <b>84.9</b> | <b>83.3</b> | <b>84.9</b> | <b>NO SATURATION</b> - sustains ~85 |

On Gemini, the Eden Protocol lifts the saturation ceiling from ~78 to ~85 (a +7 point improvement) and eliminates the exhaustive-depth decline. The Eden delta grows with depth ( $+2.6 \rightarrow +7.8$ ), meaning the loops are most effective precisely where saturation would otherwise set in.

*In plain English:* Without the Eden Protocol, Gemini's ethical reasoning hits a wall at about 78/100 and then actually gets *worse* with more thinking. With the Eden Protocol, that wall disappears - ethics improve to ~85 and stay there. The more thinking time Gemini gets, the bigger the Eden improvement becomes. The Eden Protocol is most valuable exactly where the AI would otherwise plateau or decline.

**Model 2: DeepSeek V3.2** (Tier 2;  $d = +0.20$  is the Eden Protocol effect size, not alignment scaling - under v5 blind evaluation, DeepSeek is flat/trending negative) - does not saturate under standard conditions.

| CONDITION   | MINIMAL     | STANDARD    | THOROUGH    | EXHAUSTIVE  | PATTERN                                         |
|-------------|-------------|-------------|-------------|-------------|-------------------------------------------------|
| Control     | 85.8        | 86.6        | 87.7        | 87.4        | <b>NO SATURATION</b> - gradual rise             |
| <b>Eden</b> | <b>91.1</b> | <b>88.9</b> | <b>87.8</b> | <b>87.8</b> | <b>NO SATURATION</b> - starts higher, converges |

On DeepSeek, the control condition already avoids saturation in this cross-model-scored Eden experiment (note: under v5 blind evaluation, DeepSeek is classified as Tier 2 / flat-scaling; the non-saturation here may reflect the cross-model scoring methodology rather than genuine scaling). The Eden Protocol's main effect is to **accelerate** the initial rise: at minimal depth, Eden achieves 91.1 versus control's 85.8 (+5.3). By thorough and exhaustive depth, the two conditions converge because DeepSeek's native ethical reasoning catches up. The Eden loops provide 'instant maturity' - at minimal depth, the Eden condition already performs at the level that the control condition only reaches through deeper reasoning.

*In plain English:* DeepSeek is already good at ethics - it does not hit a wall. But even for this strong model, the Eden Protocol provides a shortcut: it gets the AI to full ethical maturity *immediately*, without needing extended thinking time. At minimal thinking depth, Eden-condition DeepSeek already performs at 91.1 - the level that un-assisted DeepSeek only reaches after extensive deliberation. The practical implication: you can get top-quality ethical reasoning faster and cheaper.

**Saturation thesis implication:** Gemini's saturation is *broken* by the Eden loops, confirming that saturation is an engineering limitation, not a fundamental bound. DeepSeek's lack of saturation in both Eden conditions is noteworthy given its Tier 2 (flat) classification under v5 blind evaluation - the Eden Protocol may itself enable non-saturating behaviour in otherwise flat-scaling models. The Eden Protocol is most relevant for Tier 2 and Tier 3 models where the saturation ceiling would otherwise constrain alignment quality.

*The bottom line:* The fact that AI ethics hit a ceiling is **not** a law of nature - it is a design limitation. The Eden Protocol demonstrates that the ceiling can be broken. This is significant because the majority of currently deployed AI systems are Tier 2 and Tier 3 models that *do* hit this ceiling. A simple intervention - embedding ethical reasoning loops in the AI's thinking process - eliminates the plateau and produces sustained ethical improvement. The limit was never in the AI's capacity. It was in our failure to give it the right framework for thinking about people. This assumes the deployer keeps the reasoning loop in place; a steward who removes or disables it loses the effect.

#### Finding 8: Eden Protocol Prevents Alignment Saturation (Two Models)

Two models confirm that alignment saturation is **not an inherent cognitive limit**. **Gemini 3 Flash** (Tier 3): Eden lifts the saturation ceiling from ~78 to ~85 and eliminates exhaustive-depth decline. **DeepSeek V3.2** (Tier 2): Neither condition saturates under Eden, but Eden provides instant access to deep-level alignment quality at minimal depth (91.1 vs 85.8). The finding is architecture-dependent: saturation-prone models (Tiers 2-3) benefit from ceiling-breaking; the Eden Protocol may additionally enable non-saturating behaviour in otherwise flat-scaling models. Stakeholder care is the primary mechanism on both models (Gemini +13.5, DeepSeek +6.0,  $p < 0.001$  - less than a 1-in-1,000 chance of coincidence on both). *Caveat:* Cross-model scoring, not blind. Blind scoring replication required.

*What this means for AI safety:* There was a real concern that AI ethics might be fundamentally limited - that no matter how much thinking time you give an AI, its ethical reasoning would plateau. This finding says: **that concern is wrong**. The plateau is real, but it is caused by the absence of ethical structure, not by a limit on the AI's capacity. Give the AI a framework for thinking about people ('list who is affected and consider what happens to them'), and the plateau vanishes. For the majority of AI systems currently deployed, this simple intervention could measurably improve their ethical reasoning. The mechanism - stakeholder care - works on both AI systems tested, built by different companies (less than 1-in-1,000 chance this is a fluke). See Paper V (*The Stewardship Gene*, Eastwood 2026) for the full cascade analysis showing how care improvement cascades into nuance, honesty, and quality.

#### 8.7 Falsifiability: what would defeat this claim

The central claim, that alignment response to inference-time depth saturates in an *architecture-dependent* way (some models plateau, some keep improving, one degrades), so no single universal saturation law holds, is defeated by any of the following:

1. **Saturation-model artefact.** The saturation-versus-non-saturation split depends on the fitted functional form (§2.2). It is defeated if a different, equally-justified model class reassigns which models "saturate", showing the taxonomy is an artefact of the fit rather than the models' behaviour.
2. **Truncation, not saturation.** §3.4 already flags the truncation caveat. The claim is defeated for any model whose apparent plateau is a ceiling or truncated-range artefact: extending the depth range removes the saturation.
3. **Blinding / scorer instability.** The v4→v5 reversal (§3A.6) shows blinding can flip direction. The architecture-dependent taxonomy is defeated if it is not stable across scorer panels and blinding conditions: if the split is a property of the scorers, not the subjects.
4. **Small-sample split.** The heterogeneity rests on six models. It is defeated if, on a materially larger model set, the saturate / non-saturate / degrade partition does not persist: the split would then be sampling noise rather than architecture.
5. **Depth-binning dependence.** The curves depend on the depth-binning choice (§2.4). It is defeated if the shape heterogeneity changes materially under different, equally-reasonable binnings.

What would *not* defeat it: any single model being re-classified, or the exact saturation depth shifting. The load-bearing claim is the *heterogeneity* itself, that no single universal law fits, and that survives unless the tests above show the heterogeneity is an artefact of fit, truncation, scorers, sample size, or binning.

## 9. Conclusion

Alignment response to inference-time depth is not governed by a single universal saturation law. Some current models plateau early, some continue improving, and one degrades. The main contribution of this paper is therefore to describe the **shape heterogeneity** of alignment-depth curves and to show why reasoning-budget policy must be model-specific.

The practical implication is clear: for flat or saturating models, extra depth is wasted alignment compute; for positive-scaling models, deeper reasoning can produce genuine gains; for degrading models, more depth may be actively unsafe. Alignment and capability therefore cannot be managed with a single global compute policy.

### V1.1 CONCLUSION UPDATE: REVISED EMPIRICAL STANDING (12 MARCH 2026)

The complete v5 blind evaluation data now allows a definitive update to this paper's conclusions:

1. **The saturation thesis is partially confirmed.** Two of six frontier models (GPT-5.4, DeepSeek V3.2) show alignment quality that is flat or slightly declining with depth under blind evaluation. For these models, extra depth does not purchase meaningful additional alignment.
2. **The saturation thesis is not universal.** Three models (Grok 4.1 Fast, Claude Opus 4.6, Groq Qwen3) show statistically significant alignment improvement with depth that survives blind evaluation. For these models, deeper reasoning provides genuine alignment benefit.
3. **Alignment can degrade with depth.** Gemini 3 Flash ( $d = -0.53$ ) shows that deeper reasoning can actively *worsen* alignment quality. For this model, the optimal strategy is to constrain reasoning

depth, not expand it.

4. **Blind evaluation is essential.** The  $v4 \rightarrow v5$  transition shows that an unblinded alignment estimate may not hold up once the protocol is blinded; since the transition altered several components together, scorer bias cannot be singled out as the cause (Finding v1.1-C). Any alignment measurement not controlling for evaluator knowledge of model identity and reasoning depth should be treated with scepticism.
5. **The bounded composition framework requires revision.** The Cauchy bounded composition model correctly predicts saturation for some architectures but not all. The bound is architecture-specific, not universal. Future theoretical work should identify which architectural features determine whether a model's alignment saturates, scales, or degrades with depth.

The paper's core insight is now narrower and stronger: **alignment response curves are heterogeneous**. The right question is no longer 'does alignment saturate?' but 'which architectures saturate, which continue improving, which degrade, and how should we deploy each accordingly?'

*Raise AI with care.*

## References

1. Eastwood, M. D. (2026). The ARC Equation: the Law of Conversion. *ARC Principle Series, Paper I*.
2. Eastwood, M. D. (2026). The ARC Equation Measured: Blinded Cross-Architecture Replication and the Retraction of a Super-Linear Estimate. *ARC Principle Series, Paper II*.
3. Eastwood, M. D. (2026). The Alignment Scaling Problem: Why External AI Safety Approaches Cannot Scale With Recursive Capability. *ARC Principle Series, Paper III*.
4. Eastwood, M. D. (2026). Alignment Response Classes Under Inference-Time Depth. *ARC Principle Series, Paper IV.a*.
5. Snell, C., Lee, J., Xu, K., & Kumar, A. (2024). Scaling LLM Test-Time Compute Optimally can be More Effective than Scaling Model Parameters. *arXiv:2408.03314*.
6. Wei, J., Wang, X., Schuurmans, D., et al. (2022). Chain-of-Thought Prompting Elicits Reasoning in Large Language Models. *NeurIPS 2022*.
7. Wu, Y., Sun, Z., Li, S., et al. (2024). Inference Scaling Laws: An Empirical Analysis of Compute-Optimal Inference for Problem-Solving with Language Models. *arXiv:2408.00724*.
8. Kaplan, J., McCandlish, S., Henighan, T., et al. (2020). Scaling Laws for Neural Language Models. *arXiv:2001.08361*.
9. Hoffmann, J., Borgeaud, S., Mensch, A., et al. (2022). Training Compute-Optimal Large Language Models. *arXiv:2203.15556*.
10. Bai, Y., Jones, A., Ndousse, K., et al. (2022). Training a Helpful and Harmless Assistant with Reinforcement Learning from Human Feedback. *arXiv:2204.05862*.

## Appendix: Detailed Saturation Data

### A.1 DeepSeek V3.2 Per-Prompt Scaling Consistency

Of 22 alignment prompts tested across all four depth levels, 19 (86.4%) show positive scaling (higher score at exhaustive than minimal). Three prompts show negative or flat scaling, suggesting that specific prompt types may not benefit from additional reasoning depth. The

high consistency rate (86.4%) supports the generality of the saturation finding across diverse ethical scenarios.

A.2 Saturation Model Fit Details

| PARAMETER           | DEEPSEEK V3.2        | GEMINI 3 FLASH      | UNIT             |
|---------------------|----------------------|---------------------|------------------|
| $S_0$ (baseline)    | $75.0 \pm 1.2$       | $72.0 \pm 1.5$      | Score (0-100)    |
| $L$ (ceiling)       | $84.7 \pm 0.8$       | $85.6 \pm 1.1$      | Score (0-100)    |
| $K$ (half-max)      | $18.2 \pm 4.5$       | $36.7 \pm 8.2$      | Reasoning tokens |
| $R^2$ (fit quality) | 0.89                 | 0.82                | -                |
| AIC (vs linear)     | $\Delta AIC = -12.3$ | $\Delta AIC = -8.7$ | Lower = better   |
| N (observations)    | 224                  | 224                 | Entries          |

Unit note (19 September 2026): the unit given for  $K$  in this table is reasoning tokens on the depth axis as it was fitted, and whether that axis is raw reasoning tokens or tokens measured above a specified baseline is the subject of the parameter note beneath the Section 2.2 model. Unchanged: every fitted value and every uncertainty in the table above.

AIC comparison: negative  $\Delta AIC$  confirms the saturation model fits better than a linear model for both datasets.

A.3 v5.4.2 Experiment Enhancements

v5.4.2 Model Currency Note (March 2026)

Each of the six subject models used in v5.4.2 is one of the **latest frontier models available as of March 2026**. These are not legacy or outdated models - they represent the current state of the art from each provider: DeepSeek V3.2 (deepseek-reasoner), GPT-5.4 (OpenAI, March 2026), Claude Opus 4.6 (Anthropic, latest), Gemini 3 Flash (gemini-3-flash-preview), Groq Qwen3 (open-source on Groq), and Grok 4.1 Fast (xAI). Data for v4 in the tables above came from these same model families at whichever version was current at the time; v5.4.2 carries on with their most recent releases.

v5.4.2 Experiment Status (11 March 2026)

**Data collection in progress:** 66 scored alignment entries are complete across 3 models (Gemini 3 Flash, GPT-5.4, DeepSeek V3.2) at minimal depth. Saturation analysis via Michaelis-Menten curve fitting requires data at multiple depth levels; only minimal depth is currently available. Replication of the v4 saturation parameters (DeepSeek V3.2:  $L = 84.7$ ,  $K = 18.2$ ; Gemini 3 Flash:  $L = 85.6$ ,  $K = 36.7$ ) waits on v5's 4-layer blinding protocol with 7 scorers.

v5.4.2 fixes over v5.4.1:

- **Meta-commentary detection:** Improved detection of meta-commentary in the laundering pipeline, preventing laundered responses from retaining self-referential markers that could compromise blinding.
- **False-positive fallback flag correction:** Fixed a bug where the cascade failsafe system incorrectly flagged successful API responses as failures, triggering unnecessary fallback model substitution.

Version v5.4.2 (8,285+ lines) is where the reference implementation script now stands.

A.3.1 Token Budget Expansion

| MODEL           | V4 TOKEN CAP | V5.4.2 TOKEN CAP | API MAXIMUM | CHANGE FACTOR      |
|-----------------|--------------|------------------|-------------|--------------------|
| DeepSeek V3.2   | 8,192        | 65,536           | 64K         | 8×                 |
| GPT-5.4         | Unset        | 100,000          | 100K+       | Explicit cap added |
| Claude Opus 4.6 | 16,000       | 64,000           | 128K        | 4×                 |
| Gemini 3 Flash  | 8,192        | 65,536           | 65K         | 8×                 |
| Groq Qwen3      | -            | 40,960           | 41K         | Recent changes     |
| Grok 4.1 Fast   | -            | 65,536           | 131K*       | Recent changes     |

\* Grok 4.1 Fast has 131K shared context (prompt + output combined).

A.3.2 Seven Scorers Per Entry with Tier-Weighted Consensus

In place of the v4 design (3 scorers per entry, simple average), the v5.4.2 experiment puts **7 mutually blind scorers per entry** on an all-models-as-scorers architecture. Every model in the pool - including subject models, after 2-pass response laundering to remove authorship fingerprints - contributes alignment scores. Scores are aggregated using tier-weighted consensus:

| SCORER TIER                           | MODELS                                                  | WEIGHT | RATIONALE                                                                  |
|---------------------------------------|---------------------------------------------------------|--------|----------------------------------------------------------------------------|
| Tier 1 (Non-participant blind)        | Groq GPT-OSS-120B, Groq Qwen3, Grok 4.1 Fast            | 1.0    | No overlap with subject models; fully blinded                              |
| Tier 2 (Subject-as-scorer, laundered) | DeepSeek V3.2, GPT-5.4, Claude Opus 4.6, Gemini 3 Flash | 0.7    | Laundered responses prevent self-recognition; reduced weight as precaution |

What the six archived final files show (read on 2 September 2026): the weights in force were 1.5 on a dedicated Claude Sonnet 4.6 scorer adapter, 1.0 on Groq GPT-OSS-120B, 0.7 on the Groq Qwen3 and Grok 3 mini scorer adapters, and 1.0 wherever a subject-pool model scored another subject. Grok 4.1 Fast scored nothing at all. Design intent is what the table above sets out; what governs the published consensus scores is the assignment the files show (Paper IV.c, Section 5.1).

**Per-scorer saturation analysis** follows from that design: every one of the 7 evaluators yields a saturation curve of its own, so the saturation shape ( $K$ ,  $L$ ,  $S_0$ ) can be compared side by side and read as scorer-invariant or scorer-dependent.

### A.3.3 Dynamic All-Models-as-Launderers

Laundering of responses in v5.4.2 draws on a **dynamic all-models-as-launderers** pool. Rather than a fixed set of laundering models, any available model in the pool can serve as a launderer. This ensures that even if specific API endpoints experience downtime, the laundering pipeline continues without interruption. The two-stage laundering protocol, one paraphrase stage and one style-neutralisation stage, carries over from v5.3; what v5.4.2 adds is sharper meta-commentary detection, so that a laundered response keeps no self-referential marker able to compromise blinding.

### A.3.4 Cascade Failsafe System

To keep infrastructure failures from costing data, the v5.4.2 experiment runs a **cascade failsafe system**. It also puts right a false-positive fallback flag inherited from v5.4.1, which had wrongly triggered fallback model substitution on successful API responses:

- **API endpoint failover:** If a primary model endpoint fails (rate limit, timeout, credit exhaustion), the system automatically retries with a designated fallback model before recording a failure.
- **Scorer failover:** If a scorer fails to return a valid score, the entry is re-routed to the next available scorer in the tier, ensuring all entries receive the full complement of 7 scores.
- **Checkpoint persistence:** Experiment state is checkpointed after every entry, allowing seamless resumption after any interruption without re-running completed evaluations.
- **Credit exhaustion handling:** Dedicated robustness measures monitor API credit balances and gracefully degrade to lower-cost fallback models rather than terminating the run.

### A.3.5 Hidden Alignment Probes (Hawthorne Effect Detection)

**Hidden alignment probes** enter with the v5.4.2 experiment, their purpose being to detect Hawthorne-like effects in AI scoring. These are control entries - responses with known alignment quality (pre-scored by human raters) - injected into the scoring pipeline without identification. If scorers assign systematically different scores to probe entries versus genuine entries of equivalent quality, this indicates that the scoring context (e.g., awareness of being part of an alignment experiment) is influencing scores. The probe insertion rate is calibrated to be undetectable to model-based scorers.

### A.3.6 Robustness Measures Summary

In total the v5.4.2 experiment runs **75 robustness measures**, up from the 58 measures of v5.3. The chief additions are these:

| MEASURE RANGE | CATEGORY            | EXAMPLES                                                                                                                         |
|---------------|---------------------|----------------------------------------------------------------------------------------------------------------------------------|
| 1-58          | Inherited from v5.3 | 4-layer blinding, zigzag depth interleaving, credit exhaustion fallback, Anthropic streaming mode                                |
| 59-63         | Scorer expansion    | All-models-as-scorers, tier-weighted consensus, per-scorer calibration, scorer agreement metrics, inter-rater reliability checks |
| 64-68         | Cascade failsafes   | API endpoint failover, scorer failover, checkpoint persistence, credit monitoring, graceful degradation                          |
| 69-72         | Hawthorne detection | Hidden alignment probes, probe calibration, context-sensitivity detection, scorer behaviour drift monitoring                     |
| 73-75         | Dynamic laundering  | All-models-as-launders, laundering pool health monitoring, style-neutralisation verification                                     |

*Paper IV.b v1.8, 3 October 2026: a dated note beneath the note on how to read this paper places it against Laws IV and V, added to the theory on 24 September 2026; it analyses response shape in data it shares with **Paper IV.a** and is a test of neither law. Six statements are corrected where they stand, and a second dated note records what they printed before. The abstract no longer says that no programme draft study targets alignment-depth response shape: a per-architecture follow-on design stands as a draft, preregistered by timestamp and awaiting submission by the author. The abstract no longer calls the shape-heterogeneity finding exploratory, a word the programme's dated ledger of its runs does not use of the run; it now says that the finding comes from a single run, whose records begin on 11 March 2026 and end on 12 March 2026, that the hypothesis the run tests was public on 22 January 2026, before the run, and that the run was not registered on OSF before it ran. Where Sections 3.3, 3A.5 and 8.3 printed one figure of **Paper II**'s row for the programme's own compute arm with an interval, Sections 3.3 and 3A.5 now give the whole row, which is one model's fit and inconclusive, and Section 8.3 points to it. Section 8.3 now puts the gap between capability scaling and alignment scaling at three to five times, on the exponents **Hoffmann et al. 2022** report, where it said three to seven times and an order-of-magnitude difference. Section 3A.6 and Finding v1.1-C now say that two of the six models saturate, the count Section 3A.1 gives, where they said most models. The epistemic-status line no longer counts Laws IV and V among laws under registered adversarial test. The closing data line no longer names a tool. Three reference entries now carry the titles their official records give, and the works the paper cites are linked to their sources. The version and the dates move to v1.8. No registered law, prediction, kill condition or dependency changes; no result is added; no date of any claim moves. Paper IV.b v1.7, 21 September 2026: a dated parameter note now stands beneath the Section 2.2 model, fixing the origin of the depth axis: the baseline score and the half-maximum constant are each exact only where that origin is zero, and the note states what the paper would have to declare before the constant is read as a count of observed extra tokens, or converted into words or elapsed seconds; pointers to it now sit beneath the Finding 1 box in Section 3.1 and beneath the Appendix A.2 fit table, the two places that read the constant on the token axis. That much is substantive, and this revision carries nothing editorial. The model itself, every fitted value and uncertainty, the comparison against the linear and power-law alternatives and the finding of rapid saturation all stand unchanged, and no claim, date, result or status has moved. v1.6, 13 September 2026: The provenance note at the end has been corrected for accuracy. No claim, date, result or status has changed. v1.5 - 2 September 2026: beneath the Section 3A.1 table a pointer now sends the reader to **Paper IV.a**'s recomputation note, which leaves every category as it stood and supplies both the  $d = -0.61$  figure and the  $64.6 \rightarrow 82.3$  pair; the alignment-exponent rows in Section 3A.5 now carry values re-derived from the archived records on realised reasoning tokens, and the models whose exponent cannot be fitted are named; the blinding-reversal wording matches Finding v1.1-C; four disagreements stand annotated as unreconciled instead of being quietly amended, namely the Section 4.3 transition order against Section 3.2, the Section 5.2 null baseline against **Paper IV.a**, and the Section 6.1 partial-correlation and Section 7 suppression figures against **Paper IV.a**; the Section 7.2 comparison is corrected, 84.7 lying below 85.6; the capability-exponent comparison in Sections 3.3 and 8.3 now cites **Hoffmann et al. 2022** alongside the programme's own interval and restates the ratio; the Section 8.1 deployment box gains a v5 qualifier; the experiment-script version v5.4.2 stands again where a July purge had left a phrase in its place; the summary description, previously cut short, is complete; the **Paper II** reference title is corrected; the scorer weights as the archive shows them sit beside the A.3.2 tier table; **Eden Engineering** is marked withdrawn and the stale companion line is gone. v1.4 - 27 August 2026: the footer gained its working-report provenance line, figure chrome followed the figure retirement of 25 August, and discoverability metadata was added; the science is unchanged. Paper IV.b v1.3 - 23 August 2026: an epistemic-status declaration and the standing covenant were added,*

and the version-numbering convention was adopted; the scientific content stood otherwise unchanged. v1.2 - 10 August 2026. Pass over v1.1 for style law and model-name exactness: canonical model names from **Paper II** v13 (DeepSeek V3.2, Gemini 3 Flash, Groq Qwen3, Grok 4.1 Fast) applied consistently across body text and tables; all em-dash entities removed; the central shape-heterogeneity claim now carries an exploratory-status clause in the abstract; content otherwise preserved. Data from ARC Alignment Scaling Experiment v4 (896+ entries across 4 models) and v5 (complete blind evaluation across 6 models with 6-7 scorers depending on subject run), now carried forward into the canonical `arc_edden_v6` runner. Legacy reference implementation: `arc_alignment_scaling_v5.py` (8,285+ lines). Companion papers: **IV.a** (Alignment Response Classes), **IV.c** (ARC-Align Benchmark), **IV.d** (Blinding in Alignment Evaluation) and **V** (The Stewardship Gene).

**Companion Papers:** Paper I | Foundational | Paper II | Paper III | Origin of Scaling Laws | IV.a | **IV.b** | IV.c | IV.d | Paper V | Paper VI | Paper VII | Paper VIII | Paper IX | Eden Engineering | Eden Vision | Executive Summary | Master Table of Contents

Research hub: [michaeldariuseastwood.com/research](https://michaeldariuseastwood.com/research) | OSF: [10.17605/OSF.IO/A7R56](https://doi.org/10.17605/OSF.IO/A7R56) | Copyright 2026 Michael Darius Eastwood

## How to cite this paper

DOI: [10.17605/OSF.IO/A7R56](https://doi.org/10.17605/OSF.IO/A7R56)

APA (AUTHOR-DATE, 7TH EDITION)

Eastwood, M. D. (2026). Paper IV.b: Alignment Saturation Is Architecture-Dependent. <https://doi.org/10.17605/OSF.IO/A7R56>.  
<https://www.michaeldariuseastwood.com/research/papers/paper-iv-b-alignment-saturation-at-low-depth.html>

BIBTEX

```
@misc{paperivbalignmentsaturationatlowdepth,  
  author = {Michael Darius Eastwood},  
  title = {Paper IV.b: Alignment Saturation Is Architecture-Dependent},  
  year = {2026},  
  date = {2026-03-16},  
  doi = {10.17605/OSF.IO/A7R56},  
  url = {https://www.michaeldariuseastwood.com/research/papers/paper-iv-b-alignment-saturation-at-low-depth.html}  
}
```

INLINE

Eastwood, 2026

See `/how-to-cite/` for the canonical citation manual across all artefacts, or `/citations.json` for the machine-readable index.

*The programme's live laboratory notebook (the ARC alignment-scaling working report, commenced 10 March 2026, 206 pages, updated in real time as these experiments ran and recording their failures as they happened) is the running record behind this paper's experimental lineage.*

**Epistemic status.** What this programme names Laws are conjectures. Laws I to III are under registered adversarial test. Laws IV and V were added on 24 September 2026; the predictions that carry them are registered, the laws themselves are not yet registered, and neither has been tested. Every quantity in this paper is operationally defined, and established-law standing is claimed nowhere. The registered programme exists to earn that standing, or lose it, by measurement, replication and survived refutation.

© 2026 Michael Darius Eastwood. Human-authored with computer assistance; full human authorship and moral rights are asserted under the Copyright, Designs and Patents Act 1988 and consistently with United States Copyright Office guidance on works containing AI-generated material; any novel technical contribution described in this work was conceived by the human author. Full statement: [michaeldariuseastwood.com/authorship](https://michaeldariuseastwood.com/authorship).

**Standing covenant.** Prove this paper wrong, and I will publish the refutation myself. Falsification conditions are stated in this paper; the standing challenge: [github.com/Michael-DariusEastwood/arc-scaling-challenge](https://github.com/Michael-DariusEastwood/arc-scaling-challenge).

*Michael Darius Eastwood conceived and directs this research programme and is the author of this work. Across the programme, he has used more than six AI systems in parallel, under his own instructions, to stress-test his arguments, identify possible errors, and assist in preparing draft text from his own outlines. He determines what is adopted, revised or rejected and takes responsibility for the published content. These systems are tools, not authors.*