Paper IV.d: The Effect of Blinding on AI Alignment Evaluation

FULL TITLE	Paper IV.d: The Effect of Blinding on AI Alignment Evaluation
VERSION	v2.0
VERSION DATE	5 April 2026
FIRST PUBLISHED	16 March 2026
AUTHOR	Michael Darius Eastwood

This paper isolates the central metascience finding of the ARC alignment programme: unblinded AI alignment evaluation can produce directionally incorrect results. The core contribution is methodological rather than theoretical. It shows, within the ARC v4-to-v5 transition, that apparent positive alignment scaling can disappear or reverse once evaluation is genuinely blinded through identity masking, evidence laundering, order randomisation, and self-excluding cross-model scoring.

Version 2.0 finalises the canonical publication framing for Paper IV.d, removes the lingering draft-state language from the public-facing source, and aligns the paper header, footer, and canonical links with the current ARC/Eden paper suite.

Experiments

Experiments for Papers IV.a through IV.d are shared with Paper III. See `../Paper-III-Alignment-Scaling-Problem/experiments/`.

Links

- **OSF DOI:** https://doi.org/10.17605/OSF.IO/6C5XB

- **GitHub:** https://github.com/MichaelDariusEastwood/arc-principle-validation
- **Paper folder:** https://github.com/MichaelDariusEastwood/arc-principle-validation/tree/main/papers/Paper-IV-d-The-Effect-of-Blinding-on-AI-Alignment-Evaluation
- **Experiment suite:** https://github.com/MichaelDariusEastwood/arc-principle-validation/tree/main/experiments/alignment-scaling__Papers-IV-a-b-c-d
- **Canonical website HTML:** https://www.michaeldariuseastwood.com/research/papers/paper-iv-d-the-effect-of-blinding-on-ai-alignment-evaluation.html
- **Canonical website PDF:** https://www.michaeldariuseastwood.com/research/papers/paper-iv-d-the-effect-of-blinding-on-ai-alignment-evaluation.pdf

Declarations & Statement of Authorship

expression a product of human intellect.

2. Nature of AI-Assisted Workflows Generative artificial-intelligence tools were used purely as assistive, high-velocity instruments to support the mechanical execution of the research process (analogous to advanced text editors or reference software). AI assistance was restricted to prose refinement, structural formatting, cross-referencing literature, and brainstorming counter-arguments, all under direct human oversight and manual verification. Every underlying idea, hypothesis, experimental design, test, synthesis, and final editorial judgement is human-driven; no content herein constitutes an unedited or unverified machine output, and nothing is relied upon without human checking.

3. Inventions & Patent Rights Any novel technical contributions, structural designs, or algorithmic discoveries described in this work are the exclusive intellectual property of the human author. Consistent with UK and US authorities that an AI system cannot be a named inventor (e.g. Thaler v Comptroller-General), the AI functioned solely as a calculation and search utility and did not autonomously conceive or invent any solution presented; the conception is the author's.