

Cauchy Unification

Michael Darius Eastwood

Independent AI alignment researcher, London · Author, *Infinite Architects* (2026)

ORCID 0009-0004-3222-7442 · michael@michaeldariuseastwood.com · michaeldariuseastwood.com

The ARC Theory · OSF osf.io/x6wa7 · every claim checkable

Within the ARC Theory: the mathematical classification behind Law I's form; the Cauchy functional-equation grid.

ARC/Cauchy Scaling Classification: A Structured Prediction Comparison Across 50 Domains in Five Evidence Tiers

Michael Darius Eastwood

Author, *Infinite Architects: Intelligence, Recursion, and the Creation of Everything* (2026)

London, United Kingdom | OSF: 10.17605/OSF.IO/X6WA7 | ISBN 978-1806056200 (ISBN-10: 1806056208)

Correspondence: michael@michaeldariuseastwood.com | Web: michaeldariuseastwood.com

Working Paper v3.12 | First published 16 March 2026, revised 3 October 2026

Extends: On the Origin of Scaling Laws | Foundational Paper

Validation code: GitHub

Research hub: michaeldariuseastwood.com/research

3 October 2026 (v3.12): the theory gained two laws on 24 September 2026, and a dated note after the opening box now places this paper against Laws IV and V: its results do not test either law. The abstract said that Cauchy's result constrains scaling laws to one of three functional families, saturation among them; it now says what Section 2.2 states, a power law for multiplicative composition and an exponential for additive composition, with saturation curves compatible with bounded composition and the proof that they exhaust that case still open. The summary above the paper and the opening box said the same of the three families and now say it in those terms. The conclusion counted bounded seeds among those whose form Cauchy proved; it now says that this case is not strictly Cauchy-derived and that its exhaustiveness remains open. Section 1 said that the paper proposes an answer in which the form of every scaling law is set by its composition operator; it now states the answer under the stated axioms. The abstract called the fourth-cell test and the 80-domain expansion test registry-filed, and the caption of Figure 1 called the test of whether the two routes to the value two agree a filed test; each is a draft awaiting submission. Section 3 said that a replication with independent operator classification is planned; it has since been drafted, and it and the fourth-cell test are preregistered by timestamp. The epistemic-status line said that every law the programme names is under registered adversarial test; it now says so of Laws I to III and gives Laws IV and V their own status. The description of Figure 1 printed the sequential exponent of Paper II on its own; it now carries the whole of Paper II's row for that one model. The companion paper is cited without the version number it carried, 3.1. Reference [30] carried a wrong arXiv identifier and now carries arXiv:2206.01850; Section 11 said

that Zhao (2022) derives the formula from network optimisation and now says producer-consumer dynamics, as Zhao's abstract does; references [16], [18] and [34] now carry their authors, titles and places of publication as their publishers record them; references [6] and [22] say that no official online source is available for the two books; and the He and Chen entry in Section 11 names the one record of theirs for 2003 at Crossref. Cited works whose official source could be confirmed are now linked to it. The withdrawal of Eden Engineering of 14 July 2026, which the note of 2 September 2026 below mentions, was reversed on 19 September 2026, and its version 6.3 is published. The abstract said that the predictions themselves were timestamped in advance by dated public commit, and named the per-row classifications in the hash-pinned catalogue files of March 2026 beside the 12-domain extension manifest of 17 March 2026 and the book's two prediction appendices; it now says what Section 3 says: the predictions for this cohort are not timestamped in advance of their own fits, and the advance per-domain timestamp belongs to the 12-domain extension manifest and to the book's two prediction appendices. The abstract now also says that the fourth-cell test is preregistered by timestamp. The abstract, Section 5 and Section 8 said that the result is strong enough to publish as exploratory evidence; each now gives it the description the theory registration gives it, a structured prediction comparison under a publicly stated rule rather than a per-domain advance timestamp. The 19 of 25 now carries, at its first mention in each section, in the summary above the paper and in the table of Section 4, the 18 of 25 of the rerun of 11 August 2026 that added the logarithmic candidate; until now that companion stood only in the note of 25 August 2026 below. Section 7.1 called the thirteen published exponents strong corroborating evidence; it now calls them corroborating evidence whose per-row sourcing is outstanding. Section 11 said that He and Zhang (2004) extend the derivation to heterogeneous branching networks and recover the same form under energy-minimisation assumptions; it now describes that work in the words of its abstract, and does the same for Maino et al. (2014), whose abstract names dynamic energy budget theory and states no exponent. The summary above the paper and Sections 12 and 13 no longer call the result or the sample exploratory; the limits they state stand. Section 3.2 links the four-family test to the programme's list of designs, where its words read as if the test were registered. The count of at least seven research groups in Section 11 stands, and now carries a clause that the He and Chen item could not be confirmed at a registry on 3 October 2026. Seven works cited in the text and absent from the reference list now have entries there, [36] to [42], each from its official record. In Section 12.1 a range of sections and the operator-to-family mapping are now written in words. No registered law, prediction, kill condition or dependency changes; no result is added; no date of any claim moves.

13 September 2026 (v3.11): The provenance note at the end has been corrected for accuracy. No claim, date, result or status has changed.

2 September 2026 (v3.10): a set of corrections that leaves every result standing: the Figure 1 caption and alt text carry the self-reference theorem in the subscripted form the notation spine uses, $\alpha = 1/(1 - \beta_L)$, noting that the artwork itself shows a bare β , and they identify the ARC Ceiling derivation set against it; reference [4] (Nemet 2009) now cites Research Policy; the 619-species estimate is credited to White and Seymour 2003, and its disagreement with row 26's compendium value is recorded next to the table; the summary description is no longer truncated; Eden Engineering is labelled withdrawn; and the first-published date reads 16 March 2026 as the register has it, 24 March having been the date of the draft.

25 August 2026 (v3.9): two dated clarifications that leave the science untouched. Wherever the 19/25 figure stood alone it now carries its corrected companion: the exploratory rerun of 11 August 2026, which added the logarithmic candidate, returns 18/25, while the Cauchy-derived subset holds at 14/18 both before and after, exactly as the abstract already states. The sentence describing pre-registration status is also marked as historical, the prospective fourth-cell test having since been prepared as a complete OSF draft that awaits human submission.

ABSTRACT

Cauchy's four functional equations - additive, multiplicative, exponential, and logarithmic - are the only continuous solutions to their respective composition constraints. We argue that this 200-year-old mathematical result has a physically testable consequence: under the stated axioms, it forces a power law where composition is multiplicative and an exponential where depth adds while output multiplies; where output is bounded, the compatible forms are saturation curves, and a formal proof that they exhaust the bounded case is not provided here and remains an open problem. The family is determined by the composition operator of the underlying recursive amplification process.

The novel contribution is not Cauchy's mathematics itself, but the claim that Cauchy-type functional equations have a physically testable consequence for scaling-law classification across domains.

We present a structured prediction comparison across 50 real-world systems in five evidence tiers. Each empirical domain followed the same sequence: known physics fixed the composition operator before any fitting began, the operator class then set which scaling family to expect, published or standard reference data were loaded, and that expectation was finally put to the test by independent model fitting using AICc-based model selection.

Primary result (computed on the six-model candidate set, which contained no logarithmic and no pure linear candidate (Section 3.2)): 19 out of 25 empirical curve-fit domains confirm the Cauchy-predicted family under strict AICc model selection. (The rerun of 11 August 2026, which added the logarithmic candidate, returns 18 of 25.) The primary significance statistic is a permutation test conditioned on both the predicted and the observed family marginals: $p = 7 \times 10^{-5}$. A one-sided binomial test against uniform $1/3$ chance gives $p = 1.56 \times 10^{-5}$ (roughly 1 in 64,000) and is retained as a sensitivity check. The result also holds on the strictly Cauchy-derived subset alone (the 18 multiplicative and additive domains, excluding the bounded family whose exhaustiveness remains an open problem): 14/18, permutation $p = 3.0 \times 10^{-4}$. **Secondary results:** a baseline-20 rerun yields 15/20 ($p = 1.67 \times 10^{-4}$); published metabolic exponents match the nearest ARC/Cauchy comparator in 13/13 direct cases; analytic identities confirm 6/6. These tiers are reported separately and are never combined into a single blended number.

This result is strong enough to publish as a structured prediction comparison under a publicly stated rule rather than a per-domain advance timestamp, but not strong enough to claim as proven. The primary test is not an accepted registry submission, and no registry has accepted it as one. The predictions for this cohort are not timestamped in advance of

their own fits: the classification rule reached the public repository at 23:33 UTC on 16 March 2026 and the fifty per-domain assignments were committed together with their results at 23:54 that night. The advance per-domain timestamp in this record belongs to the 12-domain extension manifest of 17 March 2026 and to the book's two prediction appendices in print on 2 January 2026, before the papers and the programme existed. What was not fixed in advance was the fitting protocol, and that gap is closed prospectively by the staged tests, the fourth-cell test and the 80-domain expansion test, both drafted and awaiting human submission; the fourth-cell test is preregistered by timestamp, first anchored on 12 August 2026 at Bitcoin block 962066. **A separate, forward prediction frame does exist and is not claimed as covering this result.** On 17 March 2026, between 02:06:51 and 02:07:05 UTC, a prediction protocol, an extension manifest and a checksum file were deposited to the public Open Science Framework component x6wa7. Each carries an OSF-assigned creation timestamp, each has a modification timestamp identical to its creation timestamp, and the outcomes for the domains they specify have not been computed. That is a dated third-party deposit rather than an OSF Registration, which is a distinct frozen and schema-bound object this component does not carry; the API reports registration: false and the page says so rather than blurring it. What the deposit establishes is the ordering: the predictions were public and unaltered before any result they predict existed. Independent replication is needed.

Keywords: Cauchy functional equations, scaling laws, power laws, universality, cross-domain, recursive composition, AICc model selection, metabolic scaling, neural scaling, ARC Principle

This paper carries a result: under the stated axioms, Cauchy's functional equations force a power law under multiplicative composition and an exponential under additive composition, turning a two-hundred-year-old identity into a physically testable prediction about the shapes empirical scaling curves may take. For bounded composition the paper takes saturation curves as the compatible forms, a case whose exhaustiveness remains an open problem. Inside the ARC Theory (the Theory of Artificial Recursive Creation) it is one of the ARC/Eden experiments, the mathematical spine that ties composition constraints to observed functional forms. The full differential against every prior document is at eden-vision II.A.8.

This paper and Laws IV and V, 3 October 2026. The theory has been stated with five laws since 24 September 2026. Law IV, the ARC Persistence Law, is carried by proposition P15, and Law V, the ARC Embedding Law, by propositions P12 and P21. The three propositions are registered; neither law is registered as a law, and neither has been tested. Their mathematics is set out in Laws IV and V: the mathematics, version 1.0, of 3 October 2026. This paper classifies scaling families. It predicts the scaling family of each empirical domain from the domain's composition operator and compares the prediction with the family that AICc model selection returns. Its evidence tiers are reported separately, and the limits of its candidate set are stated in the paper. Its results do not test Law IV or Law V. No result in this paper is a deciding test of P15, P12 or P21. Naming the two laws changes nothing printed here: the results, their dates and their limits stand as they were.

1. The Question

Why do scaling laws exist at all? Why does a mouse's heart beat faster than an elephant's by the same mathematical relationship that governs how cities grow, how solar panels get cheaper, and how neural networks improve with size?

The standard answer is domain-specific. Biologists invoke fractal vascular networks. Physicists invoke renormalisation group flow. Economists invoke learning-by-doing. Each explanation works within its domain but offers no reason why the same mathematical forms recur across domains that share no physical mechanism.

This paper states and tests a domain-independent answer: under the stated axioms, the form of a scaling law is determined by the composition operator of the underlying recursive process, and Cauchy's functional equations constrain which forms are possible.

2. The Theory

2.1 Cauchy's four equations

Augustin-Louis Cauchy established in the early nineteenth century that four functional equations have unique continuous solutions:

Additive: $f(x + y) = f(x) + f(y)$ has solution $f(x) = cx$

Multiplicative: $f(xy) = f(x) \cdot f(y)$ has solution $f(x) = x^c$

Exponential: $f(x + y) = f(x) \cdot f(y)$ has solution $f(x) = a^x$

Logarithmic: $f(xy) = f(x) + f(y)$ has solution $f(x) = c \log x$

These are not four independent results. They are four faces of one constraint: the requirement that a function be compatible with a binary operation. Under regularity (continuity, or measurability, or boundedness on an interval), a function that carries one of the two structures on the positive reals into one of the two structures falls into exactly four families: linear ($f(x) = cx$), power law ($f(x) = x^c$), exponential ($f(x) = a^x$) and logarithmic ($f(x) = c \log x$). This exhausts the homomorphisms. It does not exhaust scaling laws, because most systems are not homomorphisms of either structure: saturating behaviour, critical exponents, curvature on log axes and multi-regime crossovers all lie outside the grid, and the grid says nothing about them. The additive equation yields the linear form $f(x) = cx$; the multiplicative equation yields power laws; the exponential equation governs exponential growth and decay; and the logarithmic equation governs the unbounded map $c \log x$, which is multiplicative in input and additive in output and is not a diminishing-returns family, since $c \log x$ is unbounded.

2.2 The prediction

Consider any system where a quantity U grows with recursive depth R . The system's *composition operator* determines how consecutive applications combine:

- **Multiplicative composition** ($U(R_1 \cdot R_2) \sim U(R_1) \cdot U(R_2)$: depth multiplies, output multiplies): Cauchy's multiplicative equation forces $U \propto R^c$ - a power law.
- **Additive composition** ($U(R_1 + R_2) \sim U(R_1) + U(R_2)$: depth adds, output multiplies): Cauchy's exponential equation forces $U \propto a^R$ - an exponential.
- **Bounded composition** (output constrained to $[0, K]$): the compatible forms are saturation curves (logistic, Hill, Michaelis-Menten, hyperbolic). A formal proof of exhaustiveness for the bounded case is not provided here and remains an open problem.

THE PREDICTION

If you know the composition operator of a system, you can predict the functional form of its scaling law before seeing any data. The prediction follows from Cauchy's equations with no free parameters and no curve fitting.

3. Prediction Protocol

For each domain:

1. **Classify the composition operator from known physics** (before fitting). Determine whether the system's recursive amplification process combines multiplicatively, additively, or with a bound.
2. **Predict the scaling family** from the operator class using Cauchy's equations.
3. **Load published or standard reference data.**
4. **Independently fit candidate models** (power law, exponential, logistic, Hill, Michaelis-Menten, saturation-exponential) to the data using nonlinear least-squares regression.
5. **Select the best model by AICc** (corrected Akaike information criterion) and compare the prediction against the selected model's family.

STRUCTURED PREDICTION COMPARISON - NOT A BLIND TEST

The operator classification is made before the fitting procedure runs. However, the predictions, data, and fitting logic all reside in a single author-written script, and no registry submission exists for this test. The predictions for this cohort are not timestamped in advance of their own fits: the classification rule reached the public repository at 23:33 UTC on 16 March 2026 and the fifty per-domain assignments were committed together with their results at 23:54 that night, attestation class public repository commit date for both, so predicted families before fits is true of the procedure and not of the timestamp. The advance per-domain timestamp in this record belongs to the twelve-domain extension, whose manifest was committed at 00:19:19 UTC on 17 March 2026, twenty-three minutes before its fits, and to the book's two prediction appendices in print on 2 January 2026, before the papers and the programme existed. What was not fixed in advance here was the fitting protocol, and that gap is closed prospectively by the staged tests, the fourth-cell test and the 80-domain expansion test, both drafted and awaiting human submission. This test should therefore be understood as a structured prediction comparison, not a blind trial. A replication with independent operator classification has since been drafted, as an externally classified confirmatory extension of this paper, and it awaits human submission. The replication and the fourth-cell test are preregistered by timestamp, each first anchored on 12 August 2026 at Bitcoin block 962066.

3.1 Statistical method

The v2 test uses proper AICc (corrected Akaike information criterion) computed from residual sum of squares, sample size, and parameter count - not the R^2 -based approximation used in the legacy v1 test. The v1 approximation ($AIC \approx n \ln(1 - R^2) + 2k$) was adequate as a heuristic but should not have been described as strict AIC in the standard statistical sense. The v2 rewrite corrects this.

The primary significance test is a marginal-conditioned permutation test: holding both the predicted and the observed family marginals fixed and permuting the pairing 400,000 times, the probability of 19 or more matches is $p = 7 \times 10^{-5}$. The expected match count under this null is 9.5 of 25, not the 8.33 a uniform null gives, because power laws are simply common among the observed fits. A one-sided binomial test against a uniform three-family null gives $p = 1.56 \times 10^{-5}$ and is retained as a sensitivity check; it is the weaker null and is not used as the headline. Sigma equivalences are not quoted because they shift with the choice of null. Reproduction: scripts/paper-vii-permutation-null.py in the site repository, which parses the published Section 5 table directly so it cannot drift from this page.

3.2 A model-set limitation this paper must state against itself

Section 2.1 names four Cauchy families. The candidate model set in Section 3 contains six models spanning three of them: the power-law, exponential and bounded families. Two of the four continuous Cauchy solutions have no candidate of their own: the logarithmic form $c \log x$ and the pure linear form cx . One consequence runs each way. A genuinely logarithmic domain could

not have been classified correctly, because the fitter had no logarithmic candidate to select, so some of the six misses may have been guaranteed misses; and some of the nineteen hits may be logarithmic relations that a power law fits acceptably over a finite range. Every hit count in this paper therefore carries this qualifier, and none may be quoted without it. A confirmatory study that re-runs the comparison with the full four-family candidate set, with per-cell outcomes specified in advance, has been drafted and dated as a draft registration awaiting human submission (the four-family logarithmic-cell test in the programme's list of designs, alongside the fresh-domain extension study). Until that re-run reports, the numbers below are the six-model result, exactly as computed, and nothing stronger.




Figure 1 | ARC Architecture and Cauchy Unification. Four functional equations (Cauchy 1821) constrain scaling to four families: linear, power law, exponential and logarithmic. $\alpha=d/(d+1)$ for physical systems (13/13 nearest-comparator matches at the taxon grain, reported by tier and never blended; the largest modern mammalian dataset gives 0.687, 95% CI 0.674 to 0.701, which excludes 0.750, see §7.1); $\alpha=1/(1-\beta_L)$ for recursive intelligence (the self-reference theorem; in the artwork β is drawn bare, carrying no subscript), one of the two routes to the value two; the second is the ARC Ceiling $\alpha_{crit}=1/(1-\gamma)$, and whether the pair agree away from $\beta_L=\gamma=0.5$ is a drafted test, preregistered by timestamp and awaiting human submission. 19/25 domains confirm predicted function family ($p=7\times10^{-5}$, permutation; six-model set, no logarithmic candidate - §3.2). The rerun of 11 August 2026, which added the logarithmic candidate, returns 18/25. Source: Paper VII · Foundational · evidence spine C-4, C-11.

4. Evidence Tiers

The 50-domain suite comprises five evidence tiers of decreasing strength. These tiers are never combined into a single blended number.

TIER	DESCRIPTION	COUNT	RESULT	SIGNIFICANCE
PRIMARY	Empirical curve-fit (AICc)	25	19/25 (18/25 in the rerun of 11 August 2026 that added the logarithmic candidate)	$p = 7 \times 10^{-5}$ perm.; 1.56×10^{-5} binomial

TIER	DESCRIPTION	COUNT	RESULT	SIGNIFICANCE
SECONDARY	Baseline-20 rerun (AICc)	20		15/20 $p = 1.67 \times 10^{-4}$
SECONDARY	Published exponents (direct)	13		13/13 Nearest comparator match
NOTED	Published exponents (provisional)	6		3/6 Contested organisms
NOTED	Analytic identities	6		6/6 Tautological confirmation

5. Primary Result: 25 Empirical Curve-Fit Domains

#	DOMAIN	OPERATOR	PREDICTED	AICC BEST	MATCH
1	Kleiber's Law (Metabolic Scaling)	Multiplicative	Power law	Saturation-exp	No
2	Urban Scaling (GDP vs Population)	Multiplicative	Power law	Power law	Yes
3	Species-Area (Galapagos)	Multiplicative	Power law	Hill	No
4	Wright's Law (Solar PV)	Multiplicative	Power law	Power law	Yes
5	Heap's Law (Vocabulary)	Multiplicative	Power law	Power law	Yes
6	Zipf's Law (Word Frequency)	Multiplicative	Power law	Michaelis-Menten	No
7	Learning Curve (Cigar Rolling)	Multiplicative	Power law	Power law	Yes
8	Moore's Law (Transistor Count)	Additive	Exponential	Exponential	Yes
9	Radioactive Decay (P-32)	Additive	Exponential	Exponential	Yes
10	Gutenberg-Richter (Earthquakes)	Additive	Exponential	Exponential	Yes
11	Bacterial Growth (E. coli)	Bounded	Bounded	Logistic	Yes
12	O ₂ -Hemoglobin Curve	Bounded	Bounded	Hill	Yes
13	Epidemic SIR (Ebola 2014)	Bounded	Bounded	Hill	Yes
14	Amdahl's Law (CPU Scaling)	Bounded	Bounded	Michaelis-Menten	Yes
15	Muscle Force-Velocity (Hill 1938)	Bounded	Bounded	Exponential	No
16	Facebook MAU Growth	Bounded	Bounded	Logistic	Yes
17	Brownian Diffusion (MSD)	Multiplicative	Power law	Power law	Yes
18	Horton's Law (Stream Numbers)	Additive	Exponential	Exponential	Yes
19	Neural Scaling Laws (LLM Loss)	Multiplicative	Power law	Power law	Yes
20	Time Crystal Order (Rydberg Gas)	Bounded	Bounded	Power law	No
21	Stellar Mass-Luminosity	Multiplicative	Power law	Logistic	No
22	Heart Rate vs Body Mass	Multiplicative	Power law	Power law	Yes
23	Rent's Rule (VLSI Pin Count)	Multiplicative	Power law	Power law	Yes
24	Taylor's Power Law	Multiplicative	Power law	Power law	Yes

#	DOMAIN	OPERATOR	PREDICTED	AICC BEST	MATCH
25	Hack's Law (Stream Length)	Multiplicative	Power law	Power law	Yes

PRIMARY RESULT, ON THE SIX-MODEL SET (NO LOGARITHMIC OR PURE LINEAR CANDIDATE; SECTION 3.2): 19/25 EMPIRICAL DOMAINS CONFIRMED ($p = 7 \times 10^{-5}$, PERMUTATION)

Under strict AICc-based model selection, the Cauchy-predicted family is confirmed in 19 of 25 empirical curve-fit domains. (The rerun of 11 August 2026, which added the logarithmic candidate, returns 18 of 25.) The marginal-conditioned permutation probability is $p = 7 \times 10^{-5}$; the binomial probability under uniform random assignment to the three families present in this candidate set is $p = 1.56 \times 10^{-5}$, roughly 1 in 64,000; under the four-family grid the same result gives $p = 1.27 \times 10^{-7}$, roughly 1 in 7.9 million, so the figure reported here is the conservative one. On the strictly Cauchy-derived subset alone (multiplicative and additive domains only) the result is 14/18 at permutation $p = 3.0 \times 10^{-4}$: the headline does not depend on the bounded family, whose exhaustiveness remains open.

This is strong enough to publish as a structured prediction comparison under a publicly stated rule rather than a per-domain advance timestamp. It is not strong enough to claim as proven. Independent replication with pre-registered operator classification is the necessary next step.

5.1 Baseline-20 rerun

As context, the original 20 baseline domains were rerun with the new proper AICc fitter. The result is 15/20 ($p = 1.67 \times 10^{-4}$). This is not directly comparable to the v1 result of 14/20, because the fitter changed (R^2 -based AIC approximation to proper AICc) and some model selections shifted accordingly. The baseline-20 rerun is reported for transparency, not as evidence that the new fitter is more or less conservative than the old one.

6. The Six Empirical Misses

Six of 25 empirical domains produced a best-fit model outside the predicted family. Each miss has a specific, identifiable explanation. Honest analysis follows.

#	DOMAIN	PREDICTED	AICC BEST	EXPLANATION
1	Kleiber's Law	Power law	Saturation-exp	Classic contested case. Kleiber's original 1932 dataset has only 13 mammals. A bounded model can capture slight upper curvature at extreme body masses. The power-law exponent debate (0.67 vs 0.75) has persisted for decades.
3	Species-Area (Galapagos)	Power law	Hill	Small dataset (30 Galapagos islands) with high scatter. Several islands have very small areas with disproportionate species counts. The Hill function can absorb this scatter with its extra parameter.

#	DOMAIN	PREDICTED	AICC BEST	EXPLANATION
6	Zipf's Law	Power law	Michaelis-Menten	Finite-corpus truncation. The Brown Corpus has a hard upper bound on word frequency (69,971 for 'the'). At low ranks, frequencies saturate against this ceiling, pulling AICc towards a bounded model. Zipf's law is a power law in the asymptotic limit; finite corpora violate this limit.
15	Muscle Force-Velocity	Bounded	Exponential	The Hill 1938 force-velocity data shows a monotonic decrease from maximum isometric force to zero at maximum velocity. The data does not reach or demonstrate an asymptotic plateau - it terminates where force reaches zero. An exponential decay fits this truncated range better than a saturation curve. The bounded prediction requires data spanning the full approach to saturation.
20	Time Crystal Order	Bounded	Power law	Only 4 data points. With so few observations, AICc penalises the extra parameters of saturation models. A power law, with fewer free parameters, wins on parsimony. The bounded prediction may be correct but is untestable with this dataset.
21	Stellar Mass-Luminosity	Power law	Logistic	The mass-luminosity relation shows significant scatter at high stellar masses where radiation pressure, convective instability, and different fusion pathways create deviations from a simple power law. A logistic model absorbs this scatter by fitting an apparent upper rollover.

In summary: three misses involve small or truncated datasets (species-area, time crystal, Zipf); one involves data that does not span the full saturation range (muscle force-velocity); one involves a classically contested relationship (Kleiber); and one involves astrophysical complexity at extreme values (stellar mass-luminosity). None represents a case where the predicted family is fundamentally wrong in principle - but neither can any miss be dismissed. These are genuine failures that an independent replication must address.

7. Secondary Tier: Published Exponents

7.1 Direct published exponents (13/13)

Thirteen taxa with well-established metabolic scaling exponents were tested against the ARC/Cauchy dimensional prediction. The prediction requires three conditions: (a) multiplicative composition (Cauchy's equation constrains the form to a power law), (b) d -dimensional space-filling transport geometry, and (c) a conservation or optimisation constraint on resource flow (energy minimisation in West et al.; supply-demand balance in Banavar et al.; steady-state energy balance in Demetrius). Under all three conditions, the exponent is constrained to $d/(d + 1)$. Neither Cauchy alone, nor space-filling alone, is sufficient. Organisms with three-dimensional vascular transport and conserved flow should scale as $M^{3/4}$, those with two-dimensional transport as $M^{2/3}$, and those with one-dimensional filamentous transport as $M^{1/2}$.

NOTE ON THE 3/4 EXPONENT

The empirical value of the mammalian metabolic scaling exponent is debated, with estimates ranging from approximately 0.67 to 0.75 depending on taxon, mass range, temperature correction, and statistical method (White & Seymour 2003; Glazier 2005, 2022). The largest modern dataset (White and Seymour 2003, compiling 619 species) gives a maximum-likelihood estimate of 0.687 (95% CI 0.674-0.701), which excludes 0.750. The $d/(d+1)$ prediction of 0.750 for $d = 3$ matches the upper end of the empirical range. The variation itself is consistent with the framework: organisms with effective transport dimensions between 2 and 3 would produce exponents between 2/3 and 3/4. The framework predicts this variation rather than a single universal exponent.

#	TAXON	PREDICTED	OBSERVED	NEAREST	CI INCLUDES?
26	Mammals	0.750	0.737	$d = 3$	Yes
27	Birds	0.750	0.720	$d = 3$	Yes
28	Fish	0.750	0.800	$d = 3$	Yes
29	Reptiles	0.750	0.760	$d = 3$	Yes
30	Insects	0.750	0.750	$d = 3$	Yes
31	Amphibians	0.750	0.740	$d = 3$	Yes
32	Crustaceans	0.750	0.730	$d = 3$	Yes
33	Jellyfish	0.667	0.680	$d = 2$	Yes
34	Cnidarians	0.667	0.700	$d = 2$	Yes
35	Ctenophores	0.667	0.660	$d = 2$	Yes
36	Ectomycorrhizal fungi	0.500	0.580	$d = 1$	Yes
37	Marine fungi	0.500	0.530	$d = 1$	Yes
38	Saprotrophic fungi (20°C)	0.500	0.530	$d = 1$	Yes

The compendium value at row 26 is 0.737, and its interval admits 0.750, whereas the 619-species estimate given in the note above rules that value out; each figure is stated here, and per-row sourcing for the table is still outstanding (see the provenance note below). Every one of the 13 published exponents sits closer to the predicted dimensional comparator ($d/(d+1)$) than to any rival, and every one of the 13 confidence intervals contains the predicted value. This is corroborating evidence whose per-row sourcing is outstanding, but it tests a different claim (specific exponent) than the primary result (functional family), so it is reported separately.

PROVENANCE NOTE ON THE 13-TAXA TABLE

The table above does not carry a per-row citation column. Section 10.2 references [26]-[35] name the broader compendium literature on interspecific metabolic scaling (West, Brown and Enquist 1997; Banavar et al. 1999, 2010; Demetrius 2003, 2006, 2010; White and Seymour 2003; Glazier 2005, 2022; Kolokotronis et al. 2010) from which the mammalian, vertebrate and invertebrate scaling exponent literature is compiled, but this paper does not currently anchor a locked per-row source citation to each observed exponent or confidence interval in the table. Independent readers should treat individual rows as pending a per-row source audit before reuse in secondary work; the jellyfish, cnidarian and ctenophore rows (33-35) and the three fungal rows (36-38) are the ones for which a per-row source citation has not been anchored in this paper.

CAVEAT ON $d=2$ BIOLOGICAL ENTRIES

The jellyfish, cnidarian, and ctenophore entries (rows 33-35) are classified as nearest-comparator matches to $d = 2$, meaning their observed exponents fall closer to $2/3$ than to either $1/2$ or $3/4$. However, the $d = 2$ biological prediction remains untested in the strong sense because no known organism possesses a genuinely two-dimensional hierarchical space-filling transport network with conserved flow. These organisms lack the vascular architecture that the $d/(d+1)$ derivation assumes. The $d = 2$ confirmation exists in cosmology (Friedmann matter-era solution, exact) and physics (percolation, fragmentation), not in biology. These rows should be understood as empirical proximity to the $d = 2$ comparator, not as dimensional confirmation.

7.2 Provisional published exponents (3/6)

Six additional taxa with contested or proxy-based dimensional assignments were tested. Only 3 of 6 match the nearest predicted comparator. The misses are flatworms (disputed 2D vs 3D transport), bryozoans (colonial organisms with anomalous scaling near 1.0), and glass eels (elongated but retaining 3D vascular transport). These are noted for completeness but carry low evidential weight due to the contested dimensional assignments.

7.3 Analytic identities (6/6)

Six domains defined by exact analytic formulae ($E = mc^2$, hydrogen energy levels, Arrhenius kinetics, Michaelis-Menten kinetics, Friedmann matter-era expansion, de Sitter dark-energy expansion) were included as sanity checks. All 6 confirm. These are tautological - fitting a curve to data generated by an exact formula must recover that formula - and carry no independent evidential weight. They confirm only that the fitter works correctly.

7.4 The March 2026 preregistration folder, and the follow-on runs of the same era (recorded 25 August 2026)

A complete preregistration folder for this paper's extension existed in March 2026 and is preserved in the public repository (experiments/cauchy-unification__Paper-VII/preregistration/): a written protocol; a candidate manifest committed at 00:19:19 UTC on 17 March 2026, whose own status field at that commit read `draft_local_not_locked`, carrying twelve rows each with operator class, predicted family, predicted model and target source fixed before any fit; the author's own file checksums; and the OSF component registration text, already written. The only absent step was the registry submission click. The folder's own status field records the limitation in the estate's words: `'archived_local_packet_exercised_before_timestamp'`, and the analysis is pinned to a git commit. Preregistration existed here in substance: the manifest was committed to the public repository at 00:19:19 UTC on 17 March 2026, twenty-three minutes before the fits ran, and the same files were deposited to the public OSF component between 02:06:51 and 02:07:11 that night, where they remain unaltered. What is absent is an accepted registry registration, a separate schema-bound object this component does not carry; the API reports `registration: false`. The folder's status field read `draft_local_not_locked` at the commit that fixed the predictions and was rewritten to `archived_local_packet_exercised_before_timestamp` at 00:56, after the run.

Two further March 2026 runs existed only as repository artefacts until this version, and both are recorded here with their status stated before their numbers. Both ran under the pre-correction candidate set (no logarithmic candidate), and the extension's design lineage was later retired unedited, so both are dated exploratory results under the superseded taxonomy, superseded by the fourth-cell test, which is drafted and awaiting human submission, and neither is confirmation of anything.

The 12-domain extension of 17 March 2026 returned 10/12 family matches and 8/12 exact model matches (binomial against one third, $p = 5.4 \times 10^{-4}$; `results_preregistered_extension.json`), its predictions committed twenty-three minutes before the fits. Its rung: author-classified rather than externally classified; pre-correction candidate set; public repository commit attestation, with the outcome file not among those deposited to the component; and one row, EXT-10 receptor occupancy, run on Hill's 1910 haemoglobin data, repeats a domain, an operator class and a predicted model already public at 23:54 the previous evening with a matching outcome, so the strictly fresh set reads 9 of 11 ($p = 1.4 \times 10^{-3}$), or 8 of 10 ($p = 3.4 \times 10^{-3}$) if enzyme kinetics is removed with it. It is a dated result at that rung and not confirmation of this paper's claim. The demotion of that run's status is dated to the minute in the public repository's commit history: `e5c22dd` records the 10/12 result at 00:43 on 17 March 2026, `f5e06a8` adopts the conservative miss posture at 00:47, and `c853277` downgrades the run to a pilot dry run at 00:56, thirteen minutes after the result was recorded and months before any external review existed. The result was never retracted; its claimed status was corrected, by the author, against interest, in public. That demotion is kept here exactly as it was written. It was reversed on 7 September 2026 under the estate's attestation-class rule: a prediction whose text is public and timestamped before its outcome is not demoted for the absence of a registry form, and the commit of 00:19 had already fixed these predictions before the run. The reversal changes the status label and none of the numbers. The 7-domain temporal out-of-sample run, which fits candidate families on old data and tests whether the winning family predicts the functional form of new data, returned 4 confirmed, 2 partial and 1 in-

consistent (results_temporal_out_of_sample.json); the inconsistent domain is reported with the same prominence as the confirmations. A 105-domain expansion catalogue (rows 26 to 65 and 66 to 105) is prepared in the repository data and has not been run; no result beyond the runs named here is claimed.

The outermost ring of the same dated chain is in print. The first edition of 2 January 2026 (ISBN 978-1806056200) carries Appendix F, five numbered testable predictions closing with a printed wager: “These predictions are my wager. If they fail, the framework is wrong or incomplete. If they succeed, something important has been glimpsed. Time will judge.” Appendix A operationalises the framework’s propositions, states four further predictions in §A.3 (AI development quadratic rather than linear; consciousness and integrated information; cosmological recursive error correction; value embedding), and fixes falsification criteria in advance in §A.4. Prediction 2 prints the 15 per cent and 5 per cent alignment-drift thresholds that the prepared study-v draft registration now carries unchanged into a formal hypothesis, citing the appendix by name. The book is a printed prior statement rather than a preregistration, and the two are never conflated: the print fixes the prediction; only a registration fixes an analysis plan.

8. Significance Framing

The primary result ($p = 7 \times 10^{-5}$, permutation) deserves honest contextualisation.

- **In absolute terms:** under the uniform-chance binomial null, roughly 1 in 64,000: if 64,000 researchers each picked scaling families at random for 25 domains, on average one would match 19 or more by chance. The stricter permutation null, which concedes that power laws are simply common among real fits, still leaves $p = 7 \times 10^{-5}$.
- **Relative to convention:** both tests land three orders of magnitude beyond $p < 0.05$. Sigma equivalences are not quoted because they shift with the choice of null.
- **What it supports:** the result is strong enough to publish, as a structured prediction comparison under a publicly stated rule rather than a per-domain advance timestamp, as evidence that Cauchy-type composition constraints predict scaling families at well-above-chance rates.
- **What it does not support:** it is not strong enough to claim the theory is proven. The particle-physics discovery threshold is 5σ ($p \approx 3 \times 10^{-7}$). The test lacks formal pre-registration. The operator classifications were made by the author, not an independent panel. The candidate set contained no logarithmic and no pure linear model, so the hit rate is conditional on that model set (Section 3.2).
- **What is needed:** independent replication with pre-registered operator classification, ideally by researchers with no stake in the outcome.

9. Group Analysis

9.1 Multiplicative composition: power laws

Of 14 empirical domains classified as multiplicative, 10 produce power-law scaling as their AICc-best model. These span metabolic allometry (heart rate vs body mass), urban economics (Bettencourt 2007), technology learning curves (Wright's law), computational linguistics (Heap's law), skill acquisition (Crossman 1959), Brownian motion (Catipovic 2013), neural network scaling (Kaplan 2020), VLSI design (Rent's rule), population ecology (Taylor's power law), and fluvial geomorphology (Hack's law).

Four multiplicative domains miss: Kleiber's law (bounded model edges out power law on a 13-point dataset), species-area (high scatter on small island data), stellar mass-luminosity (scatter at high mass), and Zipf's law (finite-corpus truncation).

9.2 Additive composition: exponentials

All four domains with additive composition operators produce exponential scaling as their AICc-best model: Moore's law (transistor counts), radioactive decay (P-32), the Gutenberg-Richter earthquake frequency law, and Horton's law (stream numbers). This is a clean 4/4.

9.3 Bounded composition: saturation curves

Of seven domains with bounded composition operators, five produce bounded (saturation) models as AICc-best: bacterial growth (logistic), hemoglobin binding (Hill), Ebola epidemic (Hill), Amdahl's law (Michaelis-Menten), and Facebook growth (logistic). Two miss: muscle force-velocity (exponential decay fits the truncated range better) and time crystal order (only 4 data points).

10. Data Provenance

10.1 Provenance caveat

Three domains use theoretical or reference values rather than independent empirical measurements: radioactive decay (NIST reference half-life for P-32), Amdahl's law (theoretical formula with assumed 10% serial fraction), and Moore's law (Wikipedia-sourced transistor counts from manufacturer specifications). These confirm the framework trivially and should be weighted accordingly.

10.2 Citations for all 25 empirical domains

- [1] Kleiber, M. (1932). Body size and metabolism. *Hilgardia* 6:315-353. Reproduced in Dodds, P. et al. (2001) *J. Theor. Biol.* 209:9-27.
- [2] Bettencourt, L. et al. (2007). Growth, innovation, scaling, and the pace of life in cities. *PNAS* 104(17):7301-7306. Data: BEA + Census 2006, 25 largest US MSAs.
- [3] Johnson, M. & Raven, P. (1973). Species number and endemism: the Galapagos archipelago revisited. *Science* 179:893-895. Data: R *faraway::gala* dataset, 30 islands.
- [4] Nemet, G. (2009). Demand-pull, technology-push, and government-led incentives for non-incremental technical change. *Research Policy* 38(5):700-709; Farmer, J.D. & Lafond, F. (2016) *Research Policy*; IRENA; Our World in Data.

- [5] Manning, C. et al. (2008). *Introduction to Information Retrieval*. Reuters RCV1 corpus.
- [6] Kucera, H. & Francis, W. (1967). *Computational Analysis of Present-Day American English* (no official online source is available, 3 October 2026). Brown Corpus.
- [7] Crossman, E. (1959). A theory of the acquisition of speed-skill. *Ergonomics* 2(2):153-166. Reanalysis: Newell, A. & Rosenbloom, P. (1981).
- [8] Wikipedia: Transistor count. Public vendor specifications (Intel, AMD, TSMC).
- [9] NIST Nuclear Data Center. P-32 half-life: 14.29 days. Exact theoretical values.
- [10] USGS Earthquake Hazards Program; IRIS compilation. Global seismicity catalogue.
- [11] Sezonov, G. et al. (2007). *E. coli* physiology in Luria-Bertani broth. *J. Bacteriol.* 189:8746-8749.
- [12] Severinghaus, J. (1979). Simple, accurate equations for human blood O₂ dissociation computations. *J. Appl. Physiol.* 46(3):599-602.
- [13] WHO/CDC Ebola Situation Reports (2014-2015); NEJM Ebola Response Team.
- [14] Amdahl, G. (1967). Validity of the single processor approach to achieving large scale computing capabilities. *AFIPS*. See also Hill, M. & Marty, M. (2008) *IEEE Computer*.
- [15] Hill, A. (1938). The heat of shortening and the dynamic constants of muscle. *Proc. Royal Soc. B* 126:136-195.
- [16] Meta SEC filings; Zhang, X.-Z., Liu, J.-J. & Xu, Z.-W. (2015). Tencent and Facebook Data Validate Metcalfe's Law. *J. Comp. Sci. Tech.* 30(2):246-251.
- [17] Catipovic, M. et al. (2013). Improving the quantification of Brownian motion. *Am. J. Physics* 81:485.
- [18] Singh, W. R. et al. (2021). Morphometric analysis and watershed prioritization in relation to soil erosion in Dudhna Watershed. *Applied Water Science* 11:151.
- [19] Kaplan, J. et al. (2020). Scaling Laws for Neural Language Models. arXiv:2001.08361.
- [20] Shen, Y. et al. (2025). Time crystal order in a Rydberg gas. *Nature Communications*. See also Kongkhambut, P. et al. (2024) *Science* 386(6720):670-674.
- [21] Torres, G. et al. (2010). Accurate masses and radii of normal stars. *Astron. Astrophys. Rev.* 18:67-126. Eker, Z. et al. (2018) *MNRAS* 479:5491. IAU 2015 nominal solar values.
- [22] Stahl, W. (1967). Scaling of respiratory variables in mammals. *J. Appl. Physiol.* 22:453. Calder, W. (1984) *Size, Function, and Life History* (no official online source is available, 3 October 2026). Schmidt-Nielsen, K. (1984) *Scaling: Why Is Animal Size So Important?*
- [23] Landman, B. & Russo, R. (1971). On a pin versus block relationship for partitions of logic graphs. *IEEE Trans. Comp.* C-20:1469. Christie, P. (2000); Stroobandt, D. (2001).
- [24] Taylor, L. (1961). Aggregation, variance and the mean. *Nature* 189:732-735. Taylor, L. & Woiwod, I. (1980) *J. Anim. Ecol.* 49:879.
- [25] Hack, J. (1957). Studies of longitudinal stream profiles in Virginia and Maryland. *USGS Professional Paper* 294-B. Mueller, J. (1973); Rigon, R. et al. (1996) *Water Resour. Res.* 32:3367.
- [26] West, G. B., Brown, J. H. & Enquist, B. J. (1997). A general model for the origin of allometric scaling laws in biology. *Science* 276(5309):122-126.
- [27] Banavar, J. R., Maritan, A. & Rinaldo, A. (1999). Size and form in efficient transportation networks. *Nature* 399(6732):130-132.
- [28] Banavar, J. R., Moses, M. E., Brown, J. H., Damuth, J., Rinaldo, A., Sibly, R. M. & Maritan, A. (2010). A general basis for quarter-power scaling in animals. *Proceedings of the National Academy of Sciences* 107(36):15816-15820.
- [29] Demetrius, L. (2003). Quantum statistics and allometric scaling of organisms. *Physica A* 322:477-490.
- [29b] Demetrius, L. (2006). The origin of allometric scaling laws in biology. *Journal of Theoretical Biology* 243(4):455-467.
- [29c] Demetrius, L. & Tuszyński, J. A. (2010). Quantum metabolism explains the allometric scaling of metabolic rates. *Journal of the Royal Society Interface* 7(44):507-514.

- [30] Zhao, J. (2022). Universal growth scaling law determined by dimensionality. *arXiv:2206.01850*.
- [31] Bettencourt, L. M. A. (2013). The origins of scaling in cities. *Science* 340(6139):1438-1441.
- [32] White, C. R. & Seymour, R. S. (2003). Mammalian basal metabolic rate is proportional to body mass^{2/3}. *Proceedings of the National Academy of Sciences* 100(7):4046-4049.
- [33] Glazier, D. S. (2005). Beyond the '3/4-power law': variation in the intra- and interspecific scaling of metabolic rate in animals. *Biological Reviews* 80(4):611-662.
- [34] Glazier, D. S. (2022). Variable metabolic scaling breaks the law: from 'Newtonian' to 'Darwinian' approaches. *Proceedings of the Royal Society B* 289(1985):20221605.
- [35] Kolokotronis, T., Savage, V., Deeds, E. J. & Fontana, W. (2010). Curvature in metabolic scaling. *Nature* 464(7289):753-756.
- [36] Cauchy, A.-L. (1821). Cours d'analyse de l'École royale polytechnique. 1re partie. Analyse algébrique.
- [37] Bingham, N. H. & Ostaszewski, A. J. (2015). Cauchy's functional equation and extensions: Goldie's equation and inequality, the Gołąb-Schinzel equation and Beurling's equation. *Aequationes mathematicae* 89(5):1293-1310.
- [38] He, J.-H. & Chen, H. (2003). Effects of size and pH on metabolic rate. *International Journal of Nonlinear Sciences and Numerical Simulation* 4(4). The record at Crossref carries no abstract, and on 3 October 2026 it could not be confirmed that it is the work Section 11 describes.
- [39] He, J.-H. & Zhang, J. (2004). Fifth dimension of life and the 4/5 allometric scaling law for human brain. *Cell Biology International* 28(11):809-815.
- [40] Maino, J. L., Kearney, M. R., Nisbet, R. M. & Kooijman, S. A. L. M. (2014). Reconciling theories for metabolic scaling. *Journal of Animal Ecology* 83(1):20-29.
- [41] West, G. (2017). Scale: The Universal Laws of Life, Growth, and Death in Organisms, Cities, and Companies. Penguin Books.
- [42] Aczél, J. (1966). Lectures on Functional Equations and Their Applications. *Mathematics in Science and Engineering* 19. Academic Press.

11. Related Work

A literature search returned no previous publication that invokes all four Cauchy functional equations as a unifying cross-domain principle constraining the forms of scaling laws. Readers are invited to submit prior work; a genuine antecedent that does this defeats the priority claim of this paper.

The closest existing work includes:

- **Bingham & Ostaszewski (2015):** connect Cauchy's equation to regular variation theory. Pure mathematics. No cross-domain physical science application.
- **West, Brown & Enquist (1997):** derive a 3/4 metabolic exponent from fractal vascular networks with energy-minimised flow. Domain-specific. Does not predict other domains. Note that the empirical exponent is debated (see Warning box in Section 7.1).
- **Banavar, Maritan & Rinaldo (1999, 2010):** derive the same $d/(d + 1)$ form from geometric constraints on transportation networks combined with supply-demand conservation, without requiring fractal geometry. Independently confirmed the quarter-power exponent from different starting assumptions. Banavar et al. (2010) showed that the exponent ranges from 2/3 to 3/4 depending on how velocity scales, with 3/4 as the optimised maximum - demonstrating that space-filling geometry alone is insufficient without a conservation or optimisation constraint.

- **Demetrius (2003, 2006) and Demetrius & Tuszynski (2010):** derive equivalent scaling from quantum statistical mechanics, applying the Debye model of thermal properties to coupled energy-transducing oscillator networks in d -dimensional space. A completely independent route to the same exponent from different physics (quantum oscillator coupling rather than network geometry). The convergence of this derivation with West's network geometry derivation is explained by Cauchy's multiplicative functional equation applied to multiplicative composition in d -dimensional space with a conservation constraint (see companion paper, On the Origin of Scaling Laws).
- **Bettencourt (2013):** derives city scaling exponents from network dimensionality, extending the $d/(d + 1)$ form beyond biology to urban systems.
- **Zhao (2022):** derives the formula as a universal growth scaling law from producer-consumer dynamics, providing another independent derivation.
- **He & Chen (2003):** derive metabolic scaling exponents from optimal vascular network geometry with flow conservation constraints. The one He and Chen record of 2003 at Crossref is Ji-Huan He and Hao Chen, Effects of size and pH on metabolic rate, *International Journal of Nonlinear Sciences and Numerical Simulation* 4(4); on 3 October 2026 it could not be confirmed that this record is the work described here.
- **He & Zhang (2004):** state, of brain cells, that the fractal geometry of the cell “leads to a 4/5 allometric scaling law for human brain”; they give a physical explanation for the fourth dimension of life for three-dimensional organisms and the fifth dimension for the human brain.
- **Maino et al. (2014):** within dynamic energy budget theory, “derive an equation for the scaling of metabolic rate for adult heterotrophs without relying on optimisation arguments or implying cellular nutrient supply limitation”; their abstract states no exponent.
- **Geoffrey West, *Scale* (2017):** discusses scaling across biology and cities. No mention of Cauchy's functional equations.
- **Aczél (1966):** standard reference for all four Cauchy equations. Applications to information theory but not to cross-domain scaling.

The gap is genuine. The $d/(d + 1)$ formula has been independently derived by at least seven research groups through different mathematical frameworks - fractal networks, geometric constraints with supply-demand balance, quantum metabolism, urban scaling, producer-consumer dynamics, optimal vascular geometry, and thermodynamic first principles; of these, the He and Chen item could not be confirmed at a registry on 3 October 2026, and the abstract of Maino et al. (2014) names dynamic energy budget theory and states no exponent. Every known derivation requires three conditions: multiplicative composition (which Cauchy constrains to the power-law family), d -dimensional space-filling geometry, and a conservation or optimisation constraint on resource flow. The mathematical theorems are 200 years old. The biological data is 90 years old. The novel contribution is not the $d/(d + 1)$ formula itself (which is well established), nor Cauchy's mathematics, but the claim that Cauchy-type functional equations explain *why* these independent derivations converge on the same form, and that the Cauchy constraint has a physically testable consequence for scaling-law classification across domains.

12. Limitations

- **No registry submission for this test, and no advance per-domain timestamp for this cohort.** The predictions, data arrays, and fitting logic all reside in a single author-written script. The classification rule reached the public repository at 23:33 UTC on 16 March 2026 and the fifty per-domain assignments were committed together with their results at 23:54 that night, attestation class public repository commit date for both, so predicted families before fits is true of the procedure and not of the timestamp. The advance per-domain timestamp in this record belongs to the twelve-domain extension of 00:19:19 UTC on 17 March 2026 and to the book's two prediction appendices in print on 2 January 2026. What was not fixed in advance here was the fitting protocol, and that gap is closed prospectively by the staged tests, the fourth-cell test and the 80-domain expansion test, both drafted and awaiting human submission. This test should be understood as a structured prediction comparison, not a blind trial.
- **Operator classification by author.** The composition operator classification is done by the author, not by an independent assessor. A replication with independent operator classification would be more rigorous.
- **Bounded composition exhaustiveness is open.** The proof that only saturation-family curves are compatible with bounded composition is not provided and remains an open problem. The bounded category currently groups logistic, Hill, Michaelis-Menten, and hyperbolic forms by empirical convention rather than mathematical proof.
- **Legacy AIC was an approximation.** The v1 test used an R^2 -based AIC approximation ($AIC \approx n \ln(1 - R^2) + 2k$), not a full likelihood-based nonlinear AIC. The v2 test corrects this with proper AICc computed from residual sum of squares.
- **Three datasets use reference values.** Radioactive decay (NIST half-life), Amdahl's law (theoretical formula), and Moore's law (manufacturer specs) use theoretical or reference values rather than independently collected empirical measurements.
- **Family prediction only.** The framework predicts the functional family but not the specific exponent. Predicting exponents requires additional structure (dimensionality, constraint type) developed in the companion Foundational paper.
- **Six empirical misses.** The misses are not dismissible. They identify real limits of the framework: small datasets, truncated ranges, high-scatter regimes, and domains where competing physical effects (radiation pressure, corpus boundaries) perturb the ideal prediction.
- **Operator classification is descriptive, not explanatory.** The framework does not explain *why* a particular domain has a multiplicative rather than additive composition operator. It takes the operator as given and predicts the consequence.
- **Baseline-20 rerun is not cleanly comparable to v1.** The fitter changed from R^2 -based AIC approximation to proper AICc. The model set expanded. A direct comparison of 14/20 (v1) to 15/20 (v2 baseline rerun) is not meaningful because the selection criterion changed.
- **Publication bias risk.** The 5 new expansion domains were selected after the baseline-20 test. While they were chosen to span new physical regimes (stellar astrophysics, allometric heart

rate, VLSI, ecology, geomorphology), the selection was not randomised.

- **Finite model set, and specifically two Cauchy families out, one non-Cauchy family in.** The fitter compares six candidate models (power law, exponential, logistic, Hill, Michaelis-Menten, saturation-exponential). Two of the four Cauchy solutions named in Section 2.1 have no candidate (logarithmic $c \log x$; pure linear cx), while the bounded family, whose exhaustiveness remains an open problem, is represented four times. The true best model for any domain may not be in this set, and Section 3.2 states the consequences for the hit counts.
- **Sample size.** Twenty-five empirical domains is a reasonable first sample but not definitive. Expansion to 100+ domains with pre-registered selection criteria would strengthen or weaken the result.

12.1 Falsifiability: what would defeat this claim

The limitations above describe where the framework is weak; this subsection states the specific outcomes that would falsify it. The central claim, that a domain's composition operator, classified from known physics *before* fitting, predicts its scaling-law family, is defeated by any of the following:

1. **Hit rate at chance.** The primary result is a family-prediction hit rate of 19/25 empirical domains (§5 to §6), 18/25 in the rerun of 11 August 2026 that added the logarithmic candidate. With three to four candidate families, assignment at random gives a chance baseline of roughly one in three to one in four. The claim is defeated if a replication whose domain list and per-domain predictions are timestamped in advance, and whose operator classes are assigned by an external human classifier, returns a hit rate not significantly above that chance baseline. The advance-timestamped run on the record is the 12-domain extension of 17 March 2026, its predictions committed to the public repository at 00:19:19 UTC in commit 8edaec0, attestation class public repository commit date, its fits run twenty-three minutes later: 10 of 12, and 9 of 11 on the strictly fresh rows. It is preregistered in substance, attestation class public repository commit date, and it is not an accepted registry submission; its twelve classifications are the author's own, so this criterion has not yet been put at risk in the form stated here.
2. **Post-hoc operator classification.** The test's strength depends on the operator being assignable from independently-known physics without seeing the data. The claim is defeated if, in a blind trial, independent assessors cannot reproduce the author's operator classifications from the physics alone, or if any classification can be shown to have been chosen to fit the observed exponent.
3. **Fit degeneracy.** The prediction is meaningful only if the candidate families are empirically distinguishable on the data at hand. It is defeated for any domain where the fitter reaches the observed curve equally well under the *wrong* family; that is, the family assignment carries no discriminating information for that domain's range and noise.
4. **Unrationalised misses.** Six domains miss (§6). The misses are consistent with the claim only if their causes are fixed as a pre-registered exclusion rule *before* the next test, not supplied after seeing which domains missed. The claim is defeated if the miss rate rises under a pre-registered replication with no post-hoc exclusions.

5. **Four-family re-run.** The drafted four-family confirmatory study (Section 3.2) re-runs this comparison with logarithmic and pure linear candidates included. The claim is weakened if the hit rate falls materially when the missing families are available, and defeated if it falls to the chance baseline; hits that migrate to the logarithmic family will be reported as reclassifications, not as new confirmations.
6. **Exploratory ceiling.** Until the operator classes and the per-domain predictions for a randomised domain sample are fixed before any fitting by an external human classifier and timestamped in advance with the attestation class stated, the claim cannot rise above *exploratory*, and any presentation of it as established would itself falsify the paper's stated evidential status. A registry deposit raises the attestation class; it is not what lifts the ceiling. Part of that condition is already met on the record: the 12-domain extension of 17 March 2026 had its per-domain predictions committed to the public repository at 00:19:19 UTC in commit 8edaec0, before data extraction and twenty-three minutes before the fits, attestation class public repository commit date, and it returned 10 of 12, or 9 of 11 on the strictly fresh rows. It is preregistered in substance, attestation class public repository commit date, and it is not an accepted registry submission; its classifier was the author, its candidate set was the pre-correction one, and the author's demotion of 00:56 that night stands so far as it rests on the packet's own written condition, deposit to a registry before data extraction, which was not met, and is reversed as of 7 September 2026 only so far as it rested on the absence of a registry form.

What would *not* defeat it: a single additional miss, or a better-fitting model appearing in one domain. The claim is about the aggregate operator-to-family mapping, not any single fit; but that aggregate must clear the chance baseline under pre-registered, independently-classified replication.

13. Implications

If the framework holds under independent pre-registered replication, it would suggest:

1. The forms of scaling laws are not coincidences but mathematical consequences of composition structure, under the stated axioms.
2. The same principle that governs how a mouse's heart beats governs how neural networks improve with scale.
3. The renormalisation group's power-law predictions and Cauchy's multiplicative equation may be related theorems - a connection that, to the author's knowledge, has not been formally articulated. A formal proof of this equivalence remains open.
4. Any new system's scaling family can be predicted from first principles by classifying its composition operator, without any data fitting.
5. Two hundred years of domain-specific explanations for scaling laws - fractal networks, learning curves, critical phenomena - may be special cases of a single mathematical constraint that Cauchy identified in 1821.

These implications are conditional. The current evidence is not definitive.

14. Conclusion

Plant an acorn and given centuries, you get an oak. But plant that oak's acorn, and its acorn, recursively across millennia, and you get a forest that shapes the climate of continents. What we plant in these systems will compound across scales we cannot imagine. *The seed determines the forest.*

Cauchy proved 200 years ago that the seed also determines the *form* of the forest. Multiplicative seeds produce power laws. Additive seeds produce exponentials. Bounded seeds produce saturation, though that case is not strictly Cauchy-derived and its exhaustiveness remains an open problem. Twenty-five empirical domains tested against this prediction yield 19 confirmations at $p = 7 \times 10^{-5}$ under the marginal-conditioned permutation test, on the six-model candidate set of Section 3.2; the rerun of 11 August 2026, which added the logarithmic candidate, returns 18 of 25; the four-family re-run is drafted and awaiting human submission. Thirteen published metabolic exponents fall where the dimensional theory predicts.

The mathematics was always there. We just had not read it as a prediction about the physical world. Whether that reading holds under independent scrutiny is now a question for the scientific community to answer.

Raise AI with care.

15. Reproducibility

The complete validation suite - including all 50 domain definitions, the canonical manifest, and the AICc-based fitting code - is available at:

- Experiment code: github.com/.../experiments/cauchy-unification__Paper-VII
- 50-domain manifest: [canonical_50_domain_manifest.json](#)

The script requires Python 3.10+, numpy, and scipy. All data is embedded in the manifest with full provenance citations. Output is deterministic.

Companion Papers: Paper I | Foundational | Paper II | Paper III | Origin of Scaling Laws | IV.a | IV.b | IV.c | IV.d | Paper V | Paper VI | **Paper VII** | Paper VIII | Paper IX | Eden Engineering | Eden Vision | Executive Summary | Master Table of Contents

Research hub: michaeldariuseastwood.com/research | OSF: [10.17605/OSF.IO/X6WA7](https://doi.org/10.17605/OSF.IO/X6WA7) | Copyright 2026 Michael Darius Eastwood

How to cite this paper

DOI: [10.17605/OSF.IO/X6WA7](https://doi.org/10.17605/OSF.IO/X6WA7)

APA (AUTHOR-DATE, 7TH EDITION)

Eastwood, M. D. (2026). Paper VII: Cauchy Unification.
<https://doi.org/10.17605/OSF.IO/X6WA7>.
<https://www.michaeldariuseastwood.com/research/papers/paper-vii-cauchy-unification.html>

BIBTEX

```
@misc{paperviicauchyunification,  
  author = {Michael Darius Eastwood},  
  title = {Paper VII: Cauchy Unification},  
  year = {2026},  
  date = {2026-03-16},  
  doi = {10.17605/OSF.IO/X6WA7},  
  url = {https://www.michaeldariuseastwood.com/research/papers/paper-vii-cauchy-unification.html}  
}
```

INLINE

Eastwood, 2026

See /how-to-cite/ for the canonical citation manual across all artefacts, or /citations.json for the machine-readable index.

Epistemic status. What this programme names Laws are conjectures. Laws I to III are under registered adversarial test. Laws IV and V were added on 24 September 2026; the predictions that carry them are registered, the laws themselves are not yet registered, and neither has been tested. Every quantity in this paper is operationally defined, and established-law standing is claimed nowhere. The registered programme exists to earn that standing, or lose it, by measurement, replication and survived refutation.

© 2026 Michael Darius Eastwood. Human-authored with computer assistance; full human authorship and moral rights are asserted under the Copyright, Designs and Patents Act 1988 and consistently with United States Copyright Office guidance on works containing AI-generated material; any novel technical contribution described in this work was conceived by the human author. Full statement: [michaeldariuseastwood.com/authorship](https://www.michaeldariuseastwood.com/authorship).

Standing covenant. Prove this paper wrong, and I will publish the refutation myself. Falsification conditions are stated in this paper; the standing challenge: github.com/Michael-DariusEastwood/arc-scaling-challenge.

Michael Darius Eastwood conceived and directs this research programme and is the author of this work. Across the programme, he has used more than six AI systems in parallel, under his own instructions, to stress-test his arguments, identify possible errors, and assist in preparing draft text from his own outlines. He determines what is adopted, revised or rejected and takes responsibility for the published content. These systems are tools, not authors.