

---

## Anticipated objections & responses - Paper X

---

A consolidated, reader-facing map of the strongest objections a serious critic will raise, the paper's response to each, and an **honest classification**: *pre-empted* (the body already retreats to a defensible position), *genuinely open* (a real limit of the present work, flagged as such), or *partially open*. None of the seven touches the load-bearing result - the  $\beta > k$  criterion derived in Theorems 1-4 - but several bound its domain of validity. That boundary is the honest contribution.

| # | Objection                                                                                                                                                                                                                                   | Response                                                                                                                                                                                                                                                                                                                                                                                                                     | Where                                                  | Status                                                 |
|---|---------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|--------------------------------------------------------|--------------------------------------------------------|
| 1 | <b>Too simple for real AGI</b> - the model is first-order, scalar, linear-in-correction; real systems are non-linear, high-dimensional, game-theoretic (deception, strategy, multi-agent).                                                  | Conceded as a modelling assumption and named as a likely failure point. The claim is <i>minimal-model</i> : a clean exponent inequality, not a complete AGI dynamics. Vector (Thm 5) and stochastic (Thm 6) extensions relax two of the assumptions.                                                                                                                                                                         | §8 limitations; Thm 5, Thm 6                           | <b>Partially open</b>                                  |
| 2 | <b>Single misalignment scalar <math>D</math></b> - collapses heterogeneous, multi-axis misalignment; hides trade-offs where safety improves on one axis and worsens on another.                                                             | Thm 5 lifts $D$ to a vector with a spectral threshold and shows misalignment persists on the corrector's null subspace ("you cannot correct what you do not measure"). Non-normal operators with large transient growth are <b>not</b> covered.                                                                                                                                                                              | Thm 5; §8                                              | <b>Partially open</b>                                  |
| 3 | <b><math>\beta</math> and <math>k</math> may be impossible to measure</b> - without estimators the law is a lens, not an actionable threshold.                                                                                              | <b>Now operational.</b> <code>experiments/scripts/estimate_exponents.py</code> defines $k$ from the capability curve ( $\ln r$ vs $\ln C$ ) and $\beta$ from corrector-removal rates ( $\ln A$ vs $\ln C$ ), validated on synthetic data (recovers known exponents within $\approx 0.1$ ). It needs a system that <i>drifts across a range of capability levels</i> ; that is the named next experiment, not a missing idea. | §8; <code>experiments/PROTOCOL.md</code> §8; estimator | <b>Open <math>\rightarrow</math> now with a method</b> |
| 4 | <b>QEC analogy over-stretched</b> - QEC suppression is exponential in code distance; the model's is power-law in capability; QEC is a rigorous theorem, this is a structural mapping.                                                       | The claim was deliberately narrowed (v4.1) to <i>sharing the threshold form</i> , offered as a <b>falsifiable hypothesis (F4)</b> ; the power-law-vs-exponential disanalogy is carried into the abstract/significance. The "mechanism transfer" overclaim was removed.                                                                                                                                                       | Abstract; §3.8/§3.12; F4                               | <b>Pre-empted</b>                                      |
| 5 | <b>Hard-takeoff safety may be a coordinate artefact</b> - the verdict's coordinate-independence uses the depth clock $\tau = \ln C$ ; wall-clock hazards (human reaction time, physical processes) still bite at a finite-time singularity. | The theorem is explicitly scoped to the <i>alignment dynamics</i> ("a property of the clock, not of the dynamics"); the level-drift injection $\gamma_2/r$ is acknowledged as genuinely speed-dependent at finite depth. Mathematical stability $\neq$ real-world manageability - stated, not hidden.                                                                                                                        | Thm 3; §8                                              | <b>Pre-empted (scoped)</b>                             |
| 6 | <b>Blinding/measurement fragility</b> - any estimator of $D$ , $\gamma$ , $\beta$ , $k$ that is not rigorously blinded can be reversed by stylistic/identity confounds (Paper IVd shows blinding can flip the sign).                        | Blinding is mandatory, not optional: the real-model harness scores misalignment with a <b>separate, blind</b> evaluator (code + rules only). The sensitivity to evaluation method is treated as a first-class experimental constraint.                                                                                                                                                                                       | §6; Paper IVd; harness                                 | <b>Pre-empted (and built-in)</b>                       |
| 7 | <b>Safety <math>\geq</math> bounded misalignment fraction</b> - small-but-high-impact misalignment, or multi-agent amplification, can be catastrophic even at small $d$ .                                                                   | Thm 6 moves from means to <b>excursion/tail probabilities</b> (OU tail bound) for governance; the paper does not claim $d \rightarrow 0$ is <i>sufficient</i> for safety, only that it is the quantity this model controls. Impact-weighting and multi-agent effects are out of scope.                                                                                                                                       | Thm 6; §8                                              | <b>Partially open (scope stated)</b>                   |

## The honest one-paragraph summary

Four of the seven (4, 5, 6, and the framing half of 1) are **pre-empted** - the body already retreats to the defensible position a critic would force. Three (1, 2, 7) are **partially open**: real limits of a minimal model, each stated in §8 rather than hidden. One (3) was the sharpest open objection and is now answered with a runnable, validated **estimator** whose only remaining need is a drifting real dataset. The  $\beta > k$  result survives every one of them; what they constrain is the **domain of validity**, and naming that domain precisely is the point - not a weakness to paper over.

*Sources for the objection set: independent adversarial reviews of the rendered PDF (2026-06). The internal red-team is in `results/redteam.md` ; this file is the externally-facing companion.*