

External review plan - Paper X

You do not need a PhD or academic rank to get this reviewed. You need to make **one narrow claim** cheap for a competent person to inspect. Do not ask anyone to "review the ARC/Eden programme" - ask them to find a fatal flaw in a single theorem or in the protocol.

Split the review into five lanes (one expert each)

Lane	What to review	The narrow ask
Control theory / dynamical systems	Theorems 1-4, the depth-clock proof, $\beta > k$	"Under the stated ODE assumptions, is the $\beta > k$ criterion correctly derived?"
Linear systems / non-normal dynamics	Theorem 5, spectral abscissa, transient growth	"Is the spectral-stability claim right, and is the non-normal-transient caveat adequate?"
Statistics / stochastic processes	Theorem 6, the OU tail bound	"Is the OU local-Gaussian framing honest, and what breaks it?"
ML evaluations / AI safety	the real-model protocol, blinding, reward-hack task design	"Would the β/k estimator give an interpretable measurement if run across six models?"
Quantum information	the QEC analogy language only	"Is the threshold-form analogy misleading? (It is explicitly not a QEC theorem.)"

Approach those whose work is *closest*, not the most famous first. The paper already cites the nearest neighbours: scalable-oversight scaling laws (Engels et al. 2025), corrigibility (Christiano), model-collapse (Shumailov et al. 2024). A control theorist familiar with gain scheduling / small-gain is ideal for lanes 1-2.

The email (narrow, no endorsement ask)

Subject: Request for adversarial technical review of a $\beta > k$ stability theorem

Dear Dr [Name],

I am an independent researcher preparing a short technical note on a minimal dynamical model of recursive self-improvement. The central claim is narrow: under the stated ODE assumptions, the misalignment fraction is asymptotically controlled iff the correction exponent β exceeds the drift-acceleration exponent k .

I am not asking for endorsement. I am asking whether you can find a fatal flaw, in one specific part: [Theorem 3 / the vector extension / the estimator protocol]. I can offer [honorarium] for a short written technical review, and I would publish your critique and my response, unless you prefer anonymity.

The paper, code, red-team report, and a claim-status ledger (which separates model theorem, internal verification, synthetic estimator validation, and a real-model pilot) are public: [links].

Respectfully, Michael Darius Eastwood

Attach, at most: (1) a 2-page mathematical note, (2) the theorem-proof appendix, (3) the code link, (4) results/redteam.md, (5) a one-page "What this does not claim" (the Scope and Claim-Status boxes). Do **not** lead with the book, the cosmic framing, or priority.

Where to post

- **arXiv** (cs.AI / cs.LG / cs.SY / eess.SY) - needs an endorser; ask one narrowly ("is this in scope and coherent?", not "do you endorse it?"). arXiv is a moderated preprint server, not peer review.
- **OpenReview / workshop** - NeurIPS / ICML / ICLR workshops on AI safety, scalable oversight, evaluations, agents, governance. Main-conference acceptance is unlikely before real-model evidence.
- **Control venues** - IEEE CDC / ACC / IFAC workshops, once the theorem polish lands.
- **Alignment Forum / LessWrong** - strong adversarial AI-safety critique (not formal academic validation).
- **OSF preprint** - already the canonical home; keep it the source of record.
- **GitHub** - open "Adversarial review requested" issues with theorem-specific labels.

After review

Publish a **response-to-review** document even if the reviews are critical - that is the single strongest credibility signal an independent researcher can give. The red-team ledger (`results/redteam.md`) and this repository's correction history already model the posture: find serious defects, fix them in the open, invite more.