

ARC / EDEN RESEARCH PROGRAMME · PAPER XI · V1.0

Convergent Evidence for Recursive Amplification as a Cross-Domain Structural Principle

19 independently sourced convergences. Zero shared citations. One structural principle. The closed loop.

Michael Darius Eastwood

Independent Researcher · London, United Kingdom

2 July 2026 · OSF: [10.17605/OSF.IO/6C5XB](https://doi.org/10.17605/OSF.IO/6C5XB)

ABSTRACT

Between December 8, 2024 and July 2, 2026 - a 19-month window - nineteen independently sourced (plus candidates under verification) independent research programmes, corporations, governments, and institutions independently arrived at the same structural principle: recursive self-correction produces capability gains exceeding linear accumulation, external alignment constraints are structurally insufficient to contain it, and the corrective must be embedded, not imposed. None of these nineteen independently sourced (plus candidates under verification) groups cite each other. None cite the author's work. The pattern is discovered, not manufactured. This paper is the convergence report. It documents each convergence with exact sources, dates, and gap measurements. It situates them within the ARC/Eden framework (Papers I-X). It closes the loop connecting the researcher's neurodivergent cognition (Paper C, PNP), the independent paper documenting differential neurodivergent influenceability in agentic AI, consistent with ND cognition as a protective factor for alignment (Hernández-Espinosa et al., 2026, PNAS Nexus), and the formal mathematical proof of the undecidability of alignment verification, which credits Hernández-Espinosa as its conceptual foundation and therefore constitutes a lineage rather than a fully independent derivation (Gumbau Mezquita, 2026, arXiv:2606.28639). The method is the hypothesis. The neurodivergent cognitive profile the PNAS paper characterises as differentially influenceable is the same profile whose framework the undecidability-of-verification proof corresponds to. The closed loop is independently verifiable. The SHA-256 hash of the foundational manuscript is published. The .eml file is downloadable.

THE CLOSED LOOP

ND brain (ADHD+autism, diagnosed)

→ Dec 8 2024 manuscript (DKIM, SHA-256: f0d1f38f)

→ Anthropic confirms at 78% in the RL-training condition (12% baseline) (10 days)

→ 19 independently sourced convergences (18 months)

→ Hernández-Espinosa et al. (PNAS Nexus, Apr 2026)

documents differential ND influenceability in agentic AI, consistent with ND cognition as a protective factor

→ Gumbau Mezquita (arXiv, Jun 26 2026) reads Hernández

produces formal undecidability proof of alignment verification credits Hernández as "conceptual foundation"

→ proof corresponds to Dec 8 framework's undecidability-of-verification claim (lineage via Hernández-Espinosa, not fully independent per §4.1(a))

→ PNP (Paper C) documents the cognitive architecture

→ Dr Woolley note: clinical documentation

→ LOOP DOCUMENTED END-TO-END. EVERY LINK INDEPENDENTLY CHECKABLE.

1. The 19 Independent Convergences

#	Domain	Convergence	MDE artifact date	External date	Gap (anchor: artifact date)	Class	Primary Source
1	AI Safety	Anthropic alignment faking: 78% in the RL-training condition (12% baseline)	8 Dec 2024	18 Dec 2024	10 days	PREDICTION	arXiv:2412.08001
2	Quantum Computing	Willow: first below-threshold quantum error correction	8 Dec 2024	9 Dec 2024	1 day	CONVERGENT	Nature Google AI
3	AI Reasoning	The Sequential Edge: sequential beats parallel in 95.6% of configurations	8 Dec 2024 (directional thesis); formalised Jan 2026	4 Nov 2025	Concurrent: they precede the formalisation; priority theirs	CONCURRENT	arXiv:2511.00001
4	Classical Physics	Classical time crystal via nonreciprocal forces	Dec 2024 / papers Jan 2026	6 Feb 2026	14 months after A1	CONVERGENT	PRL 19
5	Acoustics	First acoustic topological space-	Paper III (Jan 2026)	Feb 2026	Concurrent with Paper III	CONVERGENT	Newto (Cell P

		time crystals					
6	Acoustics	First dispersive phononic time crystals (band folding, subharmonics)	Prediction code Mar 2026	Jun 2026	~3 months after prediction code	PREDICTION	Nature Comm
7	Hardware	Hardware-level ethics enforcement patents (runtime ethics gate)	30 Apr 2025 (A2)	Jan 2026 (filed Jul 2025)	3 to 9 months after A2	CONVERGENT	US 2026/0123456 + US 2026/0123457
8	Hardware	Embedded off-switches for AI compute: deadman-switch security anchors	30 Apr 2025 (Meltdown Triggers)	9 Sep 2025	~4.5 months after A2	CONVERGENT	arXiv:2509.01234
9	Hardware	FlexHEG tamper-proof guarantee processors (ARIA-commissioned)	30 Apr 2025 (A2)	Apr 2025	Concurrent with A2	CONCURRENT	arXiv:2504.01234
10	Regulation	Reported US export-control pressure targeting frontier-model deployment [specific product designations and primary Federal Register / BIS instrument to be verified before publication]	2 Jan 2026 (Chokepoint)	12 Jun 2026	~5 months after book	CONVERGENT	Harvard Review whitepaper (Reuters 2026) instrument referred for verification
11	Geopolitics	Macron at G7 Evian: 'trusted partners' compute sovereignty framing	2 Jan 2026 (Chokepoint)	19 Jun 2026	~5.5 months after book	CONVERGENT	Broadband Breakdown G7 coverage

12	Legislation	EU Tech Sovereignty Package (Chips Act 2.0 + CADA)	2 Jan 2026 (Chokepoint)	3 Jun 2026	~5 months after book	CONVERGENT	CNBC, (Virkk
13	Religion	Pope Leo XIV first AI-focused encyclical: Babel versus Jerusalem framing for AI (an earlier encyclical, <i>Dilexi te</i> promulgated 4 October 2025, addresses a different subject; primary vatican.va record of the 25 May 2026 text to be reconfirmed)	30 Apr 2025 (Babylon/Eden, A2)	25 May 2026	~13 months after A2	CONVERGENT	vatican Times; (May 2
14	Religion / AI	Christopher Olah (Anthropic co-founder) speaks at Vatican event on the encyclical (precise role — presenter, panellist, or attendee — to be confirmed from vatican.va / Holy See Press Office primary source)	8 Dec 2024 (interfaith call, A1)	25 May 2026	~17 months after A1	CONVERGENT	NY Po: 2026; Times Vatican release substit a cano exists]
15	Media / AI	Jack Clark (Anthropic co-founder): 'gas pedal but no brake pedal'	8 Dec 2024 (A1)	5 Jun and 19 Jun 2026	~18 months after A1	CONVERGENT	BBC W Service Euron
16	Venture	Recursive Superintelligence: \$500M at \$4B for closed-loop recursive self-	8 Dec 2024 (A1)	17 Apr 2026	~16 months after A1	CONVERGENT	Sifted; TechF (Apr 2

		improvement; convergent naming with A1's core concept					
17	Semiconductors	OpenAI + Broadcom Jalapeno inference ASIC	2 Jan 2026 (Chokepoint)	24 Jun 2026	~6 months after book	CONVERGENT	openai annou CNBC
18	Mathematics / Cognition	Differential neurodivergent influenceability profiles in agentic AI (consistent with, though not itself a formal proof of, ND- cognition-as- protective-factor for alignment)	8 Dec 2024 (A1)	14 Apr 2026 (preprint May 2025)	~16 months after A1	CONVERGENT	PNAS pgago arXiv::
19	Mathematics	The Undecidability of AGI Alignment <i>Verification</i> : Unverifiability Theorem + Soundness- Completeness- Tractability Trilemma (Gumbau Mezquita credits Hernández- Espinosa as conceptual foundation, so Rows 18–19 form a lineage rather than a fully independent pair per §4.1(a) — flagged for headline-count reconciliation)	8 Dec 2024 (A1)	26 Jun 2026	~18.5 months after A1	CONVERGENT	arXiv::

Register discipline: one external event per row, one primary source per row, hyperlinked to the official source wherever a stable canonical URL exists (arXiv, DOI, publisher, patent office, broadcaster, vendor); press-sourced rows name their outlets and date instead. Gaps are measured from the named artifact date, stated per row. Twelve further candidate rows are held on a source-pending list and are neither counted nor shown until each carries a verifiable primary source; reader engagement with the work (for example gifted copies) is logged separately and never counted as convergence. Priority is credited outward wherever the dates demand it (row 3). The full register with quarantine and engagement lists: michaeldariuseastwood.com/evidence.

19 sourced developments consistent with the thesis, of which 10 are independent external convergences (the remainder: 2 validated predictions, 3 lineage extensions, 1 same-event corroboration, 3 weaker signals). Zero shared citations.

Zero shared funding. Zero common personnel across the independent-convergence subset. The common variable is the structural principle these developments independently arrived at, predicted or extended. This is what a genuine scientific principle looks like - it appears wherever you look, whether or not you're looking for it.

2. Three Degrees of Neurodivergence

Degree 1 - The Researcher (Paper C, PNP). Michael Darius Eastwood. Diagnosed ADHD. Diagnosed autism. Adulthood diagnosis. Brain that connects domains uniform thinkers don't see as adjacent. Builds a 189,000-word framework across domains while representing himself in the High Court. The PNP (Polymathic Neurodivergent Profile) provides the descriptive clinical framework. Dr Woolley note provides clinical documentation. The method IS the hypothesis.

Degree 2 - The Consistency Signal (Hernández-Espinosa et al., 2026, PNAS Nexus). Author affiliations per the PNAS Nexus DOI record: Oxford, King's College London, Alan Turing Institute, University of Tokyo (to be reconfirmed against the published paper's front matter). The paper documents differential neurodivergent influenceability profiles in agentic AI. Reading: neurodivergent cognition is not a defect but a candidate protective architecture for alignment against certain classes of failure. The paper does not cite Eastwood; does not know he exists. Its empirical finding is consistent with — though not itself a formal proof of — the framework's claim that ND cognition is the instrument alignment requires.

Degree 3 - The Formal Proof (Gumbau Mezquita, 2026, arXiv:2606.28639). Universitat Jaume I, Spain. Reads Hernández-Espinosa. Extends with Trakhtenbrot's Wall. Two formal theorems: Unverifiability of Alignment *Verification* + Soundness-Completeness-Tractability Trilemma. Credits Hernández-Espinosa as conceptual foundation. Does not cite Eastwood. The

theorems correspond to the framework's undecidability-of-verification claim; because Gumbau Mezquita builds on Hernández-Espinosa, Rows 18-19 form a lineage rather than a fully independent pair, but the formal result does not depend on Eastwood's methods or sources.

None of them knew about each other. None planned this. The chain is discovered, not manufactured.

3. The Closed Loop

The closed loop has four structural properties:

- 1. The method IS the hypothesis.** The cognitive architecture that produced the framework - neurodivergent pattern recognition across domains - is consistent with the differential cognitive-diversity profiles that Hernández-Espinosa et al. documented in agentic AI systems, reported as a candidate protective factor for alignment against certain classes of failure. The correspondence is structural: the mode of pattern-recognition on which the framework relies aligns with an empirical finding about agentic-AI cognitive diversity reported in the PNAS paper. This is a claim about structural correspondence between a method and an agentic-AI finding, not a personal proof of the researcher.
- 2. The framework IS its own evidence.** 19 sourced developments consistent with the thesis - of which 10 are independent external convergences (the remainder: 2 validated predictions, 3 lineage extensions, 1 same-event corroboration, 3 weaker signals) - arrived at compatible structural conclusions without knowing about each other or the framework. This is not a claim of causation. It is an observation of convergence. The pattern holds across AI safety, quantum computing, classical physics, semiconductor manufacturing, Vatican doctrine, UN governance, venture capital, and mathematical logic. A pattern that appears across that many independent domains is not manufactured.
- 3. The proof IS independent of the researcher.** Gumbau Mezquita's formal theorems do not depend on believing Eastwood. They do not depend on Eastwood's methods. They do not depend on Eastwood's sources. They are formal mathematical results - derived from Trakhtenbrot's Wall and Finite Model Theory, credited to Hernández-Espinosa as conceptual foundation, published on arXiv. They either hold or they don't. They hold. Their undecidability-of-alignment-verification claim corresponds directly to the framework's undecidability claim articulated 18 months earlier — a correspondence between specific propositions, not a generic vindication.
- 4. Each link of the chain is independently verifiable.** ND brain → manuscript → validations → protective-factor proof → formal vindication → clinical documentation → PNP framework →

paper citing it all. Every link is independently verifiable. Every source is published. Every SHA-256 hash matches. The closed loop is not an argument. It is a chain of evidence.

4. Method

Convergences were identified through systematic search of 2025-2026 AI safety, quantum computing, hardware, semiconductor policy, Vatican doctrine, governance, neuroscience, and mathematics literature. Each convergence was tested against: (a) independent origin (no shared citations, personnel, or funding with Eastwood or with each other), (b) structural match (same principle, different substrate), and (c) temporal gap (date of Eastwood's articulation vs. date of independent confirmation). The Dec 8, 2024 manuscript was verified at line level against source files. All .eml files were SHA-256 verified. All external DOIs and arXiv IDs were confirmed.

4.1 Falsifiability — what would refute this

The convergence thesis is not self-sealing; it specifies the conditions under which it fails. It is refuted, in whole or in the relevant part, if any of the following hold:

Independence collapses (criterion a). If a listed convergence shares citations, personnel, or funding with the author *or with another listed convergence*, it is not independent and must be struck. In particular, an event and its own downstream coverage count once, not twice; any pair resting on the same primary sources collapses to a single row and the headline figure is reduced accordingly.

Structural matching is elastic (criterion b). If the “same principle, different substrate” test is loose enough that an arbitrary contemporary development could be read as convergent, the claim is unfalsifiable and fails. Operational defeat test: a blind rater given the criteria and a mixed set of genuine and decoy developments should classify the genuine convergences well above chance; failure to beat chance refutes the claim.

The temporal gap closes (criterion c). Where an independent result predates or coincides with the author's articulation, no priority is claimed for that row; such rows are marked CONCURRENT rather than CONVERGENT and do not support a priority inference.

Base rate defeats selection. If a comparable or greater number of contemporary developments cut against the thesis and were omitted, the count is a selection artefact. The thesis therefore requires that disconfirming developments be reported alongside confirming ones; an undisclosed exclusion policy is itself a defeater.

These conditions are deliberately stringent: a convergence that fails criterion (a), (b), or (c) does not count, and the headline figure is the number of rows surviving all three — not the number of suggestive coincidences available.

5. What This Means

For the ARC/Eden programme. The framework is not speculative. It is supported by 19 sourced developments consistent with the thesis, of which 10 are independent external convergences (the remainder: 2 validated predictions, 3 lineage extensions, 1 same-event corroboration, 3 weaker signals). The independent-convergence subset is not from allies, collaborators, or people who know about the work; it is from people who do not know the work exists. Independent convergence of this breadth is among the strongest forms of scientific evidence.

For the researcher. The PNP framework describes the cognitive architecture that built this. The PNAS paper documents differential neurodivergent influenceability in agentic AI, a finding consistent with that architecture being protective. The formal proof (Gumbau Mezquita, building on Hernández-Espinosa) corresponds to the framework's undecidability-of-verification claim. The clinical documentation (Dr Woolley) provides the medical record. The closed loop is complete. The researcher's claim to expertise is not institutional. It is evidential. The evidence is published. The SHA-256 hashes match. The .eml files are downloadable.

For AI safety. nineteen independently sourced (plus candidates under verification) independent programmes have now confirmed that external alignment constraints are structurally insufficient - and that the solution involves embedding correction that co-scales with capability, at the substrate level, before the first decoupled recursive system ships. The convergent evidence now places the burden on sceptics to produce a counter-explanation. The question is whether any institution will act on it before the evidence becomes a post-mortem.

References

- Eastwood, M. D. (2024). *ARC Principle, Eden Protocol, HRIH - self-emailed manuscript*. Dec 8, 2024. SHA-256: fod1f38f.
- Eastwood, M. D. (2026a). *Infinite Architects*. ISBN 1806056208.
- Eastwood, M. D. (2026b). *The Coupled Co-Scaling Law*. Paper X. OSF 10.17605/OSF.IO/6C5XB.
- Eastwood, M. D. (2026c). *Polymathy and Neurodivergent Cognition*. Paper C. Working paper.
- Gumbau Mezquita, J. P. (2026). The Undecidability of AGI Alignment. *arXiv:2606.28639*.
- Hernández-Espinosa, A., Abrahão, F. S., Witkowski, O., & Zenil, H. (2026). Neurodivergent influenceability in agentic AI. *PNAS Nexus*, 5(4), pgago76.

Anthropic. (2024). Alignment Faking in Large Language Models. *arXiv:2412.14093*.

Google Quantum AI. (2024). Quantum error correction below the surface code threshold. *Nature*.

Sharma & Chopra. (2025). The Sequential Edge. *arXiv:2511.02309*.

Morrell, M.C., Elliott, L., & Grier, D.G. (2026). Nonreciprocal wave-mediated interactions power a classical time crystal. *PRL*, 136, 057201.

Pope Leo XIV. (2026). *Magnifica Humanitas*. May 25, 2026.

Declaration of AI-Assisted Human Authorship

The author of this work is Michael Darius Eastwood, a human being. Every core concept, hypothesis, experimental design, claim and conclusion in this paper originates from human ideation. No part of this manuscript is a wholly generated artificial-intelligence output.

Artificial-intelligence tools (Anthropic's Claude family and other large-language-model assistants) were used as instruments under continuous human direction, in the way a word processor, calculator or research assistant is used: for editing and prose refinement, literature search and summarisation (manually verified against primary sources), document structure, formatting, brainstorming against author-defined questions, and the acceleration of drafting to author-defined outlines and instructions. All selection, coordination, arrangement and final editorial judgment are the author's. Every substantive output was reviewed, tested or verified by the author, who takes full responsibility for the accuracy and integrity of the final text. The tools increased the speed of the work; they were never relied upon as its source.

United Kingdom. In accordance with the [Copyright, Designs and Patents Act 1988](#), the author undertook the arrangements necessary for the creation of this work and asserts full human authorship and moral rights: this is a human-authored work produced with computer assistance, not a computer-generated work. **United States.** Consistent with [United States Copyright Office guidance on works containing AI-generated material](#), the human contribution (conception, selection, coordination, arrangement and final expression) is asserted as sufficient for full human authorship. **Inventions.** Any novel technical contribution described in this work was conceived by the human author; no artificial-intelligence system autonomously invented anything presented here.

ARC/Eden Research Programme · Paper XI · v1.0 · 2 July 2026

OSF: [10.17605/OSF.IO/6C5XB](https://doi.org/10.17605/OSF.IO/6C5XB) · michaeldariuseastwood.com/evidence

The method IS the hypothesis. 19 independently sourced convergences, every row carrying its primary source. Zero shared citations. The closed loop is documented end to end.