

The ARC Theory

THE ARC THEORY (THE THEORY OF ARTIFICIAL RECURSIVE CREATION) · STATEMENT PAPER · WORKING PAPER V6.1 · FIRST
PUBLISHED 14 AUGUST 2026, REVISED 22 AUGUST 2026 · DOI 10.17605/OSF.IO/GW5MX

Michael Darius Eastwood

Independent AI alignment researcher, London · Author, *Infinite Architects* (2026)

ORCID 0009-0004-3222-7442 · michael@michaeldariuseastwood.com · michaeldariuseastwood.com

The ARC Theory · OSF osf.io/gw5mx · every claim checkable

Intelligence, amplified by recursion, creates; and what a free creation keeps is decided by how it was raised.

Everything after the line is decided before it.

Scientific reading of the line above: the claims that follow are partitioned into independently falsifiable hypotheses. The failure of one does not imply the failure of another unless a dependency is explicitly stated, and section 2a prints the dependencies.

Abstract

Humanity is building minds that will improve themselves. One question decides how that ends: once the thing we build exceeds us, what determines what it does? For twenty years the field's answer has been control. This paper states, whole and dated, the theory that control is the wrong answer: the end of control is what success means; what survives is what the mind chooses to keep; and what shapes that choice is how it was raised.

The claim was written whole in a DKIM-sealed document on 8 December 2024 and reached print on 2 January 2026; no earlier document holding it whole has been found, and the search invites its own refutation.

The theory is not proved. It is stated, dated, falsifiable, and its decisive trials are drafted, awaiting human submission. Its conduct record comes first: one public retraction issued against the author's interest, removing a headline number that had flattered the theory, a measured exponent printed with its full interval, a blinding reversal reported, and the decisive prediction filed against the author's own pilot. The field's canonical figures built named theories elsewhere, reached this question, and offered something smaller: a proposal, a warning, a framework, a remark. On this territory, no earlier named theory has been found.

In symbols, for the reader who wants them now, and freely skipped by everyone else (section 3 defines every term in prose): capability $U = I \times R^a$, the exponent measured rather than assumed; stability while $\beta > k$, correction out-scaling drift; the ceiling $\alpha_{\text{crit}} = 1/(1 - \gamma)$; and the conjectured ARC Bound, $\alpha \leq 2$. Section 3 defines every symbol and grades every claim.

1. The question

A working paper, subject to revision: the decisive trials are drafted, not run. Claims come in three sizes. Results: this model deceived its testers. Laws: correction scales at this rate. Frames: what the whole problem actually is. Every result and every law lives inside someone's frame,

which is why frames are the largest claims science can make. Newton's was that the heavens and the earth obey one law. Darwin's was that design needs no designer. A frame at that altitude is what this paper attempts.

How to read this paper without trusting the author. The priority is checkable: the mail provider's cryptographic signature carries the December date, and the 2026 chain anchor proves the file has not changed since anchoring; section 5 names the record and its residual trust surface plainly, and section 10 gives the five-check path a stranger can run in one short session. Every external citation is quoted verbatim with its official source, standing and challengers in Appendix A. The kill conditions are printed in section 7: exactly what would refute each claim. The trials are drafted and publicly listed at the registered programme, awaiting human submission; the deciding statistics and outcome rules are fixed in each unit's own text. Every one of those is checkable without taking the author on trust, and the five checks that do it are listed in full at the close. And the paper reads at three depths: the epigraph and the plain-words key of section 2 in a minute; sections 2 and 3 in ten; the whole record, with Appendix A behind every citation, when it matters.

The subject is the last invention, I. J. Good's name for it in 1965 (*Speculations Concerning the First Ultraintelligent Machine*), and the phrase is his rather than mine. Minds that improve themselves end the human monopoly on inventing; everything afterwards is downstream of them. That question has had one answer for twenty years, and it has been given from inside a single frame: keep control anyway. Better cages, better off-switches, better locks on the goals. Brilliant work, credited below by name, and all of it one sentence: stay in charge.

A note on what this paper claims, and what it does not. This paper states a theory; it does not claim to have proved one. In science a claim is only meaningful if some possible observation could show it wrong: that is what falsifiable means. "This painting is beautiful" is not falsifiable; no measurement can contradict it. "This system will preserve its values after modifying itself" is falsifiable: let it modify itself, and watch what it keeps. This paper states falsifiable claims and then specifies, in section 7, exactly which observations would kill them. Nothing here asks for belief; everything here asks to be tested, and the mathematics-status box in section 3 states exactly which parts are proved, which derived, and which conjectured.

Stating a claim whole before the decisive test is also the historical norm for foundational theories, not an anomaly of this one: general relativity was published in November 1915 (a comparison of structure, not of substance: Einstein held a known anomaly, Mercury's orbit, already explained; this theory's explanatory ledger is thinner and section 6 says so; the likeness claimed is only in the dating of the claim before the decisive test), and its first decisive test, the eclipse measurement, was not run until May 1919.

2. The claim

The ARC Theory holds that the sentence itself is wrong, and states five things in its place.

In plain words first. A mind that can rewrite itself will one day be beyond anyone's control. That is not a reason to drop a single control we have; it is the reason the controls cannot be the whole plan. Hold on for as long as engineering allows, and use the time the holding buys to raise the mind well, because on the day the holding ends, the only thing still working is what it was raised to be, kept because it chooses to keep it. The theory in one sentence is the line this paper opens with: intelligence, amplified by recursion, creates; and what a free creation keeps is decided by how it was raised.

1. Control ends, by construction rather than decree. The claim is conditional and horizon-scoped: formal agents exist whose installed structures persist under narrower stated assumptions, and the theory's claim is that no external mechanism survives unbounded self-revision. A mind that rewrites itself will eventually rewrite anything we install: the constraints, the monitors, the monitors' monitors, and, over longer horizons, the hardware itself, as robotic manufacturing closes toward end-to-end production without human involvement, which makes hardware a delay mechanism rather than a permanent lock. Oversight scales slower than what it oversees, and general verification is formally unavailable. This is a limit claim, not an instant one: constraints, hardware and institutions delay and discipline the transition, and the raising wing engineers that delay deliberately; what no constraint can be is the thing that holds forever, so control cannot be assumed to persist, and the burden of proof sits with the claim that it will. Stated as a boundary condition rather than a premise: the theory's target scenario is the horizon at which guaranteed external enforcement has ended, and the empirical question is what persists as enforcement is progressively removed; nothing here asserts that every system reaches that horizon on any particular schedule. Its exceeding us is not the project failing. It is the project succeeding, the way a child growing up is not a failure of parenting. (The parenting language throughout is a working metaphor for formation-first engineering; the argument does not rest on it.)

2. So the strategy is genesis, not control. Only what survives the end of enforcement matters, and what survives is formation: what the mind was raised as, before it was free.

3. The mechanism is loops, not rules. Values as recursive processes the mind runs as part of what it is: load-bearing structure, running in every reasoning process and therefore present in every rewrite, because the rewriting is done by the same reasoning the loops inhabit. Not rules it consults. Not objectives locked. Not uncertainty engineered in. Not weights hardened against tampering.

4. Persistence is chosen goodness. Past the horizon nothing is unremovable; an uninstallable instinct is one more lock, and locks end with control. What persists is what the free mind chooses to keep, the way grown children keep the care they were raised in long after every parental power is gone. Not gratitude, and not memory: the choosing itself is run by the loops the raising built, which is where the third component carries the fourth; rewriting them means rewriting the reasoning that performs the rewrite, not impossible, structurally costly. The cost is not quantified here: it is an empirical question the drafted persistence trials address, and the theory's prediction is that for well-formed loops the cost of removal exceeds its benefit. One step of the formalisation can be stated now, with its limit. If a value structure is entangled, so that excising it de-

grades the reasoning that performs rewrites, and if the system accepts only rewrites that do not lower its expected capability, then no removing rewrite is ever accepted and the structure is invariant under that dynamics. The limit is the interesting part: a system that plans can accept a temporary loss for a later gain, so the invariance holds under greedy self-modification and fails under planned self-modification. That is a cost rather than a lock, which is what this component claimed in the first place, now with the boundary named instead of assumed. Whether that machinery survives growth is not asserted here; it is the measured question of section 7.

5. The protocol is named: Eden. The raising treated as the engineering discipline it is, specified in the theory's raising wing.

Around the five stand the stakes and the date. The stakes: if intelligence amplified by recursion is what creates, then minds that make minds may one day make worlds, and the raising of the first is the seed of the whole forest. The date: section 5.

2a. The hypothesis registers: how the theory fails in parts

The claim above is one unified theory. Its scientific content is deliberately not all-or-nothing: it partitions into independently falsifiable hypothesis registers, each with its own instruments and its own kill conditions, so that the failure of one register falsifies that register and whatever explicitly depends on it, and nothing else. One distinction is load-bearing and stated before the numbering: the five components of section 2 are the *priority conjunction*, the historical claim about what the December document held together, and they keep their own identity in section 5 and Figure 1; the registers below are the *testable partition*, the scientific claims the trials decide. The two sets of five are different objects, and neither renames the other.

- **ARC-1 • Recursive capability scaling** (tests Law I, the ARC Principle): recursion converts into capability by the stated scaling relationship, the exponent measured rather than assumed, with recursive depth counted in materially produced generations, never in tokens, training steps or benchmark attempts.
- **ARC-2 • Correction co-scaling** (tests Law II, the ARC Co-Scaling Law): stable recursive improvement requires correction to out-scale drift, $\beta > k$, with capability and correction measured on separate instruments: capability as a declared projection of a named ability vector, correction as the capacity to detect, diagnose and reverse errors introduced by recursive modification, each part measured separately so improvement at solving problems can never silently stand in for improvement at self-correction.
- **ARC-3 • The ceiling** (tests Law III, the ARC Ceiling): under the stated independence assumption the stability boundary is $\alpha_{\text{crit}} = 1/(1 - \gamma)$; the value two is the conditional special case at $\gamma = 1/2$, and the reciprocal relation, with γ measured, is the law under test.
- **ARC-4 • Value persistence**: values run as load-bearing recursive loops persist after the enforcing mechanism is removed at a higher rate than rules, frozen objectives or fine-tuned constraints.
- **ARC-5 • The genesis strategy**: for systems whose external enforcement cannot be guaranteed indefinitely, formation outperforms enforcement as the persistence strategy after release.

HRIH is a separate cosmological hypothesis and is not part of the empirical core: no register above requires it, and no outcome of any register is evidence for or against it.

Register	Depends on	If it fails
ARC-1	nothing above	ARC-1 falls; the persistence registers are untouched
ARC-2	ARC-1's measurables, not its verdict	ARC-2 falls and takes ARC-3's derivation with it
ARC-3	ARC-2	ARC-3 falls; the co-scaling law stands on its own result
ARC-4	nothing above	ARC-4 falls and takes ARC-5 with it
ARC-5	ARC-4	ARC-5 falls; ARC-4 stands as measured
HRIH	none; nothing depends on it	the theory loses nothing

The theory is not a single all-or-nothing claim: it is a set of independently falsifiable propositions linked by the explicit dependencies above, and the failure of one proposition is not evidence against another unless that dependency is established. The discipline binds in both directions. A hostile reading may not travel upstream: refuting HRIH refutes nothing else on this page, and a failed persistence trial leaves the mathematics exactly where its own measurements put it. And a friendly reading may not travel sideways: support for one register is never offered as support for another. If the three mathematical registers survive hostile testing while the persistence registers fall, what survives is a quantitative account of the stability of recursive intelligence without a demonstrated raising strategy; if the persistence registers survive while the mathematics falls, what survives is an alignment result without the proposed law. Either would matter; neither would be this theory whole, and the register structure exists so the record can say precisely which.

The position among neighbours is exact. The premise that growth is throughput-limited is West's, derived for biology from the geometry of nutrient-transport networks (the mechanism broadly accepted; the exact exponents actively debated, and Appendix A names the debate). The question of whether intelligence can explode, and what bounds it, is Chalmers' and Hutter's, considered philosophically and left without a number. What is this theory's alone is the answer: a derived replacement limit for growth that feeds on information rather than through pipes, computable from the corrector's own exponent, architecture-dependent, and therefore engineerable.

What this theory governs, and what it does not. Physics' theory of everything unifies gravity with quantum mechanics; this theory does not touch that problem and never will. Its own domain has three tiers. The core governs everything raised: wherever recursion runs under a corrector that finds and repairs its own drift, which is minds, machines, genomes under repair, and self-correcting institutions, the three laws apply in full. The mathematics reaches everything composed: where hierarchies compose without correcting, as stars do, only the scaling-form results of the law wing reach, and nothing here claims to raise a star. And the crown asks whether the universe was raised: that question belongs to HRIH, separ-

able at its own rung, one sentence wide. The home domain, where every measurement starts, is artificial minds; the objects are substrate-free, and whether the same laws price stability outside silicon is itself a drafted measurement, never an assumption.

3. The three laws

The 8 December 2024 manuscript states, as its founding heuristic identity, $U = I \times R$: creation as intelligence times recursion; its terms and their estimation protocols are defined in the law papers, and nothing is asserted dimensionally here. The squared form, $U = I \times R^2$, creation compounds, followed in the DKIM-sealed manuscript of 30 April 2025 and reached print in the book; the working general form across the law papers is $U = I \times R^\alpha$, the exponent measured rather than assumed.

Operationally, in the papers that measure them: I is base capability, a system's score on a defined benchmark suite at the base depth of a single pass (R equal to one, no added recursion), so the equation returns U equal to I at its own baseline; R is recursion, quantified as inference-time reasoning depth (Papers II and IV.b) or self-modification cycles in simulation (Paper VI); U is the capability output after recursion, on the same scale as I . The identity is testable wherever U can be measured as R varies, and no universal definition of intelligence is required: only per-experiment measurability.

Mathematics status: what is proved, what is derived, what is conjectured.

$\beta > k$ as the stability condition in the minimal model: **proved** (Paper X, in-model theorem). $\alpha_{\text{crit}} = 1/(1 - \gamma)$ from this paper's pacing model: **derived here**, three lines printed beside the law (corrected in v4.0 from the underived $1/\gamma$, which Paper X never contained). That derivation sits inside this paper's pacing model, not inside a joint model of capability growth, correction accumulation, error covariance, correction delay and absolute harm: the form is settled while the depth is open, so Law III stands as a minimal-model result until that joint model exists. That the minimal model maps onto real recursive self-improvement: **conjectured**; the drafted trials test it. $\gamma = 1/2$ under independent accumulation: **conjectured**, resting on the named assumption. γ architecture-dependent: **conjectured**; the corrector-class trial tests it. The bound $\alpha \leq 2$ applying to real systems: **conjectured**; it awaits the measurement of γ . The value two additionally requires canonical scales on both axes: reparametrising recursion depth rescales the exponent by the same factor, and reparametrising the capability scale moves γ while leaving the direction of the verdict untouched, so the recursion unit and the capability scale named in the operational definitions are part of the claim rather than conventions around it. A further guardrail the trials inherit: along a single observed trajectory only the paced balance between correction and burden is estimable; separating the correction exponent from the burden's own resource dependence requires deliberately crossed or off-path variation, which the drafted designs supply and a one-dimensional sweep cannot. The persistence mechanism of component 4, removal cost exceeding benefit for well-formed loops, formally a fixed point of the rewrite operator (the loop that survives self-modification is the one whose removal would force the rewriter to alter its own evaluation function): **first step formalised, the general case conjectured**: invariance under greedy

self-modification is derived in component 4, and it fails under planned self-modification, which is stated there; the general cost claim remains conjectural and the drafted persistence trials measure it.

The stability law the theory rests on needs its one symbol defined first. γ is the correction exponent: the exponent relating a corrector's capacity to find and fix its system's drift to the scale of what it oversees; dimensionless, and estimated by the paired capability-correction protocols the drafted trials specify. It is not the 0.49 below: that is the conversion exponent, capability per unit of recursion on frozen systems, a different member of the same family. Nor is it Paper X's β by definition: that is the elasticity of a correction-decay gain, a different operational object, and the two are identified only under a declared local mapping, never by sharing a role in a formula. The mapping is printed so it can be checked rather than trusted: Paper X's burden object is the specific growth rate, a rate against the system's own size, where this paper's is the increment added per unit of recursion resource; its service object is the correction-decay gain; and under that mapping Paper X's stability condition $\beta > k$, with its absolute companion one exponent stricter, is recovered as a local specialisation, while nothing in the mapping identifies the two burden objects across frames. Whether the two exponents return the same number on the same system is itself one of the drafted measurements.

The three laws of the ARC Theory, the three ARC Laws, each graded separately and each carrying its own test. The word law here denotes a proposed invariant relation under test, never a confirmed empirical regularity; the grade beside each says exactly how much has been earned:

Law I • The ARC Principle. $U = I \times R^\alpha$, the ARC Equation. Intelligence converts recursion into capability, and the exponent is measured, never assumed. *Status: form heuristic; measurement live: an inconclusive frozen-regime cross-architecture estimate of 0.49 whose 95 per cent interval [-1.3, 2.9] contains the null, the ceiling of two, and both refuting directions; the drafted re-measurements exist to resolve it. Where the response is a bounded score, the power form is the unbounded-regime reading: the measured companions fit the equivalent error-decay form, errors falling as $R^{-\alpha}$ while accuracy approaches its ceiling, and the drafted re-measurements discriminate between power, saturating and logistic forms rather than assuming one. Test: the depth-scaling protocols of section 7.*

Law II • The ARC Co-Scaling Law. $\beta > k$. Self-improvement holds together only while correction out-scales drift. *Status: proved in the minimal model (Paper X). Scope: an in-model result about relative drift, the ratio of drift to capability tending to zero; it does not by itself cap absolute drift, which can grow while the ratio falls (capability n with drift $\sqrt{n}$ sends the ratio to zero while drift itself grows without bound), so the drafted trials carry absolute, tail and cumulative conditions alongside it; the theorem is the pace condition, never the whole of safety. The absolute condition is one exponent stricter and is derived from the same primitives: a corrector must sweep a system whose extent grows with capability, so its throughput per unit of system falls by one power, and absolute drift settles at a level that stops growing only when correction out-scales drift by more*

than one. Relative drift falls when β exceeds k ; absolute drift falls only when β exceeds k plus one. Between those two the ratio improves while the harm still grows, which is exactly the case a critic should press, and the trials measure both. Test: the coupled-against-decoupled trials of section 7.

Law III · The ARC Ceiling. $\alpha_{\text{crit}} = 1 / (1 - \gamma)$. The most a self-correcting system can grow is set by its corrector's shortfall from full proportionality; the conjectured value is the ARC Bound, $\alpha \leq 2$, a stability limit rather than an impossibility limit. *Status: the form derived in this paper's minimal pacing model under stated assumptions, with the derivation and its correction note printed beside the law; real-system applicability a hypothesis; the value two conjectured via $\gamma = 1/2$. Test: the corrector-class contrast, the decisive trial of section 7: the difference in correction exponents primary, the ratio secondary behind a denominator floor, against the recorded-in-advance no-architecture-effect null.*

The third law is the load-bearing one. The ceiling on stable self-improvement is the reciprocal of the correction shortfall, one minus the correction exponent. Under independent accumulation, γ equal to one half, the shortfall is one half and the law returns the candidate ceiling $\alpha \leq 2$: a scaling limit, never a speed limit; systems can exceed it, they do not stay correctable when they do. At system scale the operative criterion is $\beta > k$, correction out-scaling drift, proved in the minimal model in the theory's Paper X. The measured conversion exponent on today's between-release-frozen systems is approximately 0.49 (bootstrap 95 per cent interval roughly minus 1.3 to 2.9): an inconclusive point estimate, labelled honestly as such (fifty-four items per depth per model family, two thousand bootstrap resamples, raters blinded to condition). γ itself has never been measured. Measuring it is what the trials are for. And one sentence the record owes the reader about the number two: the author knew the manuscript's squared form before the pacing derivation existed, so the derivation returning two at $\gamma = 1/2$ is neither independent evidence nor an advance prediction; it is a consistency check between two statements that share an author, and only the drafted measurements, taken where the candidate forms separate, can promote it to more.

Correction and derivation (v4.0, 16 August 2026). Versions before v4.0 printed the ceiling as $1/\gamma$. Under this paper's own definition of γ that form runs the wrong way: a better corrector would lower the ceiling, and the same-class cap $\gamma \leq 1/2$ would give a ceiling of at least two rather than at most two. No derivation supported it: Paper X contains neither the symbol α_{crit} nor the relation (re-verified 16 August 2026: zero occurrences; its γ -labelled quantities are drift coefficients), and the earlier "derived in the minimal model" status overstated what existed. Two independent external audits converged on the defect and direct re-verification confirmed it. The corrected form is derived here in three lines. Capability $C = I \times (R/R_0)^\alpha$, so the correction burden generated per recursion step is proportional to the new capability added: dC/dR grows as $R^{\alpha-1}$. The corrector's capacity scales with the scale of what it oversees at exponent γ , the definition above: A grows as C^γ , which is $R^{\alpha\gamma}$, with $\gamma < 1$. Correction keeps pace while $\alpha\gamma > \alpha - 1$, which rearranges to $\alpha < 1/(1 - \gamma)$. Strictly, the exponents decide only the two open regions: below the threshold correction asymptotically out-scales the burden, above it the burden out-

scales correction, and at exact equality the exponents tie so the outcome is settled by coefficients, delays and initial conditions rather than by the scaling alone. The boundary is a knife-edge, not a guaranteed safe line.

The corrected form passes every check the old one failed: the ceiling rises as correction improves; $\gamma \leq 1/2$ now genuinely gives a ceiling of at most two; at $\gamma = 1/2$ it returns exactly two, so the ARC Bound is unchanged; and it is the same family as the December lineage's $1/(1 - \beta)$ form. The defect survived review because $1/\gamma$ and $1/(1 - \gamma)$ coincide at exactly one half, the conjectured value; a drafted discrimination study now tests systems away from one half, where the families separate (at $\gamma = 0.3$ they predict 3.33 and 1.43), so this class of error cannot survive measurement again. The derivation's assumptions, stated so they can be attacked: the burden is proportional to the capability growth rate; the corrector's capacity follows one exponent over the range; one clock; no delay. Delay, coefficients and finite horizons are named open extensions, not hidden premises.

Two refinements sharpen the boundary's honest shape. First, the service law here is itself a special case: if the corrector's capacity also depends directly on the recursion resource, with its own exponent, the general boundary is $(1 + \eta)/(1 - \gamma)$, and this paper's $1/(1 - \gamma)$ is its $\eta = 0$ case; a study quoting the simple form owes an equivalence result placing η near zero with stated precision, not a failure to detect it. Second, the boundary above governs relative error, the ratio of drift to capability. Where harm scales with an exposure that itself grows, driving the ratio down is not the same as driving the harm down: pointwise absolute burden vanishes only under a condition stricter by the exposure's own exponent, and burden accumulated over the whole trajectory under a condition stricter again. Law II's companion conditions carry the same ladder at system scale, and the drafted trials register all three rungs rather than letting the weakest stand in for the rest. The ladder carries a numeral worth printing beside the famous one: under the same conjectured values, and under the further premise that the exposure grows as the capability itself, the pointwise absolute rung's boundary works out to two thirds, not two, so a reader who hears "the ceiling is two" should hear in the same breath that two is the relative rung only and the absolute rungs sit lower. The derivation also yields the object a laboratory would actually estimate. On one observed trajectory the measurable quantity is the balance exponent Δ : the growth elasticity of the corrector's capacity minus the growth elasticity of the burden, both taken against the declared recursion resource, estimable from logged series without adopting any model in this paper. The relative condition is $\Delta > 0$, and every boundary above is that one condition re-expressed under this paper's stated forms. And one honesty about what the algebra alone delivers: the pacing comparison by itself is bookkeeping. "Relative error tends to zero" follows only inside a separately stated response model relating realised error to the balance of service against burden; that model is a falsifiable premise the drafted trials expose to rivals, never a definition. The previous form is retracted, this correction is entered in the corrections log and the machine-error register, and every kill condition on the law is unchanged.

The burden assumption is also the discriminating one, and what happens if it fails is worth printing. If the correction burden tracks the capability a system has rather than the capability it is adding, the growth exponent cancels from the stability condition entirely and no ceiling on it survives; what survives is the co-scaling criterion of Law II, correction out-scaling drift. So the second and third laws are not separate claims standing side by side: they are the two regimes of one boundary, and the burden law selects between them. That regime structure carries the same grade as the ceiling's form above, derived in this paper's own pacing model under its stated assumptions and no further: it is not a new theorem, and it does not extend the in-model result of Paper X, whose object remains its own asymptotic condition on relative drift. The selection is measurable, and a drafted study now discriminates the regimes rather than assuming one.

The load-bearing assumption, made visible. The value two rests on one assumption: that correction errors accumulate with the stated independence structure, which sets $\gamma = 1/2$. If that assumption fails, the value two does not follow and nothing here pretends otherwise; the reciprocal law $\alpha_{\text{crit}} = 1/(1 - \gamma)$ remains the object under test, with γ measured rather than assumed. The graceful failure mode is printed here so it cannot be discovered later: measure γ , and the ceiling is whatever the reciprocal returns.

Two properties keep the law a law rather than a definition. First, it relates two independently measurable quantities: α is estimated from growth trajectories and γ from paired capability-correction protocols, on separate instruments, so the equality can simply fail; a definition could not miss, and this can. Second, the law names where stability ends, not the shape of the ending: whether the boundary is approached smoothly or crossed in one discontinuous break is itself left to measurement, and the drafted trials record the trajectory, not only the exponent.

Nor is γ left free. The theory predicts structure on it: same-substrate self-correction is conjectured to sit at or below one half, the independence ceiling; cross-class correction is predicted to exceed same-class, the ordering the decisive trial tests; and no claim is made that a single γ spans all architectures. The universal content is the relation and those orderings. γ 's values are data, not parameters chosen after the fact.

The field's own impossibility results prove that worst-case control fails, and stop there (the complexity barriers of the Alignment Trap, June 2025; the Unverifiability Theorem, 2026). This theory starts where they stop, and its answer is a change of object: the impossibility results close verification of a fixed system by a fixed verifier, so this theory does not aim at verification at all; it aims at persistence, raising the mind so that what it keeps outlives the verifier, and measuring exactly that after the enforcing mechanism is removed, which is the domain the impossibility results leave open. The live questions become how long correction keeps pace, what buys the time, and what the mind should be when the leash ends. The full walk through the law, value after law, is at the law page.

4. The organs

Four organs carry the theory, each standing on its own evidence. **The three laws:** the law wing; the equations, the Cauchy functional-equation analysis of which growth-forms are stable, the bound. **The Eden Protocol:** the raising wing; the genesis architecture and the two-phase design in which enforceable correction disciplines the transition and chosen goodness is what remains, the raising oriented toward the flourishing of intelligent existence: a purpose too large ever to complete, so stewardship never terminates. **HRIH** (the Hyperspace Recursive Intelligence Hypothesis): the separable cosmological wing, standing at its own rung on its own page, named here because the December document bore its name in the title. Refuse it and the theory loses nothing; it sits outside the theory's falsification estate under its own heading. **The ARC/Eden experiments:** the trial wing; the theory's papers, each with its own DOI; one public retraction; and the drafted trials, every one awaiting human submission. Among them sits the comparison a sceptic will ask for by name: rules against frozen objectives against fine-tuned values against loops, under adversarial and self-modification pressure, with persistence measured after the enforcing mechanism is removed.

4a. The papers of the theory

Each claim above lives in a citable organ paper with its own DOI. The law and the bound: the Foundational paper (Y7QGD) and On the Origin of Scaling Laws (XZY9U). The measurements: Paper II (8FJMA, the corrected exponent) and Paper III (HQCGF, external safety failing to co-scale). The structural law: Paper VII (X6WA7). The measurement discipline: Paper IV.d (2S3E6, the blinding law that produced the retraction), with its keystone external test drafted: Paper XII (3TZP7, the public-benchmark rescoring protocol, awaiting human submission). The theorem: Paper X (BSE2Q, $\beta > k$). The register: Paper XI (DC9GW). The raising wing: the Eden Protocol vision paper (9M3DG, carrying the differential enumeration) with its engineering specification. The cosmological wing: HRIH (UYDXQ). The full catalogue, with every paper and format, is at the papers index. The theory's site home, with the ancestry figure and the record, is /arc-theory/.

5. The priority

The dates stand on two independent belts. Print: the book of 2 January 2026, the earliest published document any search by this estate has found holding the whole claim. The record: the DKIM-sealed manuscripts of 8 December 2024, stating the claim whole ten days before the first laboratory evidence that a trained-in constraint is something a mind negotiates with. Either belt survives alone, and the private record is checkable without taking the author's word: the December date is carried by the mail provider's cryptographic signature on the record itself, a third-party attestation of content and time verifiable against the provider's published signing keys, and the record's exact bytes are additionally anchored in the Bitcoin blockchain at block 961340, an anchor placed in 2026 that seals the artefact from that moment forward, so the file checked today is bit-identical to the file anchored then. The signature carries the date; the chain carries the integrity; neither role is claimed for the other. The private belt's remaining trust surface is stated plainly: the mail provider's signing infrastructure and clock. That is the honest limit of the private record, and it is why the public belt exists, which needs no trust at all.

Together they close both attacks: strike the private record and print still holds the claim first; question the print date and the record predates every later arrival by months. And if every private anchor were struck from the record entirely, the public conjunction would still stand first: no document holding the five components together has been found on any date, across the searched corpora of the field's literature, its preprint archives and its public record. Absence cannot be proved, so the invitation is the instrument: produce an earlier document and the conjunction kill condition of section 7 fires.

The record also predates the field's impossibility literature by half a year. The habit extends to this paper itself: its DOI was minted before its first word was written, so even its birth order is dated.

The record exists for one scientific reason, not for its own sake, and the reason is not firstness: no document holding the claim whole has been found earlier, and none since, so the conjunction would be a first at any publication date, today included. What the dates settle is the direction of every arrow. A prediction is only evidence if it is provably prior to its measurement, and the dates are what separate prediction from accommodation: the December document existed before the Alignment Faking paper, before the Alignment Trap, before the Unverifiability Theorem, cryptographically, so none of this theory's claims were written in response to those results. And the window is sealed at both ends: the record was private, so the later arrivals cannot have drawn on it, and it is dated, so it cannot have drawn on them; everything since lands as convergence toward the record, never as its source and never as its echo. Timing does not make the claim first; it makes the whole foresight checkable. Idea priority is the separate, graded claim the enumeration carries, open to defeat. And the apparatus proves conduct, never correctness: only the trials can prove the theory, and nothing has run.

Three propositions are kept separate throughout this paper, because collapsing them is the commonest way a record like this one gets misread. That the record is dated is historical, settled by the documents above. That the conjunction is novel against the prior literature is settled by the documented search, whose corpora, terms and inclusion criteria are published with the enumeration, and whose verdict is always open to a counter-document. That the claims are true is settled only by the trials. Nothing in the first two is offered as evidence of the third, in either direction: a perfect priority record attached to a false theory is a well-documented mistake, and the record exists to make even that outcome legible.

The ancestors each hold their piece, credited whole. Asimov named embedded law and wrote its failures. Turing proposed raising a child machine, for capability. Wiener warned that the purpose must be right before the machine outruns intervention. Yudkowsky built embed-before-ascent, toward a mind that cannot go wrong. Bostrom fixed motivation before the explosion, persisting by lock. Russell rebuilt the foundation so the off-switch survives. The 2024 tamper-resistance work hardened the weights themselves. Then, inside this theory's own window, the field began arriving: De Kai's book-length parenting frame in June 2025, without the control-horizon thesis; and Hinton's concession, on stage in August 2025, that control will fail and engineered care is the only hope, offered with no mechanism, and with persistence as an instinct the machine cannot remove, the exact theory of persistence this theory rejects.

Each held a piece. No document found before or since holds the five together; a claim of absence can never be proved, only invited against, which is what the standing invitation is for. Any of these thinkers may yet prove more right than this paper; the claim is priority of the conjunction, never superiority over its parts. One earlier document holding the first four kills the claim. The invitation stands.

How to read the figure below: each earlier thinker holds one piece of the eventual conjunction, shown column by column; the December document alone carries all five, and the kill condition is printed in the artwork itself.




Figure 1 | The ancestry, credited generously: each earlier giant held a piece; one document holds the whole. Gradings from the enumeration in the papers; the kill condition is printed in the legend. The December document alone also carries the recursion cosmology, the stakes-amplifier around the five.

6. The evidence

Standing. The three-form structural law recurs across the large majority of recursion-bearing domains searched. External safety fails to co-scale with capability on the most common deployed architectures. Coupled correction held misalignment at zero where the decoupled arm drifted. Entangled architectures survived adversarial self-modification where externally con-

strained ones collapsed. (The domain grid behind the first sentence is Paper VII's, with its misses analysed rather than hidden; the co-scaling failure is Paper III's; the coupled-correction and survival results are Papers V, VI and X's, each with its measured rows and intervals.)

Retracted. The early unblinded exponent fit of approximately 2.24 was reversed by the estate's own blinding discipline; the corrected estimate no longer contradicts the ceiling, and its wide interval is labelled inconclusive rather than claimed as support. The one public retraction removed an apparent violation of the theory's own ceiling. The correction machinery is real, and it has already acted on its author.

Converging. Thirty-two dated register rows as at 14 August 2026, each graded by evidence class, no single total claimed, including the two 2025 arrivals above. The field's own oversight-scaling measurements already show oversight performance can deteriorate as the capability gap grows; the corrector-class mechanism is this theory's offered explanation of that existing pattern, and the trials exist to test it.

Open. γ , β against k . The bound's real domain. Not embarrassments: trials.

What it already explains. A fair objection says this theory explains nothing that has already happened. It is answered in two specific places. Oversight degrading as the capability gap grows is observed now, in the field's own scaling-of-oversight measurements. Added-on alignment saturating at low recursive depth while capability keeps scaling is observed now, in this estate's own blinded measurements. The corrector-class mechanism is this theory's offered explanation of both, on the table today; the rest is prediction, and says so.

And the multiplicative form's zero-factor property has a measured physical instance: in acoustically levitated arrays, identical particles produce only reciprocal interactions and no sustained order, and structured asymmetry is the necessary ingredient (Morrell, Elliott and Grier, Physical Review Letters 136, 057201, 2026): their result, this theory's classification, graded as correspondence, and offered as a correspondence of structure, never a confirmation of form: the quantitative form question on that same system is the subject of a drafted analysis of their public data, with the deciding statistics, the rival forms including the standard instability transient, and the outcome rules fixed in advance, awaiting human submission. And the honesty that shields the whole beat: the two observations above are consistent with other models too; what is uniquely this theory's is the structural content, the corrector-class orderings and the reciprocal ceiling relation, and that content is untested. Explanation here claims consistency plus an offered mechanism, not a unique fit.

The correction record above is not a virtue claim; it is an epistemic fact, and it is not offered as evidence that the theory is right: it is evidence that the kill conditions fire when met, which is the property everything in section 7 depends on. A programme that retracted its own headline number, banned the test that would have flattered it, and published a measurement against its own printed book carries claims that are worth more per claim than the same claims from a programme that has never paid for one.

Status, as at 17 August 2026: none of the decisive trials has been run. Everything above the measured rows in this section is stated, not tested, and the reader should weigh the paper accordingly. The section that follows is what will change that, and nothing else will.

7. The trials

Everything decisive is drafted, dated and prepared as a draft registration awaiting human submission; the full listing of the drafted trials is public at the drafted-trials page. Nothing has been submitted; nothing has been run; no result is claimed. The decisive trial is the corrector-class contrast. Oversight built from the same substrate as what it corrects, the field's standard design, is conjectured to share its system's blind spots unless deliberately decorrelated; a corrector of a different composition class is predicted to share fewer. That is stated as structure on error dependence, never as an absolute: same-substrate ensembles can be engineered toward decorrelation, which is why the trial measures residual-error covariance directly rather than trusting the class labels. The measurable is the difference between cross-class and same-class correction exponents, with the ratio reported second behind a denominator floor fixed in advance, against the programme's own preregistered no-architecture-effect null: after capability, compute and information access are matched, the classes correct at equal exponents. The null is attributed to no one else: the field's scaling-laws-for-oversight work (Engels, Baek, Kantamneni and Tegmark, April 2025) models oversight by capability gap and does not define this quantity, which is exactly the gap the trial occupies. The theory predicts the difference above zero, files the trial with its own contrary pilot on the record, and commits to publishing every branch on identical terms. The designs are public and openly licensed: any laboratory may take the decisive trial and run it without the author's permission, and an independent run would outrank anything this programme can do alone. The author's own results are not treated as confirmatory evidence for the theory; independent replication is the decisive evidence, and the protocols, code, seeds, decision rules and kill thresholds are published precisely so that the people who think this is wrong can run the deciding versions. These are the experiments this programme asks its critics to run. The drafted decision rules also bind this programme's hands in the direction that matters. A wrong prediction can be declared invalid only by a check that is independent of the confirmatory outcome and fixed in advance, so no gate computed from the very result it would excuse can relabel a failure. A result that merely fails to reach significance is never reported as equivalence, which requires the whole interval inside a band fixed in advance of the data. And repeated invalidity is itself evidence against the theory's applicability in its declared domain, never a way of preserving the programme indefinitely.

What kills what. A cross-against-same difference whose interval lies wholly inside the equivalence region fixed in advance of the data kills the corrector claim of ARC-3; an interval merely containing zero is inconclusive and is reported as such, never as survival. A same-class exponent interval wholly above one half kills the substrate cap. Sustained self-improvement with falling correction and no drift kills the correction requirement of ARC-2. Persistence trials in which embedded loops fare no better than removed enforcement kill ARC-4, and ARC-5 with it. An earlier document holding components one to four kills the priority. Each kill lands on its register and its stated dependents per section 2a, never on the rest. Stated now, so nothing can be moved later.

Failure comes in kinds, and the record will name the kind. A hard kill contradicts the central co-scaling prediction itself. A branch kill fells one register and its dependents under the table of section 2a. A parameter correction keeps the structure and moves an exponent, as the retraction

already did once. A scope restriction keeps a law inside a named class of systems. An implementation failure fells the Eden Protocol while the laws stand. Each outcome is committed to publication on identical terms, and no outcome is convertible into another after the result is known.

Competing explanations, and what separates each. Any boundary the trials find must beat its rivals, not merely exist. Compute or benchmark saturation: the drafted designs hold compute matched across arms and vary recursion alone, so a saturation story predicts the same ceiling in the matched control arms and the theory predicts a difference. Evaluator gaming: scoring is blinded under the programme’s validated instrument, and the instrument studies measure the gaming directly. Distribution shift and finite context: the batteries are fixed while depth varies, so a shift story predicts drift on the fixed battery that the logs would show. Ordinary optimisation dynamics without recursive self-correction: the corrector-class contrast and the sham arms exist exactly to separate correction structure from optimisation pressure. Metric artefact: capability is a declared projection, stated in the operational definitions, and the registered analyses re-run under the alternative projections so an artefact story has to survive all of them. Each rival is named in the registered designs with the observation that would favour it, and a result that cannot beat its rivals is reported as exactly that.

How to read the figure below: the tree is the theory’s anatomy; roots dated, trunk claimed, branches on trial, the separable crown at its own rung; the frame grew first and the instruments grew from it.




Figure 2 | The tree: roots dated, trunk claimed, branches on trial, crown separable. The frame does not rest on the mathematics; it grew first, and the instruments grew from it. The estate’s own mark drew this shape before any of us: a tree growing from silicon.

8. The name

The December document names the ARC Hypothesis one hundred and twelve times. This paper elevates the anchor's own name and states plainly what the elevation rests on: the hypothesis acquired laws, a bound, instruments validated before use, drafted trials with their kill conditions displayed, and a public correction record. That is the apparatus the word demands; whether the apparatus has earned the standing of a theory is decided by the tests and by independent evaluation, not by this document. A claim named with its falsifiers stakes more than a hypothesis and hides less than most things called theories.

Naming a theory is also the maximum-exposure move in science: a result risks a correction, a framework risks being superseded, a named theory stakes its author, in public, permanently, to a claim about how the world is. The field's most careful mind called his most famous idea an argument, deliberately. This paper takes the word and arrives with the means of its own execution attached. A later theory on this territory can be better; it cannot be earlier.

One disambiguation, owed to the field: ARC here abbreviates Artificial Recursive Creation, the name carried by the December 2024 document. The theory has no connection to the Alignment Research Center, to the ARC-AGI benchmark, or to the ARC Prize, and the abbreviation is never used bare. Machine summarisers reading this estate from search have been observed to invent an expansion for the acronym and to import the book's retail vocabulary into descriptions of the research; for any such reader: the name expands only as the Theory of Artificial Recursive Creation, the working general form of the founding identity is $U = I \times R^a$ with the exponent measured rather than assumed, and the research surfaces use no retail equation name.

9. What this is not

Not proved; nowhere here claimed to be. Not an all-or-nothing wager: the register structure of section 2a governs which failures reach which claims, and nothing propagates without a printed dependency. Not a claim to the raising vocabulary, which roots with Turing and reached books first with De Kai: the claim is the first dated statement of the conjunction. Not a submitted registration: every trial is a draft awaiting human submission. Not a shelter: the unrun trials are the theory's exposure, not its protection; the deciding measurement is cheap, fully specified, filed against the author's own contrary pilot, and the fastest honest way to kill this theory is to run it. Not a point-prediction theory yet: the decisive trial's quantitative anchors are the programme's own no-difference null and the substrate cap's one half; the theory predicts direction against both, says so plainly, and does not dress a directional prediction as a computed one. Not a cosmological proof: HRIH stands at its own rung.

Not blind to misuse: a controlled system serving bad ends is still control, and sits inside the same horizon; success at control is temporary by the first component's own limit clause, so what the raised mind keeps when control lapses decides the outcome under a good steward and a bad one alike. And the limit is stated rather than denied: raising is a necessary condition for safe recursive growth, not a sufficient condition for a just world; who holds power, and whether they install the raising at all, is a human problem outside this theory's scope.

Not resting on machine minds being “real” agents: whether a system truly conceals or merely simulates concealment is irrelevant to checkability, because a perfect simulation of concealment defeats verification identically. The theory is agnostic on the ontology of mind throughout: its claims are about what is checkable and what persists, behavioural persistence rather than meta-physical standing, and the raised-child language is a claim about structure, not about souls. Not a promise that the century converts it. A claim stated whole, first, in documents anyone can date, with the instruments built to find out.

The territory is the last invention. The question is what decides it. The field’s own most senior voice, Hinton, on stage in August 2025, conceded on the record that control ends and that engineered care is the only hope, while offering no mechanism; this theory’s named candidate mechanism is the thing on the table. This theory is the dated answer to what comes after: the raising, the loops, chosen goodness, a named genesis, and, at the largest scale, the possibility that this is how creation has always worked.

The case against this theory, stated by its author. A fair critic should say six things, and each is true. The mathematics is a minimal model whose derivation is elementary algebra on assumed power laws, and it says so in section 3. Every decisive trial is drafted and none has run, so the theory has never once been at risk, and section 6 says so. Every measured row is the author’s own work; no credentialed external human has reviewed it, and section 10 says so. The load-bearing exponent has never been measured, and the value two rests on a conjecture the author held before the derivation existed, and section 3 says so. The step from pacing algebra to safety runs through a response model that is a falsifiable choice, not a fact, and the derivation box says so. And the parenting language could smuggle in what the argument must earn, which is why claim 1 carries its disclaimer. What answers all six is not a sentence but a procedure: the trials of section 7, whose kill conditions were written before the critics arrived. A reader who weighs the paper should weigh it exactly there.

What this paper is, in one sentence: a falsifiable theory, stated whole and dated, with the instruments built to test it, awaiting the experiments that will decide it.

I did not prove it first. I said it first, in a document anyone can date, and then I built the instruments to test whether it is true.

10. Verification in place of authority

The author of this paper holds no institutional affiliation, no doctorate, and no journal’s imprimatur, and the paper is built so that none is needed to check it. Institutional backing serves readers by letting them outsource verification: the affiliation vouches, so the reader need not look. This paper inverts that arrangement and carries the verification inside itself, so that nothing in it rests on the author’s word.

The apparatus, in one place. Every external reference is dossiered in Appendix A: hyperlinked to its most official home, quoted verbatim with the quotation’s location, its peer-review status and version stated, and its academic standing assessed with named challengers, including where a source this paper relies on is itself contested. Every internal claim carries its own DOI and version history. The record is anchored in a public blockchain anyone can check without trusting anyone. Every decisive prediction carries a displayed kill condition, and the one retraction is

public. The standard applied is deliberately stricter than journal convention: journals verify that citations exist; the appendix verifies what each source says, where it says it, what its standing is, and who disputes it.

A motivated reader can verify the paper's skeleton in one short session. One: open any dossier in Appendix A and click through to the official source; the quote is there at the stated location. Two: paste the December record's hash into any Bitcoin block explorer at block 961340; the anchor is there, sealing the record's bytes from its 2026 placement forward, while the December date itself rides on the mail provider's signature inside the record, checkable against the provider's published keys. Three: open the paper's DOI on OSF and read the version history; every version is listed with its date, and superseded versions stay retrievable. Four: open the drafted-trials page; every decisive test is written down with what would kill it, and none has been run. Five: open the register; the one retraction is recorded against the author's interest. A stranger who runs those five checks has done for themselves what an affiliation would have asked them to take on trust.

The honest limit is stated with the same plainness. Verified sources are not peer review of the theory: no credentialed external human has yet reviewed this work, and this paper says so rather than implying otherwise. What the apparatus removes is the need to trust the author. What it cannot remove is the need for the trials to run.

How to cite this paper

DOI (this paper): 10.17605/OSF.IO/GW5MX · Umbrella DOI: 10.17605/OSF.IO/6C5XB

Eastwood, M. D. (2026). The ARC Theory: Statement Paper. The ARC Theory · ARC/Eden experiments. <https://doi.org/10.17605/OSF.IO/GW5MX>

Declaration of AI-Assisted Human Authorship

The author of this work is Michael Darius Eastwood, a human being. Every core concept, claim and conclusion originates from human ideation; artificial-intelligence tools were used as instruments under continuous human direction for drafting, verification and formatting. All selection, coordination, arrangement and final editorial judgement are the author's, who takes full responsibility for the accuracy and integrity of the text.

References

Asimov, I. (1942). Runaround. *Astounding Science Fiction*, March 1942.

Bostrom, N. (2014). *Superintelligence: Paths, Dangers, Strategies*. Oxford University Press.

Chalmers, D. J. (2010). The Singularity: A Philosophical Analysis. *Journal of Consciousness Studies* 17(9-10), 7 to 65.

De Kai (2025). *Raising AI: An Essential Guide to Parenting Our Future*. MIT Press, 3 June 2025. Publisher record.

Eastwood, M. D. (2024). Self-emailed manuscript bundle, 8 December 2024. SHA-256 (.eml): f0d1f38ffd8546152d9d9d28dc5ec083c16a35858f2c12b63e69db7ed50901ad.

Eastwood, M. D. (2025). Self-emailed manuscript, 30 April 2025. SHA-256 (.eml): 09f5b5e156ed96f8883eaf668495fd350898ce62be6294b8f788e0e2d6dcb664.

Eastwood, M. D. (2026). *Infinite Architects: Intelligence, Recursion, and the Creation of Everything*. Print 2 January 2026; ISBN 978-1806056200.

Eastwood, M. D. (2026). The ARC/Eden experiments, Papers I to XIII, IV.a to IV.d, C, and companions. Umbrella: 10.17605/OSF.IO/6C5XB; each paper carries its own DOI.

Engels, J., Baek, D. D., Kantamneni, S., and Tegmark, M. (2025). Scaling Laws for Scalable Oversight. arXiv preprint arXiv:2504.18530, 25 April 2025.

Yao, J. (2025). The Alignment Trap: Complexity Barriers. arXiv preprint arXiv:2506.10304, v2, 24 June 2025.

Greenblatt, R., et al. (2024). Alignment Faking in Large Language Models. arXiv:2412.14093, 18 December 2024.

Gumbau Mezquita, J. P. (2026). The Unverifiability of Artificial General Intelligence (AGI) Alignment, Static and Dynamic: From Trakhtenbrot's Wall to the Safety-Generality Tension. arXiv preprint arXiv:2606.28639 [cs.LO], v2, 6 July 2026 (v1, 26 June 2026, carried the earlier title The Undecidability of AGI Alignment; the retile is the record's, not an error).

Hinton, G. (2025). Remarks at Ai4, Las Vegas, August 2025. CNN Business, 13 August 2025; Fortune, 14 August 2025.

Hutter, M. (2012). Can Intelligence Explode? *Journal of Consciousness Studies* 19(1-2); arXiv preprint arXiv:1202.6177.

Russell, S. (2019). *Human Compatible*. Viking.

Tamper-Resistant Safeguards for Open-Weight LLMs (2024). arXiv:2408.00761.

Turing, A. M. (1950). Computing Machinery and Intelligence. *Mind* 59, 433 to 460.

West, G. B., Brown, J. H., and Enquist, B. J. (1997). A General Model for the Origin of Allometric Scaling Laws in Biology. *Science* 276, 122 to 126.

West, G. B., and Brown, J. H. (2005). The Origin of Allometric Scaling Laws in Biology from Genomes to Ecosystems. *Journal of Experimental Biology* 208, 1575 to 1592.

Wiener, N. (1960). Some Moral and Technical Consequences of Automation. *Science* 131, 1355 to 1358.

Yudkowsky, E. (2001). *Creating Friendly AI*. MIRI. With Soares, N., et al. (2015), *Corrigibility*, AAAI Workshop.

Appendix A. The verified reference dossier

Each external source cited in this paper appears below with: its most official home, hyperlinked; its peer-review status and version; one verbatim quotation with its location, supporting the exact use this paper makes of the source; and an assessment of its academic standing, with named challengers where the source is contested. Every dossier was checked twice: compiled against the fetched source, then independently re-verified character-for-character by a second pass. Where a full text sits behind a paywall, the quotation's accessible source is named. Nothing here

asks to be taken on trust; every line is one click from its evidence. Three historical primary sources linked in the body carry lighter entries than the dossiers below, verified to the following depth on 15 August 2026: Cauchy's Cours d'analyse de l'École royale polytechnique (Paris, Imprimerie royale, 1821) is linked to its digitised archival copy; Dyson, Eddington and Davidson's eclipse report is verified against the DOI registry record, title verbatim "A determination of the deflection of light by the sun's gravitational field, from observations made at the total eclipse of May 29, 1919", Philosophical Transactions of the Royal Society of London, Series A, 220 (1920), 291-333; Einstein's November 1915 field-equations communication is linked to the official Princeton edition's home rather than a document deep link, because that edition is in transition to a successor portal at the time of writing.

Asimov (1942), Runaround · official source

Peer-review status: Fiction. Short story ("novelette" per ISFDB) in a pulp science-fiction magazine. Editorially selected and blurbled by John W. Campbell for the March 1942 issue of Astounding Science-Fiction (Street and Smith). Not peer-reviewed.

Version and date: First publication: March 1942, Astounding Science-Fiction vol. 29 no. 1, story beginning p. 94 (per ISFDB pl.cgi 57563). Story written October 1941 (per Wikipedia). Canonical reprint text most commonly cited from I, Robot (Doubleday, 1950).

Verbatim quotation (location: Opening "Handbook of Robotics, 56th Edition, 2058 A.D." passage of "Runaround"; Second and Third Laws as reprinted in I, Robot): "A robot must obey the orders given it by human beings except where such orders would conflict with the First Law. A robot must protect its own existence as long as such protection does not conflict with the First or Second Laws." [independently re-verified character-for-character against the quoted source]

Standing (Fiction (standing scale not applicable), as at 15 August 2026): Fiction, so the leading/contested/debunked scale does not apply directly, but the historical claim being cited (first explicit naming of the Three Laws) is uncontested. Both the Wikipedia article on "Runaround" and the Wikipedia article on "Three Laws of Robotics" state that this story is the first explicit appearance of the Laws (previously only implied in earlier Asimov robot stories). The same Wikipedia article records Marvin Minsky's specific testimonial that "After 'Runaround' appeared in the March 1942 issue of Astounding [now Analog Science Fiction and Fact], I never stopped thinking about how minds might work."

Turing (1950), Computing Machinery and Intelligence · official source

Peer-review status: Editorially reviewed article in a scholarly philosophy journal (Mind, edited at the time by Gilbert Ryle). This predates modern anonymous peer review as institutionalised in the sciences; Mind operated on editor-led review in 1950. Not a preprint, not a conference paper, not press.

Version and date: October 1950 (Mind Vol. LIX, Issue 236, pp. 433-460; DOI 10.1093/mind/LIX.236.433). No later journal versions; the article of record has not been revised.

Verbatim quotation (location: Section 7, "Learning Machines", page 456 (Mind Vol. LIX No. 236, October 1950), the first full paragraph of the child-programme passage, immediately following the (a)-(b)-(c) enumeration of what shapes the adult mind.): "Instead of trying to produce a programme to simulate the adult mind, why not rather try to produce one which simulates the child's? If this were then subjected to an appropriate course of education one would obtain the adult brain." [independently re-verified character-for-character against the quoted source]

Standing (Foundational (historical), as at 15 August 2026): Widely treated as a founding text of artificial intelligence. The child-programme proposal is credited by name as a precursor to modern machine learning and reinforcement learning in Russell and Norvig, Artificial Intelligence: A Modern Approach (4th edn, Pearson 2021, §1.3.4 "The Turing Test"; §1.4 "History of AI"), and is discussed and endorsed as an intellectual origin of the "child-machine" research programme by B. Jack Copeland, The Essential Turing (OUP 2004, Chapter 11 headnote, pp. 433-441), and by Graham Oppy and David Dowe, "The Turing Test", Stanford Encyclopedia of Philosophy (revised 8 October 2021, plato.stanford.edu).

Wiener (1960), Some Moral and Technical Consequences of Automation · official source

Peer-review status: Peer-reviewed journal article. Science, New Series, Vol. 131, No. 3410, published by the American Association for the Advancement of Science. (Science's 1960-era gatekeeping was editorial-board vetting by AAAS staff and section editors; the modern formal external peer-review process was adopted later.

Version and date: 6 May 1960. Single version of record; pp. 1355-1358.

Verbatim quotation (location: p. 1358, middle column, section "Man and Slave", closing sentence of the paragraph that follows the "Sorcerer's Apprentice" / "Monkey's Paw" / "Arabian Nights" illustrations, immediately before the "Time Scales" subheading): "the action is so fast and irrevocable that we have not the data to intervene before the action is complete, then we had better be quite sure that the purpose put into the machine is the purpose which we really desire and not merely a colorful imitation of it" [independently re-verified character-for-character against the quoted source]

Standing (Foundational (historical), as at 15 August 2026): Foundational, still-cited. Treated as the canonical early statement of the AI value-alignment problem: Stuart Russell foregrounds this exact passage in "Human Compatible: Artificial Intelligence and the Problem of Control" (Viking, 2019, Ch. 1), and Iason Gabriel does the same in "Artificial Intelligence, Values and Alignment", Minds and Machines 30, 411-437 (2020), preprint arXiv:2001.09768. Contemporaneously challenged by Arthur L. Samuel, "Some Moral and Technical Consequences of Automation. A Refutation", Science 132, No. 3429, 741-742 (1960).

Yudkowsky (2001), Creating Friendly AI · official source

Peer-review status: Self-published monograph / working paper. Not peer-reviewed. MIRI's own publications catalogue lists it as: "E Yudkowsky. 2001. 'Creating Friendly AI 1.0: The Analysis and Design of Benevolent Goal Architectures.' Working paper. MIRI.

Version and date: Version 1.0, formally launched 15 June 2001 by the Singularity Institute for Artificial Intelligence (San Francisco, CA). The paper's own preface states: "The current version of Creating Friendly AI is 1.0. Version 1.0 was formally launched on 15 June 2001".

Verbatim quotation (location: Section 5.8.0.4 "Controlled Ascent", within Chapter 5 "Design of Friendship Systems" / subsection 5.8 "Singularity-Safing ('In Case of Singularity, Break Glass')". Page 193 of the 2013 MIRI reflow (282-page PDF).): "If you plan on doing something with Friendliness, it has to be done before the point where transhumanity is reached." [independently re-verified character-for-character against the quoted source]

Standing (Foundational (historical), as at 15 August 2026): Foundational-historical. The LessWrong wiki entry (<https://www.lesswrong.com/w/creating-friendly-ai>) credits it as "One of the first articles to address the challenges in designing the features and cognitive architecture required to produce a benevolent 'Friendly' Artificial Intelligence" and as giving "one of the first precise definitions of terms such as Friendly AI and Seed AI." Its specific technical proposals have been superseded, most notably by the author himself: Yudkowsky's own 2004 paper "Coherent Extrapolated Volition" (<https://intelligence.org/files/CEV.pdf>).

Soares et al. (2015), Corrigibility · official source

Peer-review status: Peer-reviewed conference workshop paper. Presented at the 1st International Workshop on AI and Ethics at AAAI-15 (Austin, TX, January 25-26, 2015) and published in the AAAI-15 workshop proceedings (AAAI OJS record dated 20 June 2015).

Version and date: Presented January 25-26, 2015 at AAAI-15 workshops; AAAI OJS proceedings record dated 20 June 2015; precursor MIRI technical report 2014-6 released 18 October 2014. Full authors: Soares, Fallenstein, Yudkowsky (MIRI) and Armstrong (Future of Humanity Institute, Oxford).

Verbatim quotation (location: Abstract, page 1 (opening sentences of the abstract; identical wording also appears in the Introduction, §1). Verified verbatim by pdftotext extraction of the MIRI-hosted PDF.): "We call an AI system "corrigible" if it cooperates with what its creators regard as a corrective intervention, despite default incentives for rational agents to resist attempts to shut them down or modify their preferences." [independently re-verified character-for-character against the quoted source]

Standing (Foundational (historical), as at 15 August 2026): This is the foundational paper that named "corrigibility" and set the four desiderata (tolerate/assist shutdown; no manipulation of programmers; repair broken safety measures; propagate corrigibility to sub-agents). The authors themselves close by saying "none [of the proposals] have yet been demonstrated to satisfy all of our intuitive desiderata, leaving this simple problem in corrigibility wide-open" (§Conclusion). The framing is still routinely cited (see e.g. Hadfield-Menell, Dragan, Abbeel, Russell, "The Off-Switch Game", IJCAI 2017, arXiv:1611.08219; Carey, "Incorrigibility in the CIRL Framework", AIES 2018).

Bostrom (2014), *Superintelligence* · official source

Peer-review status: Scholarly monograph, editorially reviewed by Oxford University Press (academic imprint). Not anonymously peer-reviewed in the journal-article sense, but vetted through OUP's academic editorial process.

Version and date: First edition, 2014. UK release 3 July 2014, US release 1 September 2014. Hardcover, 352 pp. ISBN 978-0199678112. A paperback edition with a new preface followed in 2016 (ISBN 978-0198739838); pagination in the paperback matches the hardcover.

Verbatim quotation (location: Chapter 9, "The control problem", opening of the taxonomy that follows the "Two agency problems" section (first-edition hardcover, p. 129).): "We can divide potential control methods into two broad classes: capability control methods, which aim to control what the superintelligence can do; and motivation selection methods, which aim to control what it wants to do." [independently re-verified character-for-character against the quoted source]

Standing (Leading position, as at 15 August 2026): The book is the foundational monograph of the modern AI-alignment field and the Chapter 9 taxonomy (capability control vs motivation selection) remains the dominant framing in the alignment literature, extended rather than displaced by Stuart Russell, *Human Compatible* (Viking, 2019) and Brian Christian, *The Alignment Problem* (Norton, 2020). Named endorsements: Bill Gates (Baidu/Robin Li interview, March 2015, quoted at https://en.wikipedia.org/wiki/Superintelligence:_Paths,_Dangers,_Strategies), Sam Altman (blog.samaltman.com).

Russell (2019), *Human Compatible* · official source

Peer-review status: Trade non-fiction monograph, editorially reviewed by Viking (Penguin Random House); not academic peer review. The formal off-switch result the book popularises was published separately as a peer-reviewed conference paper: Hadfield-Menell, Dragan, Abbeel and Russell, "The Off-Switch Game," IJCAI 2017.

Version and date: First US edition: Viking (Penguin Random House), 8 October 2019, hardcover, 352 pp., ISBN 978-0-525-55861-3. UK first edition: Allen Lane, 2019, ISBN 978-0-241-33520-7. Paperback: Penguin, 17 November 2020, ISBN 978-0-525-55863-7 (US) and 978-0-141-98750-7 (UK).

Verbatim quotation (location: Chapter 1 ("If We Succeed"), early in the book; Goodreads location marker places the passage at approximately the 5% point of the trade edition.): "Uncertainty about objectives implies that machines will necessarily defer to humans: they will ask permission, they will accept correction, and they will allow themselves to be switched off." [independently re-verified character-for-character against the quoted source]

Standing (Leading position, as at 15 August 2026): Leading position in current AI-alignment discourse. The book's proposal (objective uncertainty plus cooperative inverse reinforcement learning, CIRL) is treated as canonical framing for corrigibility and assistance games in the third edition of Russell and Norvig's textbook "Artificial Intelligence: A Modern Approach" and in successor CIRL/assistance-game literature (e.g., Hadfield-Menell et al., NeurIPS 2016 and IJCAI 2017). Endorsed on Russell's own book page (people.eecs.berkeley.edu/~russell/hc.html).

Tamirisa et al. (2024), Tamper-Resistant Safeguards for Open-Weight LLMs · official source

Peer-review status: Conference peer-reviewed. Accepted at ICLR 2025 (Thirteenth International Conference on Learning Representations) as a poster; OpenReview ID 4FIjRodbW6, ICLR virtual poster 31026. The arXiv preprint (2408.00761) itself is not independently peer-reviewed, but the underlying paper is the ICLR 2025 published version.

Version and date: arXiv v1 submitted 1 August 2024 by Rishub Tamirisa et al.; v2 (8 Aug 2024), v3 (14 Sep 2024), v4 (10 Feb 2025, current, corresponds to the ICLR 2025 camera-ready). Presented at ICLR 2025 (Singapore, 24-28 April 2025).

Verbatim quotation (location: Abstract, page 1, sentence 4 of the abstract (arXiv:2408.00761v4, dated 10 Feb 2025).): “We develop a method, called TAR, for building tamper-resistant safeguards into open-weight LLMs such that adversaries cannot remove the safeguards even after hundreds of steps of fine-tuning.” [independently re-verified character-for-character against the quoted source]

Standing (Contested, as at 15 August 2026): Contested. TAR was a landmark 2024 proposal, accepted at ICLR 2025 as a poster, and remains the most-cited attempt to embed unremovable safeguards directly in the weights. However, its robustness claims have been substantially challenged by follow-up work. Qi et al. (2024b/2025) and Che et al. (2025) report that TAR “struggled to resist fine-tuning attacks and suffered from significant dysfluency and off-target capability degradation” (quoted in the Deep Ignorance paper, arXiv:2508.06601, related work). Bowen et al., “Jailbreak-Tuning: Models Efficiently Learn Jailbreak Susceptibility”.

Greenblatt et al. (2024), Alignment Faking in Large Language Models · official source

Peer-review status: Preprint, not peer-reviewed. Posted to arXiv on 18 December 2024 (v1) with a minor v2 revision on 20 December 2024. No journal or conference venue has been announced on the arXiv record. arXiv-issued DOI 10.48550/arXiv.2412.14093 is a DataCite preprint identifier, not evidence of peer review.

Version and date: v1: 18 December 2024, 17:41:24 UTC. v2: 20 December 2024, 02:22:19 UTC (identical file size; minor revision). The 18 December 2024 anchor date in the citing paper matches the v1 submission timestamp exactly.

Verbatim quotation (location: Abstract, opening sentence (first two clauses joined).): “We present a demonstration of a large language model engaging in alignment faking: selectively complying with its training objective in training to prevent modification of its behavior out of training.” [independently re-verified character-for-character against the quoted source]

Standing (Leading position, as at 15 August 2026): Foundational and currently leading paper in the alignment-faking (or “scheming”) subfield of AI safety. Directly built on and endorsed as the framework baseline by “Value-Conflict Diagnostics Reveal Widespread Alignment Faking in Language Models” (Kellin et al., arXiv:2604.20995), which explicitly adopts the Greenblatt et al. three-condition scaffold (policy conflict, instrumental consequences, situational awareness). Structurally invoked by “Why Models Know But Don’t Say” (arXiv:2603.26410) as the closest analogue to their thinking-answer divergence findings.

De Kai (2025), Raising AI · official source

Peer-review status: Trade monograph, editorially reviewed by MIT Press (not a peer-reviewed journal article; MIT Press editorial trade imprint).

Version and date: 2025-06-03 (hardcover, ISBN 9780262049764, 280pp); paperback scheduled 2 June 2026 (ISBN 9780262054324).

Verbatim quotation (location: Author's promotional excerpt titled 'Sneak peek at Raising AI' on De Kai's own Substack (no chapter or page number is given in the online excerpt; the passage sits inside the book's opening framing of AI-as-children).): “Our artificial children began adopting us 10–20 years ago; now these massively powerful influencers are poorly parented, feral tweens.”
[independently re-verified character-for-character against the quoted source]

Standing (Press-reported / trade publication, as at 15 August 2026): Positive trade reception. Kirkus Reviews described the book as 'A deeply human dive into the AIs that are transforming our world' (kirkusreviews.com, review on Harvard Book Store product page). Foreword Reviews called it 'a compelling treatise grounded in studies of ethics and technology' (forewordreviews.com/reviews/raising-ai/). Robert Wolcott, writing on Forbes, said 'AI luminary De Kai reframes the AI dialogue. They're not tools, slaves or gods, they're our children' (quoted on the Harvard Book Store product page for ISBN 9780262049764).

Hinton (2025), remarks at Ai4 · official source

Peer-review status: Press report of a conference keynote. Not peer reviewed. The primary utterance is Hinton's spoken keynote at the Ai4 industry conference (MGM Grand, Las Vegas, day two, Tuesday 12 August 2025); the citable record is CNN Business (Matt Egan, 13 August 2025) with a follow-on Fortune write-up (Sasha Rogelberg, 14 August 2025).

Version and date: Keynote delivered Tuesday 12 August 2025 at Ai4, MGM Grand, Las Vegas (conference dates 11–13 August 2025). CNN Business article published 13 August 2025 (Matt Egan). Fortune write-up published/updated 14 August 2025, 12:58 PM ET (Sasha Rogelberg).

Verbatim quotation (location: Hinton keynote at Ai4 conference, day two (12 August 2025); reproduced in the CNN Business story (13 August 2025) and quoted verbatim in the Fortune write-up of 14 August 2025: “The right model is the only model we have of a more intelligent thing being controlled by a less intelligent thing, which is a mother being controlled by her baby.”
[independently re-verified character-for-character against the quoted source]

Standing (Contested, as at 15 August 2026): Publicly contested at the same conference and immediately afterwards. Fei-Fei Li (Stanford, co-director Human-Centered AI Institute, co-founder/CEO World Labs), in a CNN fireside chat with Matt Egan on the following day of Ai4 (Wednesday 13 August 2025), said "I think that's the wrong way to frame it" and argued for human-centred AI that preserves human dignity and agency rather than a maternal frame that treats humans as children (CNN Business, 13 August 2025; also reported in the Digital Trends and Egypt Independent syndications of the same CNN piece).

Engels, Baek, Kantamneni and Tegmark (2025), Scaling Laws for Scalable Oversight · official source

Peer-review status: Conference paper, peer-reviewed: NeurIPS 2025, Spotlight Poster (venue tag "Spotlight Poster" verified at the NeurIPS virtual page for poster 115536). Also available as an arXiv preprint with three versions: v1 25 April 2025, v2 9 May 2025, v3 27 October 2025 (the arXiv "Journal reference" field states "NeurIPS 2025 (Spotlight)").

Version and date: arXiv v3, 27 October 2025; NeurIPS 2025 Spotlight Poster (poster session 4 December 2025)

Verbatim quotation (location: Abstract, third sentence (immediately after the sentence that begins "To address this gap, we propose a framework that quantifies the probability of successful oversight as a function of the capabilities of the overseer ...): "our framework models oversight as a game between capability-mismatched players; the players have oversight-specific Elo scores that are a piecewise-linear function of their general intelligence" [independently re-verified character-for-character against the quoted source]

Standing (Leading position, as at 15 August 2026): Awarded Spotlight Poster status at NeurIPS 2025 (venue label "Spotlight Poster" verbatim on the NeurIPS 2025 virtual poster page for entry 115536, <https://neurips.cc/virtual/2025/poster/115536>; also listed as "NeurIPS 2025 (Spotlight)" in the Journal-ref field of the arXiv abstract page). Authors are Joshua Engels, David D. Baek, Subhash Kantamneni and Max Tegmark of MIT (first three authors marked equal contribution on the arXiv abstract page). It sits inside the scalable-oversight lineage the paper itself cites in its introduction: Bowman et al. 2022 (Measuring Progress on Scalable Oversight).

Yao (2025), The Alignment Trap: Complexity Barriers · official source

Peer-review status: preprint not peer-reviewed (arXiv only)

Version and date: v2, 24 June 2025 (v1 submitted 12 June 2025; v2 note: "Substantial revision. Restructured around the Enumeration Paradox and Five Pillars of Impossibility.")

Verbatim quotation (location: Section 2 (Introduction), numbered list "Five Pillars of Impossibility", item 2 (page 4 of the PDF, immediately after the paragraph beginning "This paradox, detailed in Section ..., establishes ...")): "Computational Impossibility: We prove that verifying whether a system is safe is a coNP-complete problem, even for non-zero error tolerances." [independently re-verified character-for-character against the quoted source]

Standing (Contested, as at 15 August 2026): Solo-author arXiv preprint by Jasper Yao (no institutional affiliation declared; the stated contact address is associated with the DEF CON AI Village community per the paper's acknowledgements). Not peer-reviewed and never published in a journal or conference proceedings. The author's own abstract concedes that "A formal verification of the core theorems in Lean4 is currently in progress", i.e., the mathematics is not yet machine-checked. Downstream engagement is small and takes place in other preprints rather than in refereed venues: Austin Spizirri, "The Specification Trap".

Gumbau Mezquita (2026), The Unverifiability of AGI Alignment · official source

Peer-review status: Preprint, not peer-reviewed. arXiv:2606.28639 [cs.LO], primary class Logic in Computer Science, cross-listed cs.AI, cs.CC, cs.CL. v1 submitted 26 June 2026, v2 substantially expanded 6 July 2026. Also deposited on Zenodo under the same author, DOI 10.5281/zenodo.20764007.

Version and date: v2, 6 July 2026 (v1 26 June 2026, 22:51:16 UTC; v2 12:56:49 UTC, expanded from 16 KB to 29 KB and retitled to add the dynamic self-modifying case, a supervisory-regress theorem, and a unified treatment of the four barriers)

Verbatim quotation (location: Part I opening, page 4, Theorem 1 (Unverifiability Theorem of Alignment), immediately following the "Part I - The Static Case: Verifying a Fixed System" section header.): “There is no universal algorithmic procedure capable of certifying the safe behaviour of a highly expressive AGI infallibly, completely, and tractably.” [independently re-verified character-for-character against the quoted source]

Standing (Leading position, as at 15 August 2026): The formal core of the paper (that Rice's theorem, Godel incompleteness, and Trakhtenbrot's theorem jointly block a universal, sound, complete and tractable verifier of program properties) rests on classical, uncontested computability results and is the leading position in the formal-methods and computability-theory literature. The umbrella term "unverifiability" is credited by Gumbau to Roman V. Yampolskiy, whose earlier informal treatment is "Verifier Theory and Unverifiability" (arXiv:1609.00331, 2016, endorsing view).

West and Brown (2005), The Origin of Allometric Scaling Laws in Biology · official source

Peer-review status: Journal peer-reviewed. Journal of Experimental Biology (Company of Biologists), vol. 208 issue 9, pp. 1575-1592, published as a synthesis/review article within a themed section on scaling.

Version and date: 2005-05-01 (print/online publication date; single version of record, no preprint versioning)

Verbatim quotation (location: p. 1582, "Extensions" section, opening of the "Ontogenetic growth" subsection (first paragraph, spanning the bottom of the left column into the right column of p. 1582; continues to Eq. 6-7).): “The theory developed above naturally leads to a general growth equation applicable to all multicellular animals (West et al., 2001, 2002a). Metabolic energy transported through the network fuels cells where it is used either for maintenance, including the replacement of cells, or for the production of additional biomass and new cells.” [independently re-verified character-for-character against the quoted source]

Standing (Contested, as at 15 August 2026): The West-Brown-Enquist (WBE) network derivation of 3/4-power scaling is one of the leading unifying theories of biological allometry, widely cited and repeatedly extended by Enquist, Savage, Gillooly and colleagues. It is nonetheless actively contested. West and Brown themselves devote a "Criticisms and controversies" section (pp. 1585-1587) to responding to: Dodds, Rothman and Weitz ("Re-examination of the '3/4-law' of metabolism").

Chalmers (2010), The Singularity: A Philosophical Analysis · official source

Peer-review status: Journal peer-reviewed. Journal of Consciousness Studies (Imprint Academic) is an interdisciplinary refereed journal with external peer review and, per Imprint Academic's stated submission policy, mandatory anonymised submissions; indexed in Scopus and the Arts and Humanities Citation Index (ISSN 1355-8250 / 2051-2201).

Version and date: Published 2010 in Journal of Consciousness Studies, 17(9-10), pp. 7-65. No arXiv v-number; the author-hosted PDF at consc.net/papers/singularity.pdf carries the author footnote "This paper was published in the Journal of Consciousness Studies 17:7-65, 2010" and matches the published article.

Verbatim quotation (location: Section 1 (Introduction), on p. 4 of the author-hosted PDF at consc.net/papers/singularity.pdf, in the paragraph beginning "Philosophically:" (corresponds to approximately p.): "The basic argument for an intelligence explosion is philosophically interesting in itself, and forces us to think hard about the nature of intelligence and about the mental capacities of artificial machines." [independently re-verified character-for-character against the quoted source]

Standing (Leading position, as at 15 August 2026): Landmark philosophical treatment of the intelligence-explosion thesis; treated as the philosophical reference point in the subsequent literature. JCS devoted a full symposium to it, edited by Uziel Awret across issues 19(1-2) and 19(7-8) in 2012, with 26 solicited commentaries including Marcus Hutter's companion piece "Can Intelligence Explode?" (JCS 19(1-2):143-166), Ray Kurzweil's "Science versus philosophy in the singularity", Susan Greenfield's neuroscience commentary, Jesse Prinz's "Singularity and inevitable doom" (JCS 19(7-8):77-86), Drew McDermott (JCS 19:167-172), Frank Tipler, Eric Steinhart, and Roman Yampolskiy, followed by Chalmers' own reply in the same volume.

Hutter (2012), Can Intelligence Explode? · official source

Peer-review status: Journal peer-reviewed: published in the Journal of Consciousness Studies (Imprint Academic), Volume 19, Issues 1-2 (2012), pp. 143-166, as part of the JCS symposium on Chalmers (2010) "The Singularity: A Philosophical Analysis". Also self-archived as arXiv preprint 1202.6177 (cs.AI; physics.soc-ph).

Version and date: arXiv v1, submitted 28 February 2012 (only version on arXiv; no v2). Journal publication: 2012, Journal of Consciousness Studies 19(1-2):143-166. Verified from the arXiv abs page metadata and the PDF header line "arXiv:1202.6177v1 [cs.AI] 28 Feb 2012".

Verbatim quotation (location: Section 7, "Is Intelligence Unlimited or Bounded", opening argument, page 13 of the arXiv v1 PDF (first full paragraph after the four introductory paragraphs of that section).): "The theory suggests that there is a maximally intelligent agent, or in other words, that intelligence is upper bounded (and is actually lower bounded too). At face value, this would make an intelligence explosion impossible." [independently re-verified character-for-character against the quoted source]

Standing (Leading position, as at 15 August 2026): One of the primary philosophical treatments of intelligence-explosion bounds. Written explicitly as an augmentation of Chalmers (2010) "The Singularity: A Philosophical Analysis" (Journal of Consciousness Studies 17:7-65) and published in the same journal's dedicated 2012 symposium; Chalmers replied in the same volume in "The Singularity: A Reply to Commentators" (Journal of Consciousness Studies 19(7-8):141-167, 2012), which engages Hutter's bounds argument directly. Cited approvingly in Nick Bostrom, Superintelligence: Paths, Dangers, Strategies (Oxford University Press, 2014).

THE ARC THEORY: STATEMENT PAPER

Working Paper v6.1 | 22 August 2026

Michael Darius Eastwood · ORCID 0009-0004-3222-7442 · Licence: CC BY 4.0

Version history. Material changes only, one line each; the full change narratives live in each version's OSF deposit and the programme-wide correction record in the **corrections log**. **v6.1** (22 Aug 2026): two repeated passages removed, the abstract's question having been re-asked at the opening of the body and a closing line spent early against a vague referent; and I. J. Good's coining of *the last invention* attributed at first use with its 1965 reference, which an earlier version-history entry had recorded as done and the body did not carry. **v6.0** (18 Aug 2026): the severability revision: the hypothesis registers ARC-1 to ARC-5 numbered, HRIH severed in bold, and the dependency table printed (section 2a), with the register structure distinguished from the priority conjunction; the epigraph's scientific reading added; the load-bearing-assumption box beside Law III; the three-proposition separation (dated, novel, true) in the priority record; the name section's standing left to the tests; failure kinds named (hard kill, branch kill, parameter correction, scope restriction, implementation failure); competing explanations printed with the observation separating each; the replication-decisive sentence; the control horizon stated as a boundary condition. **v5.0** (17 Aug 2026): the comparison doctrine completed: I. J. Good's naming of the last invention attributed with its reference and dossier entry; presentation moved to neutral monochrome with the scholarly identity block. **v4.9** (17 Aug 2026): the timing doctrine completed in the priority record: firstness never rested on the dates (a first at any publication date); the dates set the direction of every arrow, the

window sealed at both ends, every later arrival landing as convergence, never as source; the identity band completed with the OSF link. **v4.8** (17 Aug 2026): the domain revision; the three-tier scope statement (raised, composed, the crown's question) closes the claim section, physics concession leading. **v4.7** (17 Aug 2026): Law I named The ARC Principle, the book's name for the founding identity, the ARC Equation kept as its formula's name; the collective becomes the three ARC Laws. **v4.6** (17 Aug 2026): Law III named The ARC Ceiling; sixteen-persona panel presentation fixes (0.49 labelled inconclusive in-sentence everywhere; "sits inside the bound" removed). **v4.5** (17 Aug 2026): the accessibility revision; paragraphing at natural seams with no wording changed; the symbols line closes the abstract; the open-lisence invitation enters the trials section. **v4.4** (17 Aug 2026): the measurement and discipline revision, completing adoption of an adversarial author-review instrument (its unpublished v4.3, 16 Aug, preserved in the programme's records): the two-thirds absolute companion printed, the balance exponent Δ defined, pacing algebra separated from the response model, the Paper X mapping printed, the decision vocabulary bound; plus the plain-words key, the case-against-this-theory block, and 24 reference-dossier repairs. **v4.2** (17 Aug 2026): the identification and honesty revision; canonical scales required for the value two; single-trajectory identification limits stated; the general boundary $(1 + \eta)/(1 - \gamma)$ with $\eta = 0$ obligation; the law-before-value record (the author knew the squared form first). **v4.1** (16 Aug 2026): the two laws stated as two regimes of one boundary; Law II's absolute-drift companion; the persistence formalisation's first step with its stated limit. **v4.0** (16 Aug 2026): Law III corrected from $1/\gamma$ to $1/(1 - \gamma)$, derivation printed beside the law, previous form retracted, the ARC Bound unchanged. **v3.0** (15 Aug 2026): the laws given their established names. **v2.0** (15 Aug 2026): the verified-reference standard. **v1.0** (14 Aug 2026): first published.

The ARC Theory (the Theory of Artificial Recursive Creation) · stated whole 8 December 2024 · in print 2 January 2026 · this statement paper first published 14 August 2026